

# Scientific Computation

An Out-of-Order History of Experimental Science

Bill Cochran

[wkcochran@gmail.com](mailto:wkcochran@gmail.com)

<https://github.com/wkcochran123/measurement>

July 11, 2026

# Preface: The Freedom to Answer

A question and an answer are not the same freedom.

You may choose where to put the apparatus, which dial to turn, which basis to use, which symbol to write at the top of the page. None of that entitles you to choose what the apparatus says back. The first freedom belongs to the reader of the world: the freedom to ask in one way rather than another. The second belongs, if it belongs anywhere, to the thing being read: the freedom to answer from its own state rather than echo the name placed upon it. Confuse those two and calibration becomes ventriloquism. The instrument appears to speak, but the voice coming out of it is yours.

This volume is about the distance between those freedoms. It is the fourth reading of one machine, the reading gauge and user guide. The first three volumes build and turn the instrument in the languages of construction, physics, and computation. This one follows the source, runs the reports, and asks what licenses the names by which those reports are read. It is an out-of-order history of experimental science because the order here is not the order in which ideas were discovered. It is the order in which an instrument needs them: first a question, then a selectable way to pose it, then an answer, then a record that can survive being read again.

Four words stand at that threshold: *free will*, *the axiom of choice*, *witness*, and *observation*. They are often made to blur into one another. Choice is

called freedom; seeing is called witnessing; an answer is treated as though it had been selected; a selected name is treated as though it were already an observed fact. The machine needs the four kept apart, and it needs them arranged on two axes. One axis separates a *requirement* from an *action*: what must be possible from what is actually done. The other separates the contravariant from the covariant: what comes back against a reading from what travels forward with the thing read. The whole preface fits in one small table:

|               | requirement     | action      |
|---------------|-----------------|-------------|
| contravariant | axiom of choice | witness     |
| covariant     | free will       | observation |

The table is not a proof about human freedom, and it is not a new axiom smuggled into the code. It is the reading this volume lays over the calibration. Its value is disciplinary: it tells us who owes what at every stage, and prevents a requirement from masquerading as its own fulfillment.

Begin with the contravariant requirement. Before an experiment can return a value, a requested reading determines a family of setups under which a value would count as its answer. The requirement is that this family contain at least one admissible member: a frame must be choosable, a representative nameable, a detector orientable, a coordinate available to be fixed. Change the object one way and the coordinate description must change against it so that their pairing continues to mean the same thing. That backward demand — from the claim we would like to state to the admissible setup that could justify it — is contravariant.

Promote that demand far enough and it becomes the axiom of choice. Given a family of nonempty collections, choice says that a representative may be selected from each even when no common rule tells us which representative to take. Zermelo stated the axiom in the course of proving the well-ordering theorem [50]; Gödel showed its relative consistency [24]; Cohen made forcing

into a working method for independence proofs [11]. The controversy was never over whether a person can point at one apple. It was over whether the existence of a selector may be granted globally when the selecting act has not been exhibited. In the language of this instrument, the axiom of choice is therefore a *requirement*: let there be a way to choose the representatives the reading demands. It is an open invoice, not a receipt.

The receipt is the witness. A witness is the contravariant action: not the promise that a representative exists, but the representative produced; not “there is a path,” but this path; not “some value satisfies the condition,” but this value together with the check that it does. Constructive type theory makes the distinction unusually hard to evade, because to inhabit an existential claim is to carry its witness with it [37]. A requirement points back from a proposed conclusion toward what would justify it. A witness walks that route and comes back carrying the justification.

This is why a witness bears a cost. The cost need not be mystical and it need not be money. In this device it is the finite work of constructing, elaborating, and checking the term that fills the required place. An axiom can say, in one stroke, that a selector exists. A witness must perform the selection in the case at hand. The former licenses; the latter acts. The pullback of the requirement is the literal contravariant operation. The witness is its contravariant action because it travels against the observation arrow to fill that pulled-back demand with a particular handle.

Now turn the table over. Free will, as the phrase is used here, is not the observer’s privilege of picking a coordinate. That would put it back in the contravariant column and make it merely another name for choice. Free will is the *covariant requirement*: after the question has been chosen, the system must still be able to answer. Its response must travel with its state and with the admissible change of circumstances; it must not have been filled in by the label chosen to solicit it.

This is a narrower claim than the old quarrel between necessity and lib-

erty, though it touches that quarrel. Hume asked how freedom could be made intelligible in a world of causes [30]; Kant placed freedom at the fault line between the world as described and the agency by which a description is taken up [31]. The operational answer offered here is modest. Freedom does not mean lawlessness, randomness, or exemption from cause. A compass needle is not free *from* the field; it is free to answer the field rather than the word painted beneath it. Covariance is the form of that lawful freedom. Change the admissible frame and the response changes with it according to a rule. What must remain free is not the response from law, but the response from having been decided in advance by the questioner's notation.

The covariant action is observation. Observation is not a private glance and not the passive discovery of a value that was already printed somewhere behind the apparatus. It is the act by which a system's response is gathered into a record. Bohr insisted that a quantum phenomenon includes the conditions under which it is observed [4]; von Neumann made the chain from system through apparatus to record mathematically explicit [41]. This instrument keeps the same useful caution without asking those theories to prove its construction. The chosen setup may condition or disturb the response, but it must not hard-code or copy the requested report. Observation carries the answer forward from the system into the ledger. That forward carriage is why the action is covariant.

The four terms can be written more compactly. Let  $S$  be the states the instrument might encounter and  $R$  the readings it can report. Observation is a forward map

$$O : S \longrightarrow R.$$

A requested reading pulls backward to the states that could witness it. Write  $\approx_\varepsilon$  for agreement within the device's owned floor, and define the admissible fiber and the claimable readings by

$$F_r^\varepsilon := \{s \in S : O(s) \approx_\varepsilon r\}, \quad R_{\text{adm}} := \{r \in R : F_r^\varepsilon \neq \emptyset\}.$$

Choice cannot make an empty fiber nonempty. Its requirement is only that representatives can be selected from the nonempty fibers. A global selector — an approximate section — would be a map  $W : R_{\text{adm}} \rightarrow S$  with  $W(r) \in F_r^\varepsilon$ . The contravariant action in a particular case is the local witness  $w_r := W(r)$ . Calibration then asks, point by point, for

$$O(w_r) \approx_\varepsilon r, \quad r \in R_{\text{adm}}.$$

Choose a representative, run the observation, and recover the claimed reading within the instrument's owned floor. In a finite case the device may construct the needed local witness without invoking a global choice principle. Notice what calibration does *not* ask for:  $W \circ O = \text{id}_S$ . A successful name need not reconstruct the whole world. The label is a usable handle, not the object exhausted. That asymmetry is the space in which both measurement and humility live.

One pass, however, is not enough. On one pass the selected representative and the returned reading may agree merely because the selection planted the answer. The proposed protocol therefore calibrates twice. It is the bridge now being implemented; until a build report carries it, its status is design rather than measured result. At either pass, two of the three values are returned directly as settled candidates while the third is carried as a superposition of the other two. The third is not false and it is not missing. It is present relationally rather than pointwise.

Write the three roles as  $a, b, c$ , and let  $\Sigma_i(x, y)$  denote the formal composite carried on pass  $i$ . The word *superposition* means only this declared calibration relation. It does not, by itself, mean ordinary addition, a linear or quantum state, or a mechanism behind a physical process. One admissible rotation looks schematically like

$$\begin{aligned}
t_1 : \quad & a_1, b_1 \text{ returned directly,} & \Sigma_1(a_1, b_1) \text{ carried,} \\
t_2 : \quad & b_2, c_2 \text{ returned directly,} & \Sigma_2(b_2, c_2) \text{ carried.}
\end{aligned}$$

Let  $\rho_{12}$  and  $\rho_{21}$  be the transition rules declared before either target is read. The shared role  $b$  calibrates the overlap only if

$$b_1 \approx_\varepsilon b_2, \quad \rho_{12}(\Sigma_1(a_1, b_1)) \approx_\varepsilon c_2, \quad \rho_{21}(\Sigma_2(b_2, c_2)) \approx_\varepsilon a_1.$$

The round trip must also close within the same owned floor. Only after those checks do the direct candidates earn the word *correct*. The particular rotation may change, but the discipline may not: the unresolved place must change hands, the passes must overlap, and the preregistered transition law must survive the exchange. If the same slot remains unresolved twice, nothing new has been learned. If the common slot does not agree, the frames have not been calibrated. If the transition changes merely to preserve a favored label, the name has been fitted rather than inferred.

This is what allows the machine to use names. A name does not assert that all three values were simultaneously sharp. It names the role that persists when sharpness circulates among them. On the first pass the witness fixes a frame and the observation returns two direct values plus one relation. On the second pass a new witness changes which value is held relationally, and a new observation tests whether the old relation survives. What earns the name is the invariant across the two passes: the same transition law, recovered when the direct and superposed roles exchange. The name is contravariant — ours, frame-bound, a handle chosen against the change. What licenses the name is covariant — the world's answer, carried through the change and returned repeatably.

The two-pass calibration therefore closes the table into a loop. The contravariant requirement asks for representatives from the stipulated nonempty fibers; the axiom grants a selector globally, while the witness supplies a representative locally. Free will asks whether the selected state can have a

consequence computed from that state rather than stipulated by its name. Observation supplies that consequence. The second pass then turns the consequence into a new requirement and asks the loop to survive once more. Requirement becomes action; action becomes the next requirement. Contravariant question and covariant answer meet in a report stable up to  $\approx_\epsilon$  — for this instrument, the repeatable bracket it owns.

That loop also explains how the four volumes belong together. They are not four independent votes cast for the same conclusion. They partition one serial construction and place different requirements upon it. Construction asks what must exist; physics asks what could be observed; computation asks what can be witnessed and at what compiler cost; this volume walks the code and checks how the translations are calibrated. Agreement among the gauges is useful only where the same relation survives the change of register. The second pass tests a declared translation; it does not, by itself, turn that translation into an established physical interpretation.

One final rule of honesty belongs at the door. The repository distinguishes what is *built*, what is *measured*, and what is *interpreted*. Lean definitions, theorems, evaluations, and axiom footprints are built. Deterministic elaborator reports are measured within a pinned compiler and device state. The language of free will, choice, witness, observation, and physical naming is interpretation: the reading this volume proposes, not a theorem the kernel has proved. The owned result is the bracket the device can repeat. Its safest name is the compiler structure constant,  $\alpha_c$ ; comparison with a physical constant is a calibration or interpretation performed by another gauge, never a substitution of the outside target for the device's own report.

That distinction is why the code remains open beside the prose. Where this volume leans on established mathematics or experimental science, it names the people whose work supplies the ground. Where the citations thin, the claim is ours and the reader should slow down. Generosity is part of the calibration: it lets inherited knowledge and new interpretation be told apart.



But no citation and no paragraph replaces the second pass. The reader is handed the instrument because a rerun is the final witness of repeatability; the philosophical and physical naming remains interpretation.

You may choose the question. You may not choose the answer. The world must be allowed to answer, the instrument must make a record, and the record must survive the exchange of roles before its names are earned.

Choose the frame. Produce the witness. Let the world answer. Observe it twice. Then — and only then — use the name.

# Contents

|                                                            |           |
|------------------------------------------------------------|-----------|
| <b>Preface: The Freedom to Answer</b>                      | <b>i</b>  |
| <b>I A Difference</b>                                      | <b>1</b>  |
| <b>1 A Difference Appears</b>                              | <b>2</b>  |
| 1.1 Difference is a predicate, not a particle . . . . .    | 2         |
| 1.2 The symbol and the receiver . . . . .                  | 10        |
| 1.3 From a decision to a record . . . . .                  | 14        |
| <b>2 Difference Becomes Countable</b>                      | <b>19</b> |
| 2.1 Counting is an interface . . . . .                     | 19        |
| 2.2 The update and the alleged endpoint . . . . .          | 24        |
| 2.3 A finite trace is not an infinite object . . . . .     | 28        |
| <b>3 The Residue Refuses To Disappear</b>                  | <b>35</b> |
| 3.1 The branch that cannot be silently discarded . . . . . | 35        |
| 3.2 Iterative error versus structural residue . . . . .    | 40        |
| 3.3 Observation becomes history . . . . .                  | 45        |
| <b>II The Turn</b>                                         | <b>51</b> |
| <b>4 The Tower Turns Around</b>                            | <b>52</b> |

|            |                                                   |            |
|------------|---------------------------------------------------|------------|
| 4.1        | Requirements pull backward . . . . .              | 52         |
| 4.2        | Witnesses run forward . . . . .                   | 58         |
| 4.3        | Two-pass calibration . . . . .                    | 62         |
| <b>5</b>   | <b>The Slip Becomes A Number</b>                  | <b>69</b>  |
| 5.1        | A crossing predicate . . . . .                    | 69         |
| 5.2        | Bisection is iteration by slicing . . . . .       | 75         |
| 5.3        | A report is not a certificate . . . . .           | 79         |
| <b>6</b>   | <b>The Bracketed Number</b>                       | <b>85</b>  |
| 6.1        | Dyadic data structure . . . . .                   | 85         |
| 6.2        | Resolution as an explicit parameter . . . . .     | 89         |
| 6.3        | Stop rules and budgets . . . . .                  | 94         |
| <b>III</b> | <b>The Faces</b>                                  | <b>99</b>  |
| <b>7</b>   | <b>The Number Splits Into Faces</b>               | <b>100</b> |
| 7.1        | One corridor, four faces . . . . .                | 100        |
| 7.2        | The partition point . . . . .                     | 102        |
| 7.3        | Same residue, different reading . . . . .         | 104        |
| <b>8</b>   | <b>The Electron Is Named</b>                      | <b>107</b> |
| 8.1        | Naming makes it measurable . . . . .              | 107        |
| 8.2        | The stable reading . . . . .                      | 109        |
| 8.3        | Named by counting, not fiat . . . . .             | 111        |
| <b>9</b>   | <b>The Corridor Rotates</b>                       | <b>113</b> |
| 9.1        | Rotation turns number into phase . . . . .        | 113        |
| 9.2        | Charge, mass, motion, spin, orientation . . . . . | 115        |
| 9.3        | Relabeling is angular momentum . . . . .          | 116        |

|           |                                                  |            |
|-----------|--------------------------------------------------|------------|
| <b>IV</b> | <b>The Closing</b>                               | <b>119</b> |
| <b>10</b> | <b>The Native Apparatus Presents the Residue</b> | <b>120</b> |
| 10.1      | Forcing the residue to show . . . . .            | 120        |
| 10.2      | The physical bench . . . . .                     | 122        |
| 10.3      | The computational bench . . . . .                | 124        |
| <b>11</b> | <b>The Meter</b>                                 | <b>127</b> |
| 11.1      | Elaboration cost is a reading . . . . .          | 127        |
| 11.2      | Self-reference . . . . .                         | 129        |
| 11.3      | The pulse in its own currency . . . . .          | 131        |
| <b>12</b> | <b>The Field Closes</b>                          | <b>133</b> |
| 12.1      | The second variation . . . . .                   | 133        |
| 12.2      | The magnetic side present . . . . .              | 135        |
| 12.3      | Asking, not delivering . . . . .                 | 137        |
| <b>V</b>  | <b>The Bound</b>                                 | <b>139</b> |
| <b>13</b> | <b>Alpha As A Bracket</b>                        | <b>140</b> |
| 13.1      | The coupling is asked . . . . .                  | 140        |
| 13.2      | The cut . . . . .                                | 142        |
| 13.3      | The bracket, not the real . . . . .              | 145        |
| <b>14</b> | <b>The Machine Measures Its Own Measurement</b>  | <b>148</b> |
| 14.1      | Self-application . . . . .                       | 148        |
| 14.2      | The witness bears the cost . . . . .             | 150        |
| 14.3      | Two descriptions, one object . . . . .           | 151        |
| <b>15</b> | <b>What The Machine Can And Cannot Claim</b>     | <b>154</b> |
| 15.1      | The structure it derives . . . . .               | 154        |
| 15.2      | The magnitude it cannot see . . . . .            | 156        |

*CONTENTS*

xii

|                        |     |
|------------------------|-----|
| 15.3 The jar . . . . . | 158 |
|------------------------|-----|

# Part I

## A Difference

# Chapter 1

## A Difference Appears

### 1.1 Difference is a predicate, not a particle

Scientific computation begins by fixing a grammar before assigning meaning to its names. Let  $M$  be a finite alphabet of marks:

$$\begin{aligned} \textit{mark} &::= m_0 \mid \cdots \mid m_{q-1}, \\ \textit{question} &::= \text{diff}(s, \sigma), \\ \textit{answer} &::= \text{yes}(w) \mid \text{no}(r). \end{aligned}$$

The payload  $w$  certifies that two marks differ;  $r$  certifies that they do not. The grammar does not make a mark a material object, a question a detector, or an answer an observation. Those names require separate provenance.

For a reference  $\sigma \in M$ , the governing question is

$$D_\sigma(s) := (s \neq \sigma).$$

A resolver  $\Delta$  must terminate with an evidenced judgment,

$$\frac{w : s \neq \sigma}{\Delta(s, \sigma) \Downarrow \text{yes}(w)} \quad \frac{r : \neg(s \neq \sigma)}{\Delta(s, \sigma) \Downarrow \text{no}(r)}.$$

A later routine may erase the evidence:

$$E(\text{yes}(w)) = 1, \quad E(\text{no}(r)) = 0.$$

This external Boolean encoding is lossy. Its bit records the selected branch but not the certificate that licensed it.

Reading has a direction. Let  $X$  be a finite collection of records and let  $e : X \rightarrow M$  represent each record by a mark. Construction sends records forward through  $e$ . A question  $q : M \rightarrow \{0, 1\}$ , or its evidence-bearing refinement, is read contravariantly by pullback:

$$e^*q := q \circ e : X \rightarrow \{0, 1\}.$$

For a second representation  $g : M \rightarrow N$ ,

$$(g \circ e)^* = e^* \circ g^*.$$

The order reverses: first pull the question on  $N$  through  $g$ , then through  $e$ . This is an external representation schema, not a claim that a particular map  $e$  has already been constructed. It separates the covariant forward construction of a representation from the contravariant act of reading it.

An audit record is

$$\mathcal{R} = (X, M, e, \sigma, D, \Delta, E, (A, R), J, F, P),$$

where  $J$  states invariants,  $F$  states failures, and  $P$  tags computation, physics, or interpretation. For  $q = E \circ \Delta(-, \sigma)$ ,

$$A = \{x \in X \mid (e^*q)(x) = 1\}, \quad R = \{x \in X \mid (e^*q)(x) = 0\}.$$

The words accepted and rejected are conventional branch labels, not value judgments.



The finite algorithm is explicit. Given  $L = [x_0, \dots, x_{n-1}]$  and an independent decision budget  $N$ , initialize two empty output lists and  $i \leftarrow 0$ . While  $i < n$  and  $i < N$ , compute  $m_i = e(x_i)$ , resolve  $\Delta(m_i, \sigma)$ , append the record and its evidence to the corresponding slice, and update  $i \leftarrow i + 1$ . Stop with reason COMPLETE when  $i = n$ , BUDGET when  $i = N < n$ , or RESOLVER-DOMAIN when the declared resolver does not apply to  $(m_i, \sigma)$ . If  $k$  records are visited, its cost is

$$C(L, N) = \sum_{i=0}^{k-1} (C_e(x_i) + C_\Delta(e(x_i), \sigma) + C_{\text{append}}(i)), \quad k \leq \min(n, N).$$

Its certificate contains the visited prefix  $[x_0, \dots, x_{k-1}]$ , both evidence-bearing slices, the stop reason, and every declared Boolean erasure. Coverage and exclusivity apply only to that prefix:  $|A| + |R| = k$ . A complete run has  $k = n$ ; a budget or resolver-domain stop has  $k = i$  and says nothing about the unvisited suffix. Stable append preserves relative order *within* each branch, but separating the branches loses their original interleaving. Classification error is zero relative to the declared question when every recorded certificate checks. That internal guarantee is not an empirical error rate.

**Worked example 1: printed marks (computation).** Take  $X = M = \{\mathbf{x}, \mathbf{o}\}$ ,  $e$  the identity,  $\sigma = \mathbf{x}$ , and  $L = [\mathbf{x}, \mathbf{o}, \mathbf{o}, \mathbf{x}]$ . Four decisions give

$$A = [\mathbf{o}, \mathbf{o}], \quad R = [\mathbf{x}, \mathbf{x}], \quad |A| + |R| = 4.$$

With  $N \geq 4$ , the stopping certificate is COMPLETE with  $i = 4 = |L|$ . With  $N = 2$ , it instead certifies only the prefix  $[\mathbf{x}, \mathbf{o}]$ , one entry in each slice, and stop reason BUDGET; it asserts nothing about the two remaining marks. Exchanging the two printed shapes everywhere preserves slice sizes while swapping their branch names. No metric follows: matching shapes does not define a distance between them.

**Worked example 2: an integer comparator (computation).** Let  $M = \{0, \dots, 9\}$ ,  $\sigma = 7$ , and process  $L = [7, 8, 2, 7, 9]$ . The result is

$$A = [8, 2, 9], \quad R = [7, 7].$$

The separated slices lose the original interleaving between branches and discard magnitude, while stable append preserves relative order within each branch. Values 8 and 2 take the same branch although a separately supplied arithmetic map gives  $|8-7| = 1$  and  $|2-7| = 5$ . The name difference does not supply subtraction. To obtain it, the representation must add  $d : M \times M \rightarrow \mathbb{N}$  and publish its cost and assumptions.

**Worked example 3: a finite packet alphabet (communication model).** Declare a discrete source on  $M = \{\mathbf{A}, \mathbf{B}, \mathbf{C}\}$  with

$$p(\mathbf{A}) = \frac{1}{2}, \quad p(\mathbf{B}) = p(\mathbf{C}) = \frac{1}{4}.$$

For  $p(x) > 0$ , Shannon self-information is

$$I(x) = -\log_2 p(x),$$

so the three marks carry one, two, and two bits respectively [44]. With  $\sigma = \mathbf{A}$ , both  $\mathbf{B}$  and  $\mathbf{C}$  enter the same slice. The erased decision therefore cannot recover the source mark. Shannon's channel language belongs here to a separately declared computational model, not to a physical signal.

**Worked example 4: a sampled switch (physics and interpretation).**

Suppose an experiment samples a switch at  $t_k = k\Delta t$  and records  $v_k \in \{0, 1\}$ . Physics owns the switch, sampling interval, noise, and acquisition apparatus. Post-processing may represent a record by  $e(x_k) = v_k$  and pull back the question  $D_0$ . If the apparatus has false-positive probability  $\varepsilon_+$ , false-negative

probability  $\varepsilon_-$ , and on-state prior  $\pi$ , its empirical error is

$$P_e = (1 - \pi)\varepsilon_+ + \pi\varepsilon_-.$$

This belongs to the physical route from switch state to recorded mark, not to the internal comparator. The routes meet only through an explicit comparison map and observations; shared names prove no identity.

**Worked example 5: convergence demonstrated (the instrument calibration).** Here *convergence* means agreement of two completed finite numerical routes at a declared resolution, not an asymptotic claim. Both routes begin from the same calibration data

$$C = 18, \quad \text{target} = 5, \quad d_* = \sqrt{C/\text{target}} = \sqrt{18/5}.$$

The direct route  $1 \rightarrow 3$  is a one-pass forward evaluation at  $d_*$ . In the declared scaled-integer representation it reports

$$d = 137011290548979455469 \quad \text{at scale } 10^{18}.$$

The stepped route  $1 \rightarrow 2 \rightarrow 3$  carries the residue through intermediate quasi-Newton updates. Its finite trace is

| count | $r_k$              | reported inverse value |
|-------|--------------------|------------------------|
| 0     | 2/1                | 129.6                  |
| 1     | 17/9               | 137.7                  |
| 2     | $\approx 1.897309$ | $\approx 137.016$      |
| 3     | $\approx 1.897366$ | $s/10^{18}$            |

and it stops at count three with

$$s = 137011290753751157155 \quad \text{at scale } 10^{18}.$$

These are distinct instrument algorithms. The stepped route has mediant/continued-fraction, Newton/BFGS, and residue-carrying secant realizations; these are numerical constructions whose update and stopping rules must be declared separately. Both endpoint computations are covariant forward constructions over the shared calibration data. The contravariant action enters when the comparison question is pulled back to those data; it is not a second name for either endpoint algorithm.

The finite comparator certificate is not integer equality. It records

$$|d-s| = 204771701686, \quad \frac{|d-s|}{10^{18}} = 0.000000204771701686 \approx 2.04771701686 \times 10^{-7}.$$

Thus the reported decimal values agree to nine digits. To state the reading precisely, let  $X$  be the shared calibration-input space,  $D$  the scaled-integer report space, and

$$F = (F_{\text{direct}}, F_{\text{stepped}}) : X \rightarrow D \times D.$$

At count-three resolution define the comparison question

$$q_3(a, b) := \left[ \frac{|a-b|}{10^{18}} < 8.1 \right], \quad F^*q_3 = q_3 \circ F.$$

The pullback  $F^*q_3$  is the contravariant reading of the two covariantly constructed reports. Write  $a \equiv_3 b$  when this declared-resolution question answers yes. For the shared input above,  $d \equiv_3 s$ . The house slogan  $1 \rightarrow 3 = 1 \rightarrow 2 \rightarrow 3$  abbreviates this relation; it is not integer equality and establishes path-independence only for these two routes on this calibration input.

The two-pass calibration therefore licenses a path-independent shared numerical result at the declared resolution, which is what earns the name  $\alpha_c$ . It does not identify the methods, meanings, traces, costs, or physics/computation provenance of the routes. Nor does it promote either scaled integer to the owned output. The honest result remains the ordered walls written in the

repository's house notation  $[129.6, 137.7]$ , not a point. The word *open* is the repository's metaphor for an unresolved jar, not topological interval notation; the displayed bracket records the two owned ordered walls. The values near 137.011290 are finer reports obtained past the floor.

The roles of the two passes are now exact. Pass one directly certifies  $C = 18$  and target = 5, then infers the crossing  $\sqrt{C/\text{target}}$ ; the crossing is not claimed as an independent measurement. Pass two traverses the stepped trace and compares its endpoint with the direct endpoint. Any comparison with a physical quantity requires a separate provenance chain and empirical observations.

**Conditional Richardson audit.** The direct and stepped reports are *not* a valid Richardson pair: they come from different methods, and no common discretization family  $A(h)$ , refinement ratio  $r$ , or error order  $p$  has been established. Only under the counterfactual assumptions  $D = A(h)$ ,  $S = A(h/2)$ , and leading error proportional to  $h^p$  would the extrapolate be

$$A_R = S + \frac{S - D}{2^p - 1}.$$

Applying that formula conditionally gives the sensitivity ledger

| $p$ | $A_R$                         | correction added to $S$          |
|-----|-------------------------------|----------------------------------|
| 1   | 137.011290958522858841        | $2.04771701686 \times 10^{-7}$   |
| 2   | 137.011290822008391050333 ... | $6.8257233895333 \times 10^{-8}$ |
| 4   | 137.011290767402603934066 ... | $1.3651446779067 \times 10^{-8}$ |

Treating the stepped report as the finer value is itself arbitrary and counterfactual here. Reversing the orientation reverses the correction, and no common refinement family selects either orientation. These are sensitivity values, not measurements, uncertainty bounds, or owned results. Their variation demonstrates dependence on an unestablished order assumption. Li-

censing Richardson extrapolation would require one declared family  $A(h)$ , a declared ratio  $r$ , at least three refinement levels, an observed-order estimate

$$p_{\text{obs}} = \frac{\log(|A_h - A_{h/r}|/|A_{h/r} - A_{h/r^2}|)}{\log r},$$

and stability under further refinement. Until those conditions are met, the ordered-wall jar [129.6, 137.7] remains the owned result.

Two finite consequences and losses are immediate. If  $|M| > 2$ , the erased map  $E \circ \Delta(-, \sigma) : M \rightarrow \{0, 1\}$  cannot be injective, so a binary slice cannot recover every mark. Also,  $\sigma$  does not determine the question: equality, order, membership, and threshold rules can share a reference while producing different slices. A report must publish  $(\sigma, D)$ , the representation  $e$ , update rule, stopping rule, cost, and certificate.

Symmetry supplies only

$$D_\sigma(s) \iff D_s(\sigma).$$

It preserves a relation under exchange of arguments; it proves no equality of procedures, traces, costs, or physical protocols. A two-pass calibration may test those stronger claims, but it must record both passes and state its comparison map.

The result is finite and reusable. Marks support explicit questions; terminating decisions carry evidence; evidence may be erased to bits; and pulled-back questions can slice finite records. Physics enters only when an experiment supplies an apparatus, units, uncertainty, and observations. A resonant name belongs to one provenance chain at a time. If computation and physics later report the same number or process, each route must be shown separately before comparison.

## 1.2 The symbol and the receiver

A receiver is not a second object placed beside a symbol. It is a declared action that projects a finite alphabet into the distinctions a reading can retain. Fix a mark domain  $M$ , a reference  $\sigma \in M$ , and an answer alphabet  $B = \{0, 1\}$ . The read grammar is

$$\begin{aligned} \textit{projection} &::= \pi_\sigma : M \rightarrow B, \\ \textit{accepted} &::= \text{yes}(m, w), \\ \textit{rejected} &::= \text{no}(m, r), \end{aligned}$$

where  $w$  certifies  $m \neq \sigma$  and  $r$  certifies its negation. The projection acts by

$$\pi_\sigma(m) = \begin{cases} 1, & m \neq \sigma, \\ 0, & m = \sigma. \end{cases}$$

This is one declared receiver grammar. Another grammar may replace its question and resolver; all slices below concern the displayed inequality rule. Its zero fiber is  $N_\sigma = \pi_\sigma^{-1}(\{0\})$ . The accepted-bit view suppresses distinctions among marks in that fiber by mapping them all to zero, while the rejected evidence ledger retains the visited occurrences. More generally, information may be lost in either many-to-one fiber of  $\pi_\sigma$ . This is a finite projection, not a claim about a linear kernel.

Direction remains part of the grammar. A forward representation  $e : X \rightarrow M$  carries records into marks, and  $\pi_\sigma : M \rightarrow B$  carries marks forward into bits. Let  $q : B \rightarrow \mathbf{Answer}$  be a question asked of a projected bit. Reading pulls  $q$  back:

$$\pi_\sigma^* q := q \circ \pi_\sigma, \quad e^*(\pi_\sigma^* q) = q \circ \pi_\sigma \circ e.$$

Equivalently,

$$(\pi_\sigma \circ e)^* q = e^*(\pi_\sigma^* q).$$

The order of reading reverses the order of forward maps. The receiver there-

fore consists of a finite forward projection followed by contravariant reading of the projected value; the projection itself is not contravariant.

The numerical method is a budgeted stable partition. Given  $L = [x_0, \dots, x_{n-1}]$ , a budget  $N$ , representation  $e$ , reference  $\sigma$ , and resolver  $\Delta$ , initialize  $A \leftarrow []$ ,  $R \leftarrow []$ ,  $i \leftarrow 0$ . While  $i < n$  and  $i < N$ :

1.  $m_i \leftarrow e(x_i)$ ,
2.  $a_i \leftarrow \Delta(m_i, \sigma)$ ,
3. append  $(x_i, m_i, a_i)$  to  $A$  for yes, otherwise to  $R$ ,
4.  $i \leftarrow i + 1$ .

The run stops with COMPLETE at  $i = n$ , BUDGET at  $i = N < n$ , or RESOLVER-DOMAIN when the declared decision does not cover the encountered pair. Its certificate contains the visited prefix, accepted and rejected evidence, erased projection bits, and stop reason. On the visited prefix,

$$|A| + |R| = i, \quad A \cap R = \emptyset.$$

Stable append preserves relative order within each slice but loses the interleaving between slices. A partial certificate says nothing about the unvisited suffix.

The audit schema is

$$\mathcal{Q} = (X, M, e, \sigma, \pi_\sigma, \Delta, N, A, R, J, F, P).$$

Here  $J$  records invariants,  $F$  records failure and information loss, and  $P$  records provenance. The schema is external scientific bookkeeping. The computational route owns finite marks, decisions, projection, traversal, and slices. Physics owns apparatus variables, units, noise, and observations. Interpretation owns any proposed bridge.



**Worked example 1: repeated finite marks (computation).** Let  $M = \{\mathbf{a}, \mathbf{b}, \mathbf{c}\}$ ,  $\sigma = \mathbf{a}$ , and

$$L = [\mathbf{a}, \mathbf{c}, \mathbf{b}, \mathbf{a}, \mathbf{c}].$$

With  $N \geq 5$ , the complete run returns

$$A = [\mathbf{c}, \mathbf{b}, \mathbf{c}], \quad R = [\mathbf{a}, \mathbf{a}].$$

The projected trace is  $[0, 1, 1, 0, 1]$ . It cannot distinguish  $\mathbf{b}$  from  $\mathbf{c}$ , although the evidence-bearing accepted slice can retain their marks. If only the projected trace is stored, the original input order has several compatible preimages. In this example the loss occurs in the accepted fiber  $\{\mathbf{b}, \mathbf{c}\}$ , not in the zero fiber  $\{\mathbf{a}\}$ ; either fiber can lose distinctions when it has more than one member.

**Worked example 2: budgeted packet reading (computation).** Take packet marks  $M = \{\text{idle}, \text{data}, \text{check}\}$ , reference  $\sigma = \text{idle}$ , six records, and  $N = 4$ . Suppose their first four marks are

$$[\text{idle}, \text{data}, \text{check}, \text{data}].$$

The run stops for budget with one rejected and three accepted records. Its coverage equation is  $1 + 3 = 4$ , not six. Neither remaining packet may be classified from this certificate. If each decision costs  $c_i$ , the certified work is  $\sum_{i=0}^3 c_i$ ; a claim about six decisions would exceed the ledger.

**Worked example 3: Shannon loss under projection (communication model).** Let a discrete source satisfy

$$p(\mathbf{a}) = \frac{1}{2}, \quad p(\mathbf{b}) = p(\mathbf{c}) = \frac{1}{4}.$$

For  $p(m) > 0$ , self-information is  $I(m) = -\log_2 p(m)$  [44]. The source entropy is

$$H(M) = -\sum_m p(m) \log_2 p(m) = \frac{3}{2} \text{ bits.}$$

Projection against  $\sigma = \mathbf{a}$  yields an equiprobable bit, hence  $H(B) = 1$  bit. The lost  $\frac{1}{2}$  bit is the average distinction between  $\mathbf{b}$  and  $\mathbf{c}$  erased by the projection. This calculation requires the displayed discrete distribution; the receiver grammar alone supplies no probabilities.

**Worked example 4: threshold readout (physics, then interpretation).** Suppose a finite experiment records integer converter readings  $z_k \in \{0, \dots, 1023\}$ . A declared threshold  $\tau$  gives

$$\pi_\tau(z) = \mathbf{1}\{z > \tau\}.$$

Physics owns the sensor, converter, units relating the converter reading to the measured quantity, calibration of  $\tau$ , and observed false-positive and false-negative frequencies. Computation owns the finite comparison and resulting slices. Reading the projection as a detector action is licensed only by the experimental protocol. Equal output bits do not identify two sensor states, and a perfect internal comparison does not imply zero empirical error.

The receiver therefore earns a precise but limited name. It is a finite projection followed by contravariant reading of the projected value, not a physical detector. Its accepted and rejected slices certify a finite partition of the visited prefix. Its many-to-one fibers state exactly what the erased bit cannot recover. Any later claim that a symbol has physical meaning must supply a separate apparatus route and comparison map; resemblance of names is not a calibration.

### 1.3 From a decision to a record

A durable record begins with a receipt, not with a familiar arithmetic interpretation. The neutral grammar is

$$\begin{aligned}
 \textit{numeral-record} &::= \text{anchor}(\textit{receipt}) \mid \text{numeral-step}(\textit{receipt}, \textit{numeral-record}), \\
 \textit{tally} &::= \text{halt}(\textit{receipt}) \mid \text{tally-step}(\textit{receipt}, \textit{numeral-record}, \textit{tally}), \\
 \textit{piece} &::= \text{origin}(\textit{receipt}) \mid \text{piece-step}(\textit{receipt}, \textit{tally}, \textit{piece}), \\
 \textit{record} &::= \text{end}(\textit{receipt}) \mid \text{record-step}(\textit{receipt}, \textit{piece}, \textit{record}).
 \end{aligned}$$

The numeral record is the earlier finite recursive payload, given here only by its anchor and extension forms. Its arithmetic meaning is reserved rather than supplied by the name. The grammar posits finite forms and nothing else. A receipt accompanies every form, and each recursive argument stores the remainder. No magnitude, chronology, or physical event is supplied by these productions.

Forward construction prepends one step:

$$\begin{aligned}
 T' &= \text{tally-step}(\rho, \nu, T), \\
 P' &= \text{piece-step}(\rho, T, P), \\
 S' &= \text{record-step}(\rho, P, S).
 \end{aligned}$$

Repeated prepending places the newest addition at the head. Tail traversal proceeds by matching the head, recording its receipt and payload, and continuing with the recursive remainder. Given budget  $N$ , initialize the visit count  $j \leftarrow 1$ . Continue while a step remains and  $j \leq N$ ; stop with END, BUDGET, or RESOLVER-DOMAIN. The certificate contains the visited prefix in traversal orientation, the final unvisited tail when present, all encountered receipts, and the stop reason.

Traversal indices are analysis-ledger entries, not members of the grammar. The ledger assigns  $j = 1, 2, \dots$  as a run visits heads. A later run may assign the same indices only when it declares the same starting form and traversal

orientation. Prepend construction and tail traversal point in opposite list orientations: construction adds newest first; complete traversal reads newest to oldest. Arrival orientation is recovered only by an explicit reversal.

Three finite slices are available after traversal. A shape slice separates terminal forms from step forms. An index slice selects ledger positions satisfying a declared question, such as  $j$  odd. Given a declared comparison receipt  $\rho_*$ , a receipt slice selects one value of the derived relation  $\text{parity}(\rho, \rho_*)$ . Each slice must retain evidence, visited occurrences, and orientation. Stable append preserves relative traversal order within each output but loses interleaving between outputs. For  $v$  visited occurrences,

$$|A| + |R| = v, \quad A \cap R = \emptyset.$$

A budget stop certifies nothing about the unvisited tail.

**Worked example 1: prepend and budgeted traversal (computation).** Let  $E = \text{end}(\rho_e)$ , and prepend three record steps carrying pieces  $P_1, P_2, P_3$  in that arrival order:

$$S_1 = \text{record-step}(\rho_1, P_1, E), \quad S_2 = \text{record-step}(\rho_2, P_2, S_1), \quad S_3 = \text{record-step}(\rho_3, P_3, S_2).$$

Complete tail traversal reports  $P_3, P_2, P_1$ , followed by the terminal receipt. With  $N = 2$ , it certifies only ledger entries  $(1, \rho_3, P_3)$  and  $(2, \rho_2, P_2)$ , then returns the remaining tail  $S_1$  with BUDGET. It may not report anything about  $P_1$  without another visit.

A head projection deliberately ignores recursive tails. For step forms define

$$\pi_{\text{head}}(\text{record-step}(\rho, P, S)) = (\rho, P).$$

For a terminal form it returns its terminal receipt and a terminal tag. This is a many-to-one forward map. Its fibers contain every form with the same head summary, regardless of tail.

**Worked example 2: same head, different tail (computation).** Let

$$S_A = \text{record-step}(\rho, P, S_1), \quad S_B = \text{record-step}(\rho, P, S_2), \quad S_1 \neq S_2.$$

Then

$$\pi_{\text{head}}(S_A) = \pi_{\text{head}}(S_B) = (\rho, P).$$

The head read cannot reconstruct which tail was present. The pair lies in one fiber of  $\pi_{\text{head}}$ , and the missing distinction is exactly the tail information suppressed by projection. Retaining the original forms in an evidence ledger preserves them; retaining only the projected head does not.

The relation between two receipts also directs a finite comparison. Declare a current receipt  $\rho$ , a comparison receipt  $\rho_*$ , a derived relation  $\text{parity}(\rho, \rho_*) \in \{\text{agree}, \text{disagree}\}$ , current payload  $a$ , comparison payload  $b$ , and a binary question  $C$ . The action is

$$\text{compare}_{\rho, \rho_*}(a, b) = \begin{cases} C(a, b), & \text{parity}(\rho, \rho_*) = \text{agree}, \\ C(b, a), & \text{parity}(\rho, \rho_*) = \text{disagree}. \end{cases}$$

Derived agreement sends the current payload into the forward comparison; derived disagreement reverses the arguments. This rule establishes no order law. If  $C$  is asymmetric, reversal may change the answer; if  $C$  is symmetric, it may not. The relation between receipts chooses argument orientation but does not prove anything about the comparison relation.

**Worked example 3: receipt-directed comparison (computation).**

Let two numeral-record payloads be denoted  $\nu_L, \nu_R$ , and declare  $C(a, b)$  to answer yes exactly for the ordered pair  $(\nu_L, \nu_R)$ . Then

$$\text{compare}_{\rho_1, \rho_*}(\nu_L, \nu_R) = \text{yes} \quad \text{when} \quad \text{parity}(\rho_1, \rho_*) = \text{agree}, \quad \text{compare}_{\rho_2, \rho_*}(\nu_L, \nu_R) = \text{no} \quad \text{when} \quad \text{pa}$$

The outputs differ solely because the second receipt relation reverses the arguments. Neither output assigns an arithmetic or physical meaning to  $\nu_L, \nu_R$ . Those resonances remain reserved.

Contravariance belongs to questions pulled through forward projections. If  $q : H \rightarrow \mathbf{Answer}$  asks about a head summary and  $\pi_{\text{head}} : S \rightarrow H$ , then

$$\pi_{\text{head}}^* q = q \circ \pi_{\text{head}}.$$

For an index analysis, let  $\pi_{\text{idx}}$  project each visited ledger entry to its assigned index. An index question  $q_I$  is read as  $\pi_{\text{idx}}^* q_I$ . The projections move summaries forward; reading pulls questions backward. The receipt-directed argument reversal above is a separate finite action and is not identified with this pullback law.

**Worked example 4: shape, index, and receipt slices (computation).**

Relative to declared comparison receipt  $\rho_*$ , suppose a complete traversal yields four ledger rows whose middle column is the derived parity

| index | derived receipt relation                            | payload |
|-------|-----------------------------------------------------|---------|
| 1     | $\text{parity}(\rho_4, \rho_*) = \mathbf{agree}$    | $P_4$   |
| 2     | $\text{parity}(\rho_3, \rho_*) = \mathbf{disagree}$ | $P_3$   |
| 3     | $\text{parity}(\rho_2, \rho_*) = \mathbf{agree}$    | $P_2$   |
| 4     | $\text{parity}(\rho_1, \rho_*) = \mathbf{agree}$    | $P_1$   |

and then a terminal row. The odd-index slice selects rows one and three. The derived-agree receipt slice selects rows one, three, and four. The shape slice separates the four step rows from the terminal row. These are three questions over one traversal, not three reconstructions of the original arrival orientation.

A computation-side resonance with Cohen may now be stated carefully.

From a completed traversal  $S$ , define the finite partial assignment

$$p_S(k) = \text{payload at traversal index } k$$

for visited indices  $k$ . Extending the budget can extend the assignment when earlier rows remain unchanged. This resembles Cohen’s finite-condition discipline: compatible finite commitments may be extended [11]. The resemblance concerns finite partial assignments only. Because traversal reads prepend-built forms newest first,  $p_S(k)$  is indexed in traversal orientation, not arrival orientation. No generic extension, forcing relation, independence result, or claim about every possible continuation follows.

The action ledger closes the section. The receipt grammar and prepend rules are *posited*. A finite construction trace, visited-prefix certificate, projection equality, or checked slice is *demonstrated*. Head and index questions are *read* contravariantly through declared projections. Arithmetic meanings for numeral records and physical meanings for receipt relations, sides, or arrival events remain *reserved*. A head projection does not reconstruct a tail; a finite slice does not certify an unread remainder; argument reversal does not select which side is which; and a compatible partial assignment does not produce a generic extension.

## Chapter 2

# Difference Becomes Countable

### 2.1 Counting is an interface

Counting begins with four finite capabilities, not with a completed enumeration. A neutral grammar is

$$\begin{aligned} \textit{counting-capability} &::= \text{counting}(\textit{carrier}, \textit{tally}, U), \\ \textit{admission-capability} &::= \text{admission}(\textit{counting-capability}, A), \\ \textit{indexing-capability} &::= \text{indexing}(\textit{counting-capability}, \textit{origin}, V), \\ \textit{bounding-capability} &::= \text{bounding}(\textit{indexing-capability}, B). \end{aligned}$$

Here  $U$  updates a tally-shaped record,  $A$  asks an admission question,  $V$  updates an indexed record from a declared origin, and  $B$  asks a bound question. The capabilities provide places for finite data and questions. They do not guarantee that repeated updates are distinct, that every candidate appears, or that an index assignment is one-to-one or onto.

Use the neutral forms

$$\text{halt}(\rho), \quad \text{step}(\rho, \nu, T),$$

with receipt  $\rho$ , numeral record  $\nu$ , and recursive tail  $T$ . Arithmetic meaning



for  $\nu$  remains reserved. The default counting update has a head-update/tail-reset shape:

$$\begin{aligned} U(\text{halt}(\rho)) &= \text{halt}(\rho_*), \\ U(\text{step}(\rho, \nu, T)) &= \text{step}(\rho_*, \nu_*, T_*), \end{aligned}$$

where  $\rho_*, \nu_*, T_*$  are the current receipt, carrier payload, and stored tally supplied by the capability. The input tail does not survive the second branch. The default indexing update likewise uses a fixed origin  $O$ :

$$\begin{aligned} V(\text{halt}(\rho)) &= \text{step}(\rho, c, O), \\ V(\text{step}(\rho, \nu, T)) &= \text{step}(\rho, U(\nu), O). \end{aligned}$$

It retains the input head receipt, updates the head payload, and resets the tail to  $O$ .

**Worked example 1: same head, different tail (computation).** Let

$$R_1 = \text{step}(\rho, \nu, T_1), \quad R_2 = \text{step}(\rho, \nu, T_2), \quad T_1 \neq T_2.$$

The default recurrence gives

$$V(R_1) = V(R_2) = \text{step}(\rho, U(\nu), O).$$

Thus the update is many-to-one: records with the same head and different tails share an output fiber. Iteration cannot reconstruct the discarded tail, so this recurrence alone cannot certify enumeration of its input records.

The default admission question is also only a question. On neutral numeral forms it has three branches:

$$A(s, t) = \begin{cases} \text{yes}, & s \text{ is an anchor,} \\ \text{no}, & s \text{ is a step and } t \text{ is an anchor,} \\ R(\rho_t) \vee Q(\nu_s, \nu_t), & s, t \text{ are steps.} \end{cases}$$

Here  $R$  is the declared question carried by the threshold receipt; an affirmative receipt answer admits the pair without consulting  $Q$ . The default bound question has a related orientation rule:

$$B(s, t) = \begin{cases} \text{yes}, & s \text{ is an anchor,} \\ \text{no}, & s \text{ is a step and } t \text{ is an anchor,} \\ [\text{parity}(\rho_s, \rho_t) = \text{agree} \wedge Q(i_s, i_t)] \\ \vee [\text{parity}(\rho_s, \rho_t) = \text{disagree} \wedge Q(i_t, i_s)], & s, t \text{ are steps.} \end{cases}$$

Agreement preserves the comparison arguments; disagreement reverses them. Neither capability supplies a resolver for  $A$ ,  $B$ , or  $Q$ , and none proves reflexivity, transitivity, antisymmetry, totality, or any other order law. The questions and defaults may be replaced in another declared capability; subsequent slices then concern the replacements.

The actual iteration is the forward orbit

$$r_0 = O, \quad r_{k+1} = V(r_k).$$

For an actual finite run with transition budget  $N$ , record  $r_0 = O$ , then apply  $V$  for  $k = 0, \dots, N - 1$ . Stop with BUDGET after the  $N$ -th transition. The certificate contains the declared origin, budget  $N$ , all  $N$  transition pairs  $(r_k, r_{k+1})$ , the receipts retained inside the recorded values, the prefix  $[r_0, \dots, r_N]$  of  $N + 1$  values, every declared erasure, and the stop reason. It certifies only those  $N$  transitions and says nothing about a further update.

This is a covariant construction. Given a question  $q$  on orbit values, reading is contravariant pullback:

$$V^*q = q \circ V.$$

For composable forward maps  $W$  and  $V$ ,

$$(V \circ W)^*q = W^*(V^*q).$$

Receipt-directed reversal inside  $B$  is not categorical variance; it is merely a branch that exchanges two arguments.

**Worked example 2: an orbit need not enumerate (computation).**

Let a permitted update satisfy  $V(a) = b$  and  $V(b) = a$ . Starting at  $a$ , the orbit is

$$[r_0, r_1, r_2, r_3, r_4, r_5, r_6] = [a, b, a, b, a, b, a] \quad (N = 6).$$

A six-transition run has seven recorded values but only two distinct values. It demonstrates iteration and a two-cycle, not seven distinct indexed objects. Therefore an update interface, even with a budget and origin, is not an enumeration certificate.

A separate finite indexing experiment can be posited, but it must be named as external. Given  $L = [x_1, \dots, x_n]$ , capacity  $K$ , budget  $N$ , and an explicitly declared resolver  $\Delta_A$  for the chosen admission question, visit while  $j \leq n, K, N$ . Resolve  $\Delta_A(x_j)$ , append admitted rows to a ledger and rejected rows to a complement slice, then increment  $j$ . Stop with COMPLETE, CAPACITY, BUDGET, or RESOLVER-DOMAIN. Its certificate contains the visited prefix, both slices, assigned and unused indices, unvisited suffix, and stop reason. For  $v$  visits,

$$|\text{admitted}| + |\text{rejected}| = v.$$

This generic  $L/K/N$  traversal, its assignment  $\iota$ , and its resolver are *posited* analysis machinery; they are not contained in the four capability names.

A genuine finite enumeration claim needs an encoder and decoder. For

finite sets  $X$  and  $I$ , require both round trips:

$$\text{dec}(\text{enc}(x)) = x \quad (x \in X), \quad \text{enc}(\text{dec}(i)) = i \quad (i \in I).$$

**Worked example 3: good and bad round trips (computation).** For  $X = \{\mathbf{a}, \mathbf{b}, \mathbf{c}\}$  and  $I = \{1, 2, 3\}$ , a good pair is

| $x$             | $\mathbf{a}$ | $\mathbf{b}$ | $\mathbf{c}$ | $i$             | 1            | 2            | 3              |
|-----------------|--------------|--------------|--------------|-----------------|--------------|--------------|----------------|
| $\text{enc}(x)$ | 1            | 2            | 3            | $\text{dec}(i)$ | $\mathbf{a}$ | $\mathbf{b}$ | $\mathbf{c}$ . |

Both round trips return their inputs. A bad assignment is

| $x$                          | $\mathbf{a}$ | $\mathbf{b}$ | $\mathbf{c}$ |
|------------------------------|--------------|--------------|--------------|
| $\text{enc}_{\text{bad}}(x)$ | 1            | 1            | 2            |

because the collision at 1 forces any decoder to choose at most one of  $\mathbf{a}, \mathbf{b}$ . Index 3 is also unused. No decoder can satisfy both round trips, regardless of the interface name.

Shannon's codebook is another map,  $C : I \rightarrow \mathcal{W}$ , not the index assignment. With a declared discrete distribution  $p(i) > 0$ ,

$$I(i) = -\log_2 p(i), \quad \bar{\ell} = \sum_i p(i) \ell(C(i)).$$

These require probabilities and word lengths [44]. Unambiguous decoding requires a declared uniquely decodable or prefix condition.

**Worked example 4: an index is not its word (communication model).**

Let  $p(1) = 1/2$ ,  $p(2) = p(3) = 1/4$ , with  $C(1) = 0$ ,  $C(2) = 10$ , and  $C(3) = 11$ . The words are prefix-free, and direct calculation gives entropy  $3/2$  bits and expected length  $3/2$  bits. This equality is demonstrated for this distribution and codebook only. It identifies neither an index with its word nor a capability orbit with an encoded message.

**Worked example 5: counted threshold crossings (physics, then interpretation).** Suppose an apparatus yields a finite acquisition record of threshold crossings. Physics owns the sensor, threshold, dead time, sampling interval, missed-event rate, and false-event rate. The external traversal may index accepted records. Its certificate counts those accepted records, not unobserved physical events. Interpreting the ledger as an event count remains reserved until the experimental error model is supplied.

The action ledger now separates the claims. The four capabilities, their defaults, and the external  $L/K/N$  experiment are *posited*. A finite orbit trace, fiber equality, cycle, slice, or pair of checked round trips is *demonstrated*. Questions on orbit values or indices are *read* contravariantly through forward maps. Full enumeration, bijection, cardinality, and physical event count remain *reserved*. Counting is an interface for finite updates and questions, not a guarantee hidden in a name.

## 2.2 The update and the alleged endpoint

The local capability stores an indexing capability, a named endpoint record, a current chain, and a chain update. The step capability stores that local capability, a current piece, a named accumulation, and a second update. A neutral grammar is

$$\begin{aligned}
 \text{chain} &::= \text{end}(\text{receipt}) \mid \text{entry}(\text{receipt}, \text{piece}, \text{chain}), \\
 \text{accumulation} &::= \text{empty}(\text{receipt}) \mid \text{indexed}(\text{receipt}, \text{chain}, \text{accumulation}), \\
 \text{local-process} &::= \text{local}(\text{indexing}, \text{named-endpoint}, \text{current-chain}, W), \\
 \text{step-process} &::= \text{step}(L_{\text{self}}, \text{current-piece}, \text{named-accumulation}, U).
 \end{aligned}$$

These capabilities posit records and updates. The words endpoint and accumulation attach names; they supply no metric, error bound, or statement about an infinite tail.

Let  $L_{\text{ext}}$  be a separately supplied local process, with stored chain  $S_{\text{ext}}$  and

update  $W_{\text{ext}}$ . Let  $A_*$  be the stored accumulation,  $O_\rho$  the origin piece made from receipt  $\rho$ , and  $V$  the indexing update on pieces. The external local update is

$$\begin{aligned} W_{\text{ext}}(\text{end}(\rho)) &= \text{entry}(\rho, O_\rho, S_{\text{ext}}), \\ W_{\text{ext}}(\text{entry}(\rho, p, T)) &= \text{entry}(\rho, V(p), S_{\text{ext}}). \end{aligned}$$

Both branches reset the output tail to  $S_{\text{ext}}$ ; the input tail  $T$  is discarded. The default step update is

$$\begin{aligned} U(\text{empty}(\rho)) &= \text{indexed}(\rho, S_{\text{ext}}, A_*), \\ U(\text{indexed}(\rho, S, T)) &= \text{indexed}(\rho, W_{\text{ext}}(S), A_*). \end{aligned}$$

Both branches reset the outer tail to  $A_*$ . This update also ignores the stored current piece in the step capability and the stored  $L_{\text{self}}$ ; it uses the chain and update from external dependency  $L_{\text{ext}}$ . A name suggesting approximation cannot restore any discarded tail or make either unused record operative.

**Worked example 1: equal head, unequal tails (default recurrence).**

Let

$$C_1 = \text{indexed}(\rho, S, T_1), \quad C_2 = \text{indexed}(\rho, S, T_2), \quad T_1 \neq T_2.$$

Then

$$U(C_1) = U(C_2) = \text{indexed}(\rho, W_{\text{ext}}(S), A_*).$$

The two inputs occupy one output fiber. This tail-reset equality is demonstrated directly from the recurrence. It does not establish that either input approaches the named accumulation.

For a finite run, choose origin  $x_0$  and budget  $N$ , then compute

$$x_{k+1} = U(x_k), \quad k = 0, \dots, N-1.$$

The budget certificate contains  $x_0$ ,  $N$ , all transition pairs, the  $N+1$  recorded

values, receipts retained in them, declared erasures, and BUDGET. It certifies only those transitions. A slice may ask any declared question  $q$  of the recorded values. Construction is covariant along  $U$ ; reading is contravariant:

$$U^*q = q \circ U, \quad (U \circ Z)^*q = Z^*(U^*q).$$

Stable accepted and rejected slices preserve within-slice order and lose cross-slice interleaving unless ledger positions are retained.

A Cauchy assertion requires independent mathematical data. On a declared metric space  $(X, d)$ , it is

$$\forall \varepsilon > 0 \exists M \forall m, n \geq M, \quad d(x_m, x_n) < \varepsilon.$$

A modulus gives a rule  $M(\varepsilon)$ . The local and step capabilities contain no metric or modulus, and their default recurrences do not establish this quantified statement.

**Worked example 2: fixed update (external metric experiment).**

Outside the capability grammar, declare  $X = \{3\}$ ,  $U(3) = 3$ , and the discrete metric. With  $N = 4$ , the finite trace is

$$[3, 3, 3, 3, 3].$$

All four transition pairs check. The separately declared rule and metric also justify the Cauchy assertion, because all pairwise distances vanish. The justification comes from those external premises, not from a named endpoint.

**Worked example 3: alternating update (external metric experiment).**

Declare  $X = \{-1, 1\}$ ,  $U(-1) = 1$ ,  $U(1) = -1$ , ordinary distance, and  $N = 6$ . The exact trace is

$$[-1, 1, -1, 1, -1, 1, -1].$$

For  $\varepsilon = 1$ , values arbitrarily far along the infinite-tail rule remain distance 2 apart. This supplies a counterexample to the Cauchy assertion while preserving a valid finite update trace.

**Worked example 4: first tolerance hit (external numerical experiment).** Declare  $x_0 = 1$ ,  $U(x) = x/2$ , alleged endpoint  $a = 0$ , ordinary distance, and  $\varepsilon = 1/8$ . Four transitions give

$$[1, \frac{1}{2}, \frac{1}{4}, \frac{1}{8}, \frac{1}{16}].$$

The question  $q_\varepsilon(x) = [d(x, a) \leq \varepsilon]$  first answers yes at ledger position three, producing

$$R_\varepsilon = [1, \frac{1}{2}, \frac{1}{4}], \quad A_\varepsilon = [\frac{1}{8}, \frac{1}{16}].$$

This demonstrates a first hit in the finite trace. Persistence requires an additional invariant; the stopping question alone does not certify the unread infinite tail.

Finite descriptions provide a separate computation-side resonance. Fix a declared prefix-free machine  $M$ . Its description complexity is

$$K_M(x) = \min\{|p| : M(p) = x\},$$

when at least one halting description  $p$  produces  $x$ . The value depends on the chosen machine and encoding. Chaitin's results connect such measures with algorithmic complexity and irreducibility [10]; they do not turn every short update into an uncomputability result.

**Worked example 5: description budget and displayed storage (external accounting model).** Assume a fixed-width display in which the update description occupies  $h$  symbols and each recorded value occupies ex-



actly  $b$  symbols. Displaying an  $N$ -transition trace then uses

$$h + (N + 1)b$$

symbols before receipts and formatting overhead. This is an accounting formula under the fixed-width assumptions, not a universal lower bound and not a value of  $K_M$ . A compact update may coexist with a long displayed trace.

The action ledger is exact. The local and step capabilities, their stored records, and the default  $W, U$  recurrences are *posited*. The two tail resets, ignored current piece, ignored stored local process, use of the external dependency, finite traces, output-fiber equality, and checked slices are *demonstrated*. Questions are *read* contravariantly through forward updates. A Cauchy property, modulus, persistent enclosure, mathematical endpoint, physical steady state, or algorithmic-complexity claim remains *reserved* until its additional premises are supplied. The capabilities update finite records; their names do not certify an infinite tail.

## 2.3 A finite trace is not an infinite object

A finite run and a request without a preset bound are different grammatical forms:

$$\begin{aligned} \textit{finite-run} &::= \text{run}(U, x_0, N), \\ \textit{open-request} &::= \text{continue}(U, x_0, \textit{stop-question}), \\ \textit{trace} &::= [x_0, x_1, \dots, x_k], \quad x_{j+1} = U(x_j). \end{aligned}$$

The first supplies a transition budget  $N$ . The second supplies a stopping question but no advance claim that the question will ever answer yes. Neither grammar stores an infinite object. Words such as iterative, infinite, or unbounded can name capabilities or intentions; they do not establish termi-

nation, nontermination, distinctness, or coverage.

For a budgeted run, apply  $U$  for  $j = 0, \dots, N-1$ . The certificate contains the origin, budget, all transition pairs, the  $N+1$  recorded values, retained receipts, declared erasures, and stop reason. It certifies exactly the visited prefix. Construction carries states forward through  $U$ . A question  $q$  about a state is read contravariantly:

$$U^*q = q \circ U, \quad (U \circ V)^*q = V^*(U^*q).$$

A slice may separate first occurrences from repeated occurrences, accepted from rejected states, or positions before and after a detected collision. Stable slicing preserves within-slice traversal order and loses interleaving unless indices remain attached.

Pigeonhole reasoning must be scoped. If a finite set  $X$  has  $m$  states, any list of  $m+1$  states contains a collision. That statement alone proves only

$$\exists i < j \leq m, \quad x_i = x_j.$$

It says nothing about how the list was generated. To infer an eventual periodic orbit, one additionally needs a single deterministic, time-independent update satisfying  $x_{k+1} = U(x_k)$ . Then equality at  $i < j$  implies by induction

$$x_{i+t} = x_{j+t} \quad (t \geq 0),$$

so the orbit is periodic from  $i$  with candidate period  $\lambda = j - i$ . Calling  $i$  the least preperiod  $\mu$ , or  $\lambda$  the least period, additionally requires first-repeat and minimality checks. Collision plus transition compatibility earns a periodic tail; collision alone does not.

**Worked example 1: collision without dynamics (computation).** Take the finite list

$$[a, b, a, c].$$

Pigeonhole identifies the repeated  $a$ . No update rule is declared, so there is no licensed next state after either occurrence and no cycle certificate. Calling the list an orbit would add a premise not present in the data.

**Worked example 1b: projected collision is not state collision (computation).** Let  $S = \{a, b, c\}$ , with

$$U(a) = b, \quad U(b) = c, \quad U(c) = c,$$

and observation projection

$$\pi(a) = 0, \quad \pi(b) = 0, \quad \pi(c) = 1.$$

The first two observations collide even though  $a \neq b$ ; their successors also project differently. Thus a projection collision proves neither state equality nor a cycle, and it does not select which underlying states should be identified. For  $m$  observed values in an alphabet of size  $r < m$ , finite pigeonhole reasoning proves only

$$\exists i < j < m, \quad \pi(x_i) = \pi(x_j).$$

The projection carries observations forward. A question  $q$  on the readout alphabet is read contravariantly:

$$\pi^* q = q \circ \pi.$$

Forward observation and pulled-back reading do not upgrade equality of readouts to equality of states.

A visited-map detector makes the missing premise explicit. Initialize a finite map  $H[x_0] = 0$ . For  $j = 1, \dots, N$ , compute  $x_j = U(x_{j-1})$ . If  $x_j \notin H$ , store  $H[x_j] = j$ . If  $x_j \in H$ , return

$$\mu = H[x_j], \quad \lambda = j - \mu,$$

together with the two equal states and all transitions between them. Stop with CYCLE, BUDGET, or RESOLVER-DOMAIN. This requires decidable state equality and a declared deterministic update.

**Worked example 2: visited-map cycle certificate (computation).**

Let  $U$  act on  $\{0, 1, 2, 3, 4\}$  by

$$U(0) = 1, \quad U(1) = 2, \quad U(2) = 3, \quad U(3) = 2, \quad U(4) = 4.$$

From  $x_0 = 0$ , the detector records

$$[0, 1, 2, 3, 2]$$

and stops at  $j = 4$ , because state 2 first appeared at index 2. It returns

$$\mu = 2, \quad \lambda = 2.$$

The transition pairs  $2 \mapsto 3 \mapsto 2$ , together with determinism, certify the repeating tail  $2, 3, 2, 3, \dots$ . State 4 is irrelevant to this orbit; a cycle certificate is not coverage of the state set.

Floyd's detector trades stored history for additional updates. Set the tortoise to  $U(x_0)$  and the hare to  $U(U(x_0))$ . Repeatedly update the tortoise once and the hare twice until they meet or a transition budget is exhausted. A meeting is a collision under the same deterministic update. Separate finite passes then recover  $\mu$  and  $\lambda$ . Its certificate must record every evaluated transition, the meeting state, budgets for all passes, and the recovered cy-

cle parameters. Without decidable equality or total updates on the visited states, the method has no licensed comparison.

**Worked example 3: Floyd detection on a finite orbit (computation).**

For the update in the previous example and origin 0,

| round    | 0 | 1 |
|----------|---|---|
| tortoise | 1 | 2 |
| hare     | 2 | 2 |

so the first meeting is at round one. A subsequent entry-point pass recovers  $\mu = 2$ , and a lap from state 2 recovers  $\lambda = 2$ . The arithmetic agrees with the visited-map result, but the algorithms have different traces and storage costs. Agreement of  $(\mu, \lambda)$  does not identify the methods.

**Worked example 4: budget exhaustion is inconclusive (computation).** Let  $U(n) = n + 1$  on nonnegative integer records, origin 0, and  $N = 5$ . The certificate is

$$[0, 1, 2, 3, 4, 5], \quad \text{BUDGET.}$$

No collision appears in the prefix. This proves only six distinct visited values. It proves neither that a later collision occurs nor that the run continues forever. The open request and the finite certificate must remain separate.

Turing’s halting theorem concerns a precisely encoded universal decision problem: no single algorithm decides for every encoded machine and input whether that machine halts [46]. It does not imply that a particular finite run lacks a stopping proof, and a budget stop is not evidence of undecidability. Conversely, a local cycle certificate settles the behavior of that declared deterministic orbit without solving the universal problem.

Shannon supplies a probability fence. Information calculations require a declared finite alphabet and readout map; empirical frequencies in one finite

trace are estimates, not automatically source probabilities. Entropy needs a probability distribution, while capacity additionally needs a channel law and an optimization domain [44]. Repetition introduced by an encoding or deterministic update is not, by itself, a stochastic process.

Chaitin supplies a second computation-side fence. Fix a declared prefix-free machine  $M$  and define

$$K_M(x) = \min\{|p| : M(p) = x\}.$$

The value depends on  $M$  up to the usual machine-dependent additive constants, and claims about algorithmic randomness require an infinite bit specification and the relevant complexity bounds [10]. A finite trace, a repeated state, or a capability labeled infinite supplies none of those assumptions. No uncomputability or randomness conclusion follows from naming alone.

**Worked example 5: finite trace storage (external accounting model).**

Under a fixed-width display with  $b$  symbols per state and  $h$  symbols for the update description, an  $N$ -transition trace occupies

$$h + (N + 1)b$$

symbols before receipts, indices, and formatting. This is a conditional storage calculation, not a universal lower bound and not  $K_M$ . A cycle certificate may compress future description to  $(\mu, \lambda)$  plus one period, but that compression is licensed only after the cycle has been demonstrated.

The action ledger closes the chapter. The update, origin, equality question, budgets, and stopping policies are *posited*. Finite traces, slices, collisions, transition compatibility, and checked cycle parameters are *demonstrated*. Questions are *read* contravariantly through covariant forward updates. An infinite object, universal termination decision, algorithmic randomness claim, or physical periodic process remains *reserved* until its ad-

ditional premises and provenance are supplied. A finite trace may certify a cycle; it never becomes an infinite object merely by being given an unbounded name.

## Chapter 3

# The Residue Refuses To Disappear

### 3.1 The branch that cannot be silently discarded

The residue capability consists of a step process and a replaceable representative question  $\mathcal{R}(a, b)$ . For neutral accumulation forms  $\text{empty}(\rho)$  and  $\text{indexed}(\rho, S, L)$ , its default branch table is

| $a$                                | $b$                                | $\mathcal{R}(a, b)$                                                    |
|------------------------------------|------------------------------------|------------------------------------------------------------------------|
| $\text{empty}(\rho_1)$             | $\text{empty}(\rho_2)$             | $\rho_1 = \rho_2$                                                      |
| $\text{empty}(\rho_1)$             | $\text{indexed}(\rho_2, S_2, L_2)$ | <b>yes</b>                                                             |
| $\text{indexed}(\rho_1, S_1, L_1)$ | $\text{empty}(\rho_2)$             | <b>no</b>                                                              |
| $\text{indexed}(\rho_1, S_1, L_1)$ | $\text{indexed}(\rho_2, S_2, L_2)$ | $[\rho_1 = \rho_2] \wedge Q_{\leq}(S_1, S_2) \wedge Q_{<}(L_1, L_2)$ . |

The capability supplies no resolver for this question and proves no reflexivity, symmetry, transitivity, or equivalence law. Another declaration may replace  $\mathcal{R}$ ; every subsequent read must then name the replacement.



A history-bearing sample has the grammar

$$sample ::= \text{initial}(\rho, L) \mid \text{response}(\rho_s, S, \rho_r, R, H).$$

The response stores a first receipt and stimulus, a second receipt and response, and prior history  $H$ . Its default comparison question  $\mathcal{C}(x, y)$  has four branches:

| $x$                                               | $y$                                               | $\mathcal{C}(x, y)$                                                                                      |
|---------------------------------------------------|---------------------------------------------------|----------------------------------------------------------------------------------------------------------|
| $\text{initial}(\rho_1, L_1)$                     | $\text{initial}(\rho_2, L_2)$                     | $[\rho_1 \sim \rho_2 \wedge Q_{\leq}(L_1, L_2)] \vee [\rho_1 \not\sim \rho_2 \wedge Q_{\leq}(L_1, L_2)]$ |
| $\text{response}(\rho_1, S_1, \rho'_1, R_1, H_1)$ | $\text{response}(\rho_2, S_2, \rho'_2, R_2, H_2)$ | $[\rho_1 \sim \rho_2 \wedge Q_{\leq}(R_1, R_2)] \vee [\rho_1 \not\sim \rho_2 \wedge Q_{\leq}(S_1, S_2)]$ |
| $\text{initial}(\rho_1, L_1)$                     | $\text{response}(\rho_2, S_2, \rho'_2, R_2, H_2)$ | $[\rho_1 \sim \rho_2 \wedge Q_{\leq}(L_1, S_2)] \vee [\rho_1 \not\sim \rho_2 \wedge Q_{\leq}(S_1, S_2)]$ |
| $\text{response}(\rho_1, S_1, \rho'_1, R_1, H_1)$ | $\text{initial}(\rho_2, L_2)$                     | no.                                                                                                      |

Here  $\rho_1 \sim \rho_2$  means derived receipt agreement. In the response/response branch, agreement selects the response comparison, while disagreement selects the stimulus comparison. The second receipts  $\rho'_1, \rho'_2$  and both history tails are ignored. Again, the table is a question, not an established order. The sample grammar supplies no resolver for  $\mathcal{C}$ .

**Worked example 1: receipt-selected comparison (computation).**

Take response samples  $X, Y$ . If their first receipts agree, the question reduces to  $Q_{\leq}(R_X, R_Y)$ , regardless of their stimuli, second receipts, or histories. If the first receipts disagree, it reduces to  $Q_{\leq}(S_X, S_Y)$ , regardless of their responses. Two samples can therefore receive the same comparison answer while carrying different ignored histories. Receipt relation selects an argument pair; it does not select which pair has physical meaning.

Forward history construction prepends responses:

$$\begin{aligned} H_0 &= \text{initial}(\rho_0, L_0), \\ H_1 &= \text{response}(\rho_1, S_1, \rho'_1, R_1, H_0), \\ H_2 &= \text{response}(\rho_2, S_2, \rho'_2, R_2, H_1), \\ H_3 &= \text{response}(\rho_3, S_3, \rho'_3, R_3, H_2). \end{aligned}$$

Traversal reads newest to oldest:  $H_3, H_2, H_1, H_0$ .

**Worked example 2: budgeted history traversal (computation).** With visit budget  $N = 2$ , traversal records the heads of  $H_3$  and  $H_2$ , then stops with BUDGET and returns  $H_1$  as the residual prior tail. Its certificate contains those two visited heads, their receipts and payloads, the untouched tail  $H_1$ , and the stop reason. It makes no claim about the contents of  $H_1$  or  $H_0$  beyond retaining their finite record.

Euclidean division is an external numerical model for explicit remainder accounting. For  $n \geq 0$ ,  $d > 0$ , initialize  $q = 0, r = n$ . While  $r \geq d$  and the budget remains, update

$$q \leftarrow q + 1, \quad r \leftarrow r - d.$$

The invariant is

$$n = dq + r.$$

A complete stop requires  $0 \leq r < d$ ; a budget stop preserves the invariant but does not certify a terminal Euclidean remainder.

**Worked example 3: iterative division (external numerical experiment).** For  $n = 23, d = 5$ , the trace is

$$(q, r) : (0, 23) \rightarrow (1, 18) \rightarrow (2, 13) \rightarrow (3, 8) \rightarrow (4, 3).$$

Each row satisfies  $23 = 5q + r$ . The complete certificate gives  $q = 4, r = 3$ , four transitions, the invariant checks,  $d > 0$ , and  $0 \leq 3 < 5$ . With budget two, the run stops at  $(2, 13)$ . It certifies  $23 = 5 \cdot 2 + 13$ , but not a completed division because  $13 \geq 5$ . This algorithm is an external arithmetic experiment, not a property supplied by the residue capability.

Projection makes information loss measurable. Define

$$\pi_{\text{now}}(\text{response}(\rho, S, \rho', R, H)) = (\rho, S, \rho', R), \quad \pi_{\text{full}}(H) = H.$$

The full projection is injective. The current projection is many-to-one whenever histories differ behind an equal head.

**Worked example 4: four histories, two summaries (Shannon model).**

Let four histories  $H_1, H_2, H_3, H_4$  be equiprobable, with

$$\pi_{\text{now}}(H_1) = \pi_{\text{now}}(H_2) = u, \quad \pi_{\text{now}}(H_3) = \pi_{\text{now}}(H_4) = v.$$

Under this declared discrete distribution,

$$H(\text{history}) = 2 \text{ bits}, \quad H(\text{summary}) = 1 \text{ bit}, \quad H(\text{history} \mid \text{summary}) = 1 \text{ bit}$$

by Shannon's definitions [44]. The missing bit measures unresolved history within each two-member fiber. It is not heat, energy, or a physical erasure cost.

If  $H_A \neq H_B$  but  $\pi_{\text{now}}(H_A) = \pi_{\text{now}}(H_B)$ , a separating question  $q_{\text{full}} : \text{History} \rightarrow \text{Answer}$  exists on the full finite history set, for example

$$q_{\text{full}}(H) = [H = H_A].$$

Then  $q_{\text{full}} \circ \pi_{\text{full}}$  distinguishes them. For any  $q_{\text{now}} : \text{Summary} \rightarrow \text{Answer}$ , no  $q_{\text{now}}$  can separate equal summaries after  $\pi_{\text{now}}$ , because both inputs have the same image. Construction carries histories forward through either projection;

reading pulls the separately typed questions back:

$$\pi_{\text{full}}^* q_{\text{full}} = q_{\text{full}} \circ \pi_{\text{full}}, \quad \pi_{\text{now}}^* q_{\text{now}} = q_{\text{now}} \circ \pi_{\text{now}}.$$

A checksum gives another lossy external projection. For bytes  $b_1, \dots, b_m$ , let  $c = (\sum_j b_j) \bmod 7$ . Records  $[1, 2, 3]$  and  $[0, 0, 6]$  share checksum six while remaining distinct. Equality of checksums is not equality of records; the collision certificate identifies only their common projection.

**Worked example 5: truncated acquisition (physics, then interpretation).** Suppose an apparatus acquires five threshold-event records but a display retains three. Computation can certify a three-record prefix, two-record dropped suffix, and coverage  $3 + 2 = 5$ . Physics owns the sensor, threshold, acquisition window, dead time, and missed-event model. The displayed three is not a physical event count until that apparatus map and its uncertainty are supplied.

The action ledger is explicit. The residue capability, exact default representative and sample questions, history updates, projections, budgets, and external arithmetic experiments are *posited*. The branch reductions, ignored second receipts and histories, three-response trace, residual prior tail, division invariant, Shannon fiber loss, checksum collision, and truncation coverage are *demonstrated*. Questions are *read* contravariantly through covariant projections and updates. Order laws, numerical remainder meaning, physical residue, and erasure energy remain *reserved*. A branch may be discarded intentionally, but the certificate must state which branch, which fiber, and what information can no longer be recovered.

### 3.2 Iterative error versus structural residue

A state and a structural residue answer different questions. The recurrence evolves states,

$$x_{k+1} = T(x_k).$$

A state error exists only relative to a declared reference  $x_*$ , as  $e(x) = x - x_*$ . A structural residue is a retained branch, tail, receipt, or complement slice that a projection does not display. The former may shrink, grow, alternate, or reach a fixed point. The latter may have no numerical magnitude at all. Calling both a remainder does not identify them.

For state  $x$ , alleged target  $x_*$ , update  $T$ , finite approximation  $s_N$ , exact target  $s_*$ , and rounding map  $Q$ , the categories are

| category             | definition                          | required extra data |
|----------------------|-------------------------------------|---------------------|
| state error          | $e(x) = x - x_*$                    | $x_*$               |
| fixed-point residual | $r(x) = T(x) - x$                   | $T$                 |
| truncation error     | $\tau_N = s_* - s_N$                | $s_*$               |
| rounding error       | $\eta_Q(x) = Q(x) - x$              | $Q$                 |
| structural residue   | retained branch or projection fiber | record projection   |

A dropped suffix is not automatically numerical. Without a valuation into an additive quantity, its numerical error entry is *absent*, not guessed.

Use the finite audit grammar

$$\begin{aligned}
\textit{state-run} &::= \text{run}(x_0, T, N, \textit{stop-question}), \\
\textit{loss-ledger} &::= \text{ledger}(\textit{stopping-loss}, \textit{rounding-loss}, \textit{truncation-loss}, \textit{retained-branch}), \\
\textit{stop} &::= \text{BUDGET} \mid \text{TOLERANCE} \mid \text{FIXED} \mid \text{RESOLVER-DOMAIN}.
\end{aligned}$$

For budget  $N$ , compute  $x_{k+1} = T(x_k)$  for  $k = 0, \dots, N - 1$ . The certificate contains  $x_0$ ,  $N$ , all transition pairs, the  $N + 1$  visited states, declared projections and erasures, and stop reason. State errors  $e$  and fixed-point residuals

$r$  are computed ledger columns when their required data are supplied. The certificate certifies no unvisited update.

If an externally justified target  $x_*$  exists, a stopping report may include state error  $d(x_N, x_*)$  under a declared distance. For a finite record  $L$  sliced at  $K$ , structural truncation loss is the dropped suffix  $\text{drop}_K(L)$ , or a declared summary of it. None of these is supplied merely by the recurrence.

Fixed-point and tolerance questions are reads on forward updates. For

$$q_{\text{fix}}(x) = [T(x) = x], \quad q_\varepsilon(x) = [d(x, x_*) \leq \varepsilon],$$

the pulled-back questions are

$$T^*q = q \circ T.$$

The update is covariant; reading is contravariant. A yes-answer to  $q_{\text{fix}}$  certifies one fixed value. A first yes-answer to  $q_\varepsilon$  certifies one visited crossing. Neither answer alone proves attraction, uniqueness, or persistence.

**Worked example 0: fixed residual, nonzero state error (external numerical experiment).** Let  $T$  be the identity,  $x_* = 0$ , and  $x_0 = 7$ . Then

$$r(x_0) = T(7) - 7 = 0, \quad e(x_0) = 7 - 0 = 7.$$

The update is at a fixed point, but that point is not the alleged target. A zero fixed-point residual does not imply zero state error without uniqueness or attraction premises.

**Worked example 1: decreasing update error (external numerical experiment).** Let  $x_0 = 8$ ,  $T(x) = x/2$ , target  $x_* = 0$ , the ordinary norm, and  $N = 4$ . Then

$$x_k = \frac{8}{2^k}, \quad e(x_k) = \frac{8}{2^k}, \quad r(x_k) = -\frac{4}{2^k},$$

with trace  $[8, 4, 2, 1, \frac{1}{2}]$ . If  $T$  is a contraction satisfying  $\|T(x) - T(y)\| \leq L\|x - y\|$ ,  $0 \leq L < 1$ , with fixed point  $x_* = T(x_*)$ , then

$$\|e(x)\| \leq \frac{\|r(x)\|}{1 - L}.$$

For halving  $L = 1/2$ , so the bound is exact:  $\|e(x_k)\| = 2\|r(x_k)\|$ . For tolerance  $\varepsilon = 1$ , the first accepted value is  $x_3 = 1$ . A first-hit policy stops there with TOLERANCE; a budget policy records one more update. The trace demonstrates finite decrease. Persistence follows only after separately establishing  $0 \leq T(x) \leq x$  for nonnegative  $x$ .

**Worked example 2: fixed magnitude without error reduction (external numerical experiment).** Let  $T(e) = -e$ ,  $e_0 = 1$ , and  $N = 4$ . Then

$$[1, -1, 1, -1, 1].$$

The magnitude  $|e_k| = 1$  is fixed while the signed error alternates. A question about magnitude returns a constant answer; a question about signed equality does not. Selecting the former projection cannot justify a claim that the full error reached a fixed point.

**Worked example 3: nearest-hundredth ledger (external numerical experiment).** Declare  $Q_2$  to round to the nearest hundredth, with exact halfway cases sent to the candidate whose final decimal digit is even. Then

$$Q_2(\frac{1}{3}) = 0.33 = \frac{33}{100}, \quad \eta_{Q_2}(\frac{1}{3}) = \frac{33}{100} - \frac{1}{3} = -\frac{1}{300}.$$

The certificate records exact input, rounded output, scale, tie rule, and signed error. Changing the tie rule can change halfway cases; the name rounding does not select a policy.

**Worked example 4: truncation and retained complement (computation).** For

$$L = [\mathbf{a}, \mathbf{b}, \mathbf{c}, \mathbf{d}, \mathbf{e}], \quad K = 3,$$

stable slicing gives

$$\text{keep}_3(L) = [\mathbf{a}, \mathbf{b}, \mathbf{c}], \quad \text{drop}_3(L) = [\mathbf{d}, \mathbf{e}].$$

Coverage is  $3 + 2 = 5$ . The dropped slice is structural residue relative to the retained-prefix projection. It is not a scalar error until a separate valuation maps records to numbers. Erasing the suffix destroys reconstruction; keeping it in a complement ledger makes the loss explicit.

A history-bearing response provides the same distinction. Let

$$H_A = \text{response}(\rho, S, \rho', R, H_1), \quad H_B = \text{response}(\rho, S, \rho', R, H_2), \quad H_1 \neq H_2.$$

The current projection maps both to  $(\rho, S, \rho', R)$ . Their tails form a structural distinction even when the current numerical error is equal. A full-history question can separate them; no question on the common current summary can.

**Worked example 4b: projected fixed point, distinct histories (computation).** Let a current-summary update be the identity. The shared summary  $z = \pi_{\text{now}}(H_A) = \pi_{\text{now}}(H_B)$  has projected residual

$$T(z) - z = 0.$$

Nevertheless  $H_A \neq H_B$  because their tails differ. The zero projected residual certifies stability of the summary only; it does not eliminate the structural residue in the history fiber.



**Worked example 4c: finite geometric truncation (external numerical experiment).** For

$$s_N = \sum_{k=1}^N 2^{-k} = 1 - 2^{-N},$$

declare exact target  $s_* = 1$ . Then

$$\tau_N = s_* - s_N = 2^{-N}.$$

The finite run owns the computed  $s_N$ . The symbolic error formula additionally depends on the declared target and the separately established geometric identity; a finite prefix alone does not own an unrecorded infinite tail.

**Worked example 5: finite information loss (Shannon model).** Suppose four histories are equiprobable and a summary projection groups them into two equal fibers. Then, under that declared distribution,

$$H(\text{history}) = 2 \text{ bits}, \quad H(\text{summary}) = 1 \text{ bit}, \quad H(\text{history} \mid \text{summary}) = 1 \text{ bit}$$

by Shannon's definitions [44]. The conditional bit quantifies unresolved history in this model. It is not a temperature, energy, or universal cost of erasure.

A physical error requires another provenance chain. If a sensor report is  $y_k$  and a calibrated reference is  $y_*$ , physics may define  $e_k = y_k - y_*$  with units and uncertainty. Computation can update, slice, and round the finite recorded values. It cannot infer that  $y_*$  is physically correct, that noise is independent, or that an unobserved environment received discarded information. Those are experimental premises.

The error ledger for a run is therefore

$$\mathcal{E} = (x_0, T, N, d, x_*, Q, K, \text{trace}, \text{slices}, e, r, \tau, \eta, \text{structural branches}, \text{valuations}, \text{stop}).$$

Every unjustified entry is tagged *absent* rather than guessed. Every supplied target, norm, contraction constant, rounding rule, and valuation is tagged *external* or *demonstrated*. A budget stop reports the last visited state and any error or residual columns justified by supplied reference data, together with the unexecuted suffix of the request. A tolerance stop reports its first accepted index. A fixed stop reports the checked pair  $(x, T(x))$ . A truncation report records both slice sizes and the fate of the complement.

The action ledger closes the distinction. Updates, metrics, targets, tolerances, contraction premises, rounding rules, truncation budgets, valuations, and projections are *posited*. Finite traces, first hits, fixed pairs, residual identities, contraction bounds under their premises, local rounding losses, geometric finite sums, coverage identities, and projection fibers are *demonstrated*. Error and history questions are *read* contravariantly through covariant updates and projections. Attraction, persistent enclosure, physical accuracy, erasure energy, and numerical meaning for a structural branch remain *reserved*. Iterative error may be reduced; structural residue must be accounted for even when it is not a number.

### 3.3 Observation becomes history

An observation capability stores a step process, records named before  $B$  and after  $A$ , a relative question  $Q_{\text{rel}}$ , and a sample update  $W$ . For accumulation forms  $\text{empty}(\rho)$  and  $\text{indexed}(\rho, S, L)$ , the relative question has the exact shape table

| $B$                                | $A$                                | $Q_{\text{rel}}(B, A)$ |
|------------------------------------|------------------------------------|------------------------|
| $\text{empty}(\rho_1)$             | $\text{empty}(\rho_2)$             | $R(\rho_1, \rho_2)$    |
| $\text{indexed}(\rho_1, S_1, L_1)$ | $\text{indexed}(\rho_2, S_2, L_2)$ | $R(\rho_1, \rho_2)$    |
| $\text{empty}(\rho_1)$             | $\text{indexed}(\rho_2, S_2, L_2)$ | no                     |
| $\text{indexed}(\rho_1, S_1, L_1)$ | $\text{empty}(\rho_2)$             | no.                    |

Here  $R$  asks whether the first receipts agree in the declared sense. The question ignores both chains and tails, and the capability supplies no resolver. Its name proves no temporal order, causality, or covariance.

Samples have the grammar

$$sample ::= \text{initial}(\rho, L) \mid \text{response}(\rho_s, S, \rho_r, R_s, sample).$$

With fixed receipt  $\rho_*$ , the default observation update is

$$\begin{aligned} W(\text{initial}(\rho, L)) &= \text{response}(\rho, L, \rho, A, \text{initial}(\rho, L)), \\ W(\text{response}(\rho_s, S, \rho_r, R_s, H)) &= \text{response}(\rho_s, S, \rho_*, A, \text{response}(\rho_s, S, \rho_r, R_s, H)). \end{aligned}$$

The update retains the entire prior sample. It preserves the first receipt and stimulus and replaces the current response receipt and response. It ignores the stored  $B$  and the stored step process.

For origin  $H_0$  and budget  $N$ , compute  $H_{k+1} = W(H_k)$  for  $k = 0, \dots, N - 1$ . The certificate contains all transition pairs,  $N + 1$  heads, retained receipts, declared erasures, and stop reason. Tail traversal reads newest to oldest and returns the unvisited prior tail on BUDGET.

**Worked example 1: accumulating observation history (computation).** For  $H_0 = \text{initial}(\rho_0, L_0)$ ,

$$\begin{aligned} H_1 &= \text{response}(\rho_0, L_0, \rho_0, A, H_0), \\ H_2 &= \text{response}(\rho_0, L_0, \rho_*, A, H_1), \\ H_3 &= \text{response}(\rho_0, L_0, \rho_*, A, H_2). \end{aligned}$$

A traversal budget two records the heads of  $H_3, H_2$ , returns  $H_1$ , and certifies nothing about the unvisited tail beyond retaining it.

A trial layer uses a different grammar:

$$\begin{aligned}
 \textit{trial} &::= \text{hypothesis}(\rho, \textit{sample}) \\
 &\quad | \text{trial-response}(\rho_1, \textit{sample}, \rho_2, \textit{sample}, \textit{trial}), \\
 \textit{repeat-process} &::= \text{repeating}(\textit{observation-capability}, \textit{stored-stimulus}, \textit{stored-expectation}, P), \\
 \textit{repeat-interface} &::= \text{repeatability}(\textit{repeat-process}, \textit{typical-question}).
 \end{aligned}$$

Let  $S_*$  be the stored stimulus,  $E_*$  the stored expectation, and  $\rho_T, \rho_D$  fixed receipts. The default trial update is

$$\begin{aligned}
 P(\text{hypothesis}(\rho, S)) &= \text{trial-response}(\rho, S, \rho_T, S_*, E_*), \\
 P(\text{trial-response}(\rho_1, S_1, \rho_2, S_2, T)) &= \text{trial-response}(\rho_2, S_*, \rho_D, S_2, E_*).
 \end{aligned}$$

The second branch ignores the first receipt  $\rho_1$ , first sample  $S_1$ , and prior tail  $T$ . Both branches reset the prior tail to the stored expectation  $E_*$ . Thus  $W$  accumulates old samples, while  $P$  resets old trial history.

**Worked example 2: trial reset fiber (computation).** Let

$$T_1 = \text{trial-response}(\rho_a, S_a, \rho_2, S_2, U_1), \quad T_2 = \text{trial-response}(\rho_b, S_b, \rho_2, S_2, U_2),$$

where the ignored first fields and tails may all differ. Then

$$P(T_1) = P(T_2) = \text{trial-response}(\rho_2, S_*, \rho_D, S_2, E_*).$$

The update has a many-to-one fiber. Equality after the reset does not reconstruct either prior trial.

The outer repeat interface supplies the default typical-response question

$$\text{typical}(s, t) := [P(s) = t].$$

No resolver is supplied for it. This unresolved question proves neither that two runs agree, that repeated applications return one result, that a cadence

is constant, nor that a laboratory experiment is repeatable. Those require finite run certificates and, for physics, an apparatus protocol.

Current and full-history projections remain distinct:

$$\pi_{\text{current}}(\text{response}(\rho_s, S, \rho_r, R_s, H)) = (\rho_s, S, \rho_r, R_s), \quad \pi_{\text{full}}(H) = H.$$

The first is many-to-one; the second is injective. Construction moves histories forward through  $W$  or projections. Questions are read contravariantly:

$$\pi_{\text{current}}^* q_c = q_c \circ \pi_{\text{current}}, \quad \pi_{\text{full}}^* q_h = q_h \circ \pi_{\text{full}}.$$

No current-summary question separates equal summaries with different tails.

**Worked example 3: current slice versus accumulated history (computation).** Let two histories have equal current heads and unequal tails. A current slice puts them in one fiber, while the full-history question  $q_A(H) = [H = H_A]$  separates them. With a traversal budget  $N$ , coverage and exclusivity apply only to the  $N$  visited rows; a stopped traversal makes no statement about its returned tail.

A sufficient statistic requires a declared probability family and target. Let  $X, Y$  be independent uniform bits and let the target be their parity

$$\Theta = X \text{ xor } Y.$$

Use summary  $Z = \Theta$ . The four histories are equiprobable, so

$$H(X, Y) = 2, \quad H(Z) = 1, \quad H(X, Y \mid Z) = 1$$

in bits. The summary discards one bit of history but is sufficient for the declared parity target because

$$H(\Theta \mid Z) = 0.$$

By contrast, the first bit alone is not sufficient:

$$H(\Theta \mid X) = 1.$$

**Worked example 4: sufficiency is target-relative (Shannon model).**

The parity summary perfectly answers the parity question but cannot reconstruct which of the two histories in its fiber occurred. These entropy identities require the displayed joint distribution [44]. Changing the target or distribution requires a new sufficiency calculation. The discarded bit is information in this model, not heat or energy.

**Worked example 5: apparatus history (physics, then interpretation).**

Suppose a detector creates time-stamped threshold records. Physics owns the sensor, clock calibration, threshold, dead time, missed-record rate, and uncertainty. Computation can apply  $W$ ,  $P$ , project the latest record, and traverse finite history. Neither prepend orientation nor a typical-response question establishes physical time, constant cadence, causality, or experimental repeatability.

The observation ledger is

$$\mathcal{O} = (B, A, Q_{\text{rel}}, W, \text{repeat-process}, P, \text{repeat-interface}, S_*, E_*, N, \text{traces}, \pi_{\text{current}}, \pi_{\text{full}}, \text{slices}, \text{probability})$$

Absent resolvers, distributions, apparatus maps, cadence models, and uncertainty are marked absent.

The action ledger is exact. The observation capability, trial grammar, inner repeat process, outer repeat interface, default  $Q_{\text{rel}}, W, P$  and typical question, budgets, projections, probability family, and target are *posited*. The ignored stored records and arguments, accumulating observation trace, resetting trial fiber, finite slices, entropy identities, and target-relative parity sufficiency are *demonstrated*. Current and history questions are *read* contravariantly through covariant updates and projections. Temporal order,

cadence, causal direction, whole-history recovery from a current summary, laboratory repeatability, and sufficiency outside the declared model remain *reserved*. Observation becomes history under  $W$ ; the later trial update  $P$  deliberately resets that history to its stored expectation.

## Part II

### The Turn



# Chapter 4

## The Tower Turns Around

### 4.1 Requirements pull backward

A requirement is a question on reported values. Let  $S$  be a finite setup domain,  $R$  a report domain, and

$$O : S \rightarrow R$$

a forward observation map. Let  $Q$  be requirement syntax on  $R$ , and separately posit a partial evaluator

$$\text{eval}_Q : R \rightarrow \text{Option}\{\text{yes}, \text{no}\}.$$

Its resolved reading pulls backward to setups:

$$O^* \text{eval}_Q = \text{eval}_Q \circ O.$$

Construction or observation carries a setup forward to a report; reading asks which setups satisfy a report-side question. The requirement does not choose a setup, construct a witness, or perform an experiment. The evaluator may return evidenced yes, evidenced no, or no value; the last case leaves  $Q$

unresolved.

Composition reverses:

$$S \xrightarrow{O} R \xrightarrow{G} T, \quad (G \circ O)^* \text{eval}_P = O^*(G^* \text{eval}_P)$$

for question syntax  $P$  on  $T$  and its partial evaluator. The outer evaluator is pulled through  $G$  first and then through  $O$ . This contravariant reading is distinct from reversing arguments inside a comparison rule.

An admissible fiber is

$$F_Q = \{s \in S : \text{eval}_Q(O(s)) = \text{some}(\text{yes})\}.$$

Using the word admissible requires a declared  $Q$ . The set  $F_Q$  is defined even when  $S = \emptyset$ , and may itself be empty for nonempty  $S$ . Claiming that a selectable setup exists requires

$$F_Q \neq \emptyset.$$

Even that statement is only an existence requirement. It does not select an element of  $F_Q$ ; selection and witness construction belong to the next action.

For a family  $(F_i)_{i \in I}$ , a global selection claim requires

$$\forall i \in I, \quad F_i \neq \emptyset.$$

For a finite explicitly listed family with decidable membership, exhaustive search selects the first member of each fiber and certifies every choice; no axiom of choice is needed. An arbitrary family requires a separate selection principle and still cannot select from an empty fiber.

Tolerance must also be explicit. A raw metric test

$$d(r, r') \leq \varepsilon$$

is reflexive and symmetric but generally not transitive, so it is not automatically an equivalence relation. For a genuine tolerance equivalence, declare a finite bin map  $b_{h,c} : R \rightarrow B$ , with scale  $h$  and boundary convention  $c$ , and define

$$r \equiv_{h,c} r' \iff b_{h,c}(r) = b_{h,c}(r').$$

Equality in  $B$  makes this relation reflexive, symmetric, and transitive. A desired report bin  $\beta \in B$  gives

$$\text{eval}_{Q_\beta}(r) = \begin{cases} \text{some}(\text{yes}), & b_{h,c}(r) = \beta, \\ \text{some}(\text{no}), & b_{h,c}(r) \neq \beta, \end{cases} \quad F_\beta = O^{-1}(b_{h,c}^{-1}(\{\beta\})).$$

The scale  $h$ , cell diameter, boundary convention, and units are part of the requirement. A bin scale is not automatically the metric tolerance  $\varepsilon$ ; their relation must be demonstrated for the chosen geometry.

**Worked example 1: a nonempty tolerance fiber (external numerical experiment).** Let

$$S = \{\mathbf{a}, \mathbf{b}, \mathbf{c}\}, \quad O(\mathbf{a}) = 1.02, \quad O(\mathbf{b}) = 1.11, \quad O(\mathbf{c}) = 1.27,$$

and use left-closed bins  $b_{0.1,\text{left}}(r) = \lfloor 10r \rfloor$ , whose cells have diameter 0.1. For desired bin  $\beta = 11$ ,

$$F_{11} = \{\mathbf{b}\}.$$

The fiber is nonempty, so an admissible setup exists. Its singleton form identifies  $\mathbf{b}$  as the sole admissible candidate, but no subsequent setup-selection action has yet been performed.

A bounded filter realizes a finite tolerance read. Given candidates  $S_N = [s_1, \dots, s_N]$ , target  $r_*$ , tolerance  $\varepsilon$ , and budget  $B \leq N$ , visit  $i = 1, \dots, B$ , compute

$$y_i = O(s_i), \quad \delta_i = d(y_i, r_*),$$

and retain  $s_i$  exactly when  $\delta_i \leq \varepsilon$ . Stop with COMPLETE when  $B = N$ , BUDGET when  $B < N$ , or RESOLVER-DOMAIN. The certificate contains candidate indices, readings, residuals, evidenced decisions, retained and rejected slices,  $N, B$ , completeness, erasures, and stop reason. An empty retained slice means only that no visited candidate passed. It proves an empty fiber only when  $S_N$  is certified as a complete enumeration of  $S$  and the run is complete.

Projecting a residual to a pass bit or bin erases information. Equal residuals do not identify candidates or readings; equal pass bits erase still more. The report records the fate of residual, pass, and bin values.

**Worked example 2: a requirement can be unsatisfied (external numerical experiment).** With the same domain and observation map, desired bin  $\beta = 14$  gives

$$F_{14} = \emptyset.$$

The requirement remains a well-formed question even though no supplied setup meets it. Neither a name nor a choice principle manufactures a member of an empty fiber.

Requirements may themselves depend on prior requirements. Represent a finite dependency record as nodes  $(v, Q_v)$  and directed edges  $v \rightarrow u$ , meaning that resolving  $Q_v$  requires the report for  $u$ . A budgeted backward traversal uses a worklist initialized with the requested node. At each visit:

1. remove one unresolved requirement  $v$ ,
2. append each unvisited dependency  $u$  to the worklist,
3. when dependencies are resolved, evaluate the declared resolver for  $Q_v$ ,
4. record  $v$  in the yes, no, or unresolved slice.

Stop with COMPLETE, BUDGET, RESOLVER-DOMAIN, or DEPENDENCY-CYCLE. The certificate contains the requested node, visited nodes and edges,

traversal order, all resolver outputs, unresolved worklist, slices, budget, and stop reason. It certifies only the visited dependency subgraph.

**Worked example 3: finite backward dependency traversal (computation).** Let requirement  $A$  depend on  $B, C$ , and  $C$  depend on  $D$ :

$$A \rightarrow \{B, C\}, \quad C \rightarrow D.$$

A depth-first backward traversal may visit  $A, B, C, D$ . With budget three, it stops after  $A, B, C$ , returns  $D$  on the unresolved worklist, and cannot resolve  $C$  or  $A$ . The certificate demonstrates the explored edges and budget stop; it does not infer the missing answer.

**Worked example 4: composed reading (computation).** Let  $O : S \rightarrow R$  report an integer and  $G : R \rightarrow \{\text{even}, \text{odd}\}$  report parity. For the question  $P$  whose evaluator returns `some(yes)` on even and `some(no)` on odd,

$$(G \circ O)^* \text{eval}_P(s) = \text{eval}_P(G(O(s))) = O^*(G^* \text{eval}_P)(s).$$

The equality demonstrates two descriptions of one pulled-back question. It does not identify the maps  $O, G$ , their costs, or their report domains.

Cohen’s finite-condition method offers a computation-side resonance. A finite partial assignment

$$p = \{v_1 \mapsto a_1, \dots, v_k \mapsto a_k\}$$

states requirements on finitely many coordinates. An extension  $q \supseteq p$  must preserve earlier assignments. This resembles a dependency ledger whose resolved requirements cannot be silently changed [11]. In forcing vocabulary, a condition set  $D_Q$  is dense precisely when

$$\forall p \exists q \supseteq p, \quad q \in D_Q,$$

where membership in  $D_Q$  means that  $q$  decides or satisfies the declared requirement  $Q$ . The finite worklist does not demonstrate this quantified density statement. The analogy stops there: no forcing relation, generic filter, generic extension, or independence result follows from a finite worklist.

**Worked example 5: physical acceptance window (physics, then interpretation).** Suppose an apparatus reports a voltage  $O(s)$  for each finite setup. Physics owns the apparatus, voltage units, calibration, noise model, and repeated observations. A bin requirement can classify recorded voltages, but it does not show that a setup will reproduce the bin on another trial. Nonempty empirical fibers and physical reproducibility need their own uncertainty and sampling certificates.

The requirement audit is

$$\mathcal{D} = (S, R, O, Q, \text{eval}_Q, \varepsilon, b_{h,c}, h, c, \text{cell diameter}, \text{dependency graph}, \text{resolvers}, S_N, \text{candidate indices},$$

An absent resolver, empty domain, empty admissible fiber, or unvisited dependency is reported directly.

The action ledger closes the section. Domains, forward maps, question syntax, partial evaluators, metric tolerances, bin maps, dependency edges, candidate lists, and budgets are *posited*. Pullback identities, bin-equivalence laws, coverage and exclusivity invariants, completeness implications, finite exhaustive-selection correctness, and quantified nonemptiness checks are *demonstrated*. Actual traversal order, visited nodes, worklist, stop reason, resolver outputs, fiber membership, residual comparisons, and dependency answers are *read* through contravariant questions on covariant forward maps. Arbitrary or global selection, witness action, physical reproducibility, forcing conclusions, and answers outside the visited subgraph remain *reserved*. Requirements pull backward; actions that satisfy them come later.

## 4.2 Witnesses run forward

A requirement, an existence statement, a selection, and a witness are four different records. Retain the setup and report domains  $S, R$ , forward map  $O : S \rightarrow R$ , question syntax  $Q$ , and partial evaluator

$$\text{eval}_Q : R \rightarrow \text{Option}\{\text{yes}, \text{no}\}.$$

The requirement read is contravariant:

$$O^* \text{eval}_Q = \text{eval}_Q \circ O.$$

Its admissible fiber is

$$F_Q = \{s \in S : \text{eval}_Q(O(s)) = \text{some}(\text{yes})\}.$$

The statement  $F_Q \neq \emptyset$  establishes existence. A selector action must then produce a particular  $s \in F_Q$ . Only after that action can forward construction package a witness

$$w = \text{witness}(s, y, c_O, c_Q, c_D),$$

where  $y = O(s)$ ,  $c_O$  certifies the forward evaluation,  $c_Q$  is the evidenced yes-answer, and  $c_D$  records dependency certificates. Neither the requirement nor nonemptiness contains this record.

For a finite candidate list  $S_N = [s_1, \dots, s_N]$ , a budgeted witness search visits at most  $B \leq N$  candidates. For each visited  $s_i$ :

1.  $y_i \leftarrow O(s_i)$ ,
2.  $a_i \leftarrow \text{eval}_Q(y_i)$ ,
3. append  $(i, s_i, y_i, a_i)$  to yes, no, or unresolved slice,
4. on the declared first-yes policy, construct  $w_i$  and stop.

Stop reasons are WITNESS, COMPLETE-NO, BUDGET, and RESOLVER-DOMAIN.

The certificate contains candidate indices, forward reports, evaluator answers, all slices, selected index when present, budget, completeness status, erasures, and stop reason. A budget stop with no yes-answer does not prove  $F_Q$  empty. A complete-no result proves emptiness only when  $S_N$  is a certified complete enumeration of  $S$ .

**Worked example 1: finite witness search (computation).** Let  $S = \{2, 3, 4, 5\}$ ,  $O(s) = s^2$ , and let  $Q$  ask whether the report exceeds 15. In list order, the reports are 4, 9, 16, 25. A first-yes policy selects  $s = 4$  and returns

$$w = \text{witness}(4, 16, [16 = 4^2], [16 > 15], c_D).$$

The backward requirement identifies the fiber  $\{4, 5\}$ ; the forward action selects one member according to the declared list order. A different selection policy may return 5 without changing the requirement.

This finite exhaustive selection uses the explicit list and evidenced evaluator; it needs no axiom of choice. Selecting simultaneously from an arbitrary family of nonempty fibers is a separate global premise and remains reserved.

Dependencies require an exact two-pass finite algorithm. First traverse edges backward from the requested target, record the reverse dependency slice, reject missing predecessors, test acyclicity, and compute a topological order with predecessors before consumers. Then, for each node in that order:

1. wait until every predecessor certificate is present,
2. form a cache key from inputs, predecessor results, policies, and versions,
3. if an entry exists, run its declared checker and reuse only on success,
4. otherwise spend construction budget, run the node builder, run the checker,
5. memoize only a checked result and append the node to the new or reused slice.

Stops MISSING-DEPENDENCY, CHECKER-FAILURE, BUDGET, and DEPENDENCY-CYCLE return the unresolved frontier. None proves an admissible fiber empty.



The run ledger slices nodes into newly constructed, verified-reused, checker-failed, and unresolved rows.

**Worked example 2: diamond dependency and memoized reuse (computation).** Let

$$A \leftarrow \{B, C\}, \quad B \leftarrow D, \quad C \leftarrow D.$$

The reverse slice has four nodes, and one valid forward order is

$$D, B, C, A.$$

With an empty cache and sufficient budget, the run invokes four builders;  $D$  is constructed once and its checked certificate is consumed by both  $B$  and  $C$ . If a valid  $D$  entry is preloaded, the run performs one cache verification and three new constructions, for  $B, C, A$ . The ledger reports exactly three new results, one verified reuse, and the dependency edges each consumer checked.

**Worked example 3: cache invalidation and erasure (computation).**

A cache entry is invalid when any dependency result, builder version, checker version, policy, configuration, forward map  $O$ , or evaluator changes. Matching a node name alone is unsound. A compact hash key is acceptable only under an explicit collision-resistance or collision-checking assumption; hash equality is not mathematical identity. Storing only the target result may preserve a value while erasing construction trace, checker trace, and cost. The audit must record each such erasure rather than treating a cache hit as a complete historical witness.

Turing’s framework separates executable evidence from oracle assertion. For a declared deterministic transition system and encoded input, a finite halting trace

$$c_0 \rightarrow c_1 \rightarrow \cdots \rightarrow c_m = \text{halt}$$

is an executable run-specific witness: every transition and the final halt condition can be checked. An assertion  $\mathcal{H}(M, x)$  supplied by an oracle without such a trace is openly posited, not constructed. For an effective universal machine model, Turing’s result rules out a total computable procedure that correctly decides halting for every encoded machine-input pair [46]. It does not invalidate a finite halting trace for one run, and such a trace witnesses only the positive halting case.

**Worked example 4: trace witness versus oracle assertion (computation).** Consider a finite transition rule that decrements a positive counter and halts at zero. From three, the trace

$$3 \rightarrow 2 \rightarrow 1 \rightarrow 0 \rightarrow \text{halt}$$

is a checkable witness under the declared rule. The bare assertion *this run halts* may be correct, but without the rule and trace it is only an assertion. Conversely, a finite budget stop before halt is not evidence that the run never halts.

**Worked example 5: apparatus witness (physics, then interpretation).** Suppose a setup produces a voltage inside a declared acceptance bin. Physics owns the apparatus, calibration, units, noise model, and repeated measurements. A computational witness can package the selected setup, recorded voltage, bin decision, and acquisition certificate. It does not establish that every future trial will agree or that the voltage has another physical name. Those conclusions need separate experimental evidence.

The witness audit is

$\mathcal{W} = (S, R, O, Q, \text{eval}_Q, S_N, N, B, \text{selection policy}, \text{target node}, \text{dependency graph}, \text{reverse slice}, \text{acyclic})$

Missing dependencies, unresolved answers, stale memo entries, partial can-

candidate coverage, and oracle assumptions are reported directly.

The action ledger closes the section. Domains, maps, question syntax, evaluators, candidate order, selection policy, dependency graph, builders, checkers, budgets, versions, hash assumptions, and memo policy are *posited*. Topological-order validity, traversal invariants, cache-verification implications, certificate composition, and the fact that a completed witness satisfies its pulled-back requirement are *demonstrated*. Actual reverse slice, traversal order, visited nodes, frontier, stop reason, reports, answers, cache hits, checker outputs, costs, and fiber membership are *read*. The witness is the covariant action and result satisfying a contravariantly pulled requirement; it is not the requirement itself. Arbitrary global selection, unsupported oracle answers, future physical reproducibility, and claims beyond the visited dependencies remain *reserved*. Requirements pull backward; witnesses run forward only after selection acts.

### 4.3 Two-pass calibration

A two-pass calibration earns a name by surviving an exchange of roles. Let three labels be  $X, Y, Z$ . Pass one returns and verifies two direct values  $x_1, y_1$ , then forms an explicit composite

$$z_\Sigma = \Sigma_1(x_1, y_1).$$

Pass two rotates the roles: it returns and verifies  $y_2, z_2$ , then reconstructs

$$x_\Sigma = \Sigma_2(y_2, z_2).$$

The maps  $\Sigma_1, \Sigma_2$ , their parameters, budgets, stopping rules, scales, and comparison resolution must be registered before either pass is evaluated and must not depend on an outside target.

Three comparisons close the protocol:

$$y_1 \equiv_\varepsilon y_2, \quad z_\Sigma \equiv_\varepsilon z_2, \quad x_\Sigma \equiv_\varepsilon x_1.$$

The first checks the direct overlap. The second checks the pass-one reconstruction against a pass-two direct value. The third checks the pass-two reconstruction against the pass-one direct value and closes the round trip. A shared name is licensed only after all three comparisons pass. Two direct values and one composite are recorded at each pass; the composite is never silently promoted to a direct reading.

Construction runs forward:

$$(x_1, y_1) \xrightarrow{\Sigma_1} z_\Sigma, \quad (y_2, z_2) \xrightarrow{\Sigma_2} x_\Sigma.$$

Requirements read contravariantly. If  $F$  maps the complete pass inputs to the three comparison pairs and  $q_\varepsilon$  asks whether a pair agrees at the declared resolution, then

$$F^* q_\varepsilon = q_\varepsilon \circ F.$$

The comparison question does not construct a report, and the forward maps do not choose their acceptance criterion.

Neither composite map is assumed injective. Distinct direct pairs may occupy one composite fiber:

$$\Sigma_1(x, y) = \Sigma_1(x', y') \not\Rightarrow (x, y) = (x', y').$$

The pass ledger therefore retains direct pairs, composite values, and fiber evidence separately, and records any erasure caused by keeping only a composite. This is a computational handoff between finite records, not a physical or quantum superposition.

**Worked check 1: exact finite role rotation (external numerical design).** Take direct values  $x_1 = 2, y_1 = 3$  and preregister

$$\Sigma_1(x, y) = x + y.$$

Then  $z_\Sigma = 5$ . Pass two directly returns  $y_2 = 3, z_2 = 5$  and uses the preregistered inverse reconstruction

$$\Sigma_2(y, z) = z - y,$$

giving  $x_\Sigma = 2$ . All three comparisons close exactly. This demonstrates the protocol shape; it is not the instrument's measured calibration and carries no physical interpretation.

The actual instrument uses two finite numerical routes over shared calibration data

$$C = 18, \quad \text{target} = 5, \quad d_* = \sqrt{C/\text{target}} = \sqrt{18/5}.$$

The current report directly certifies the two route endpoints and the jar. It does not by itself print every  $X, Y, Z$  role-rotation entry above; the full three-role certificate remains a calibration requirement until those direct/composite records are emitted. The direct  $1 \rightarrow 3$  route evaluates the measured second-variation formula at  $d_*$  and reports

$$d = 137011290548979455469 \quad \text{at scale } 10^{18}.$$

The stepped  $1 \rightarrow 2 \rightarrow 3$  route carries its residue through an intermediate quasi-Newton descent and reports

$$s = 137011290753751157155 \quad \text{at scale } 10^{18}.$$

Its finite count trace is

| count | distance report    | inverse report    |
|-------|--------------------|-------------------|
| 0     | 2/1                | 129.6             |
| 1     | 17/9               | 137.7             |
| 2     | $\approx 1.897309$ | $\approx 137.016$ |
| 3     | $\approx 1.897366$ | $s/10^{18}$ .     |

Mediant/continued-fraction, Newton/BFGS, and residue-carrying secant realizations provide separately declared stepped algorithms. Agreement of their displayed endpoint does not identify their traces or costs.

**Worked check 2: path comparison at declared resolution (instrument computation).** The scaled discrepancy is

$$|d - s| = 204771701686, \quad \frac{|d - s|}{10^{18}} = 2.04771701686 \times 10^{-7}.$$

The scaled integers are not equal. Their decimal reports agree to nine digits, and the discrepancy is below the count-three resolution. Let  $b_3$  be the preregistered finite partition whose owned cell is the count-three jar, and write

$$a \equiv_3 b \iff b_3(a) = b_3(b).$$

This is an equivalence induced by equality of bin labels; for these reports,  $d \equiv_3 s$ . The house slogan  $1 \rightarrow 3 = 1 \rightarrow 2 \rightarrow 3$  abbreviates  $d \equiv_3 s$ , not integer equality and not identity of algorithms.

The one discrepancy  $\Delta = |d - s|/10^{18}$  is not a variance estimate. Statistical variance requires at least  $n \geq 2$  comparable repeated discrepancies  $\Delta_1, \dots, \Delta_n$ :

$$\bar{\Delta} = \frac{1}{n} \sum_{j=1}^n \Delta_j, \quad s_{\Delta}^2 = \frac{1}{n-1} \sum_{j=1}^n (\Delta_j - \bar{\Delta})^2.$$

Inference from  $s_{\Delta}^2$  additionally requires a declared sampling model, stability assumptions, and treatment of dependence. The direct and stepped routes share calibration inputs and are not automatically independent. For random reports  $D, S$ ,

$$\text{Var}(D - S) = \text{Var}(D) + \text{Var}(S) - 2 \text{Cov}(D, S).$$

This probabilistic covariance is distinct from categorical covariance of forward maps. No statistical variance, covariance, or uncertainty interval has been computed by the current two-route comparison.

The owned result is the ordered-wall jar written

$$[129.6, 137.7].$$

The word open is the repository's metaphor for unresolved content, not topological interval notation. Both finer reports near 137.011290 fall between the ordered walls, but the instrument hands over the jar rather than either point.

**Worked check 3: wall and ownership audit (instrument computation).** The lower and upper walls come from the first two finite stepped reports,  $2/1 \mapsto 129.6$  and  $17/9 \mapsto 137.7$ . Their order is checked, the count-three stop is recorded, and both direct and stepped endpoint reports are compared against the walls. These checks establish the owned bracket and path agreement at its resolution. They do not establish agreement with an external physical constant.

Richardson extrapolation remains conditional. The direct and stepped values arise from different algorithms; no common family  $A(h)$ , refinement ratio  $r$ , or error order  $p$  is established. Under the counterfactual assumptions  $D = A(h)$ ,  $S = A(h/2)$ , one could write

$$A_R = S + \frac{S - D}{2^p - 1}.$$

Those sensitivity values depend on the unestablished order and orientation. Licensing the method requires one implemented ordered  $h$ -family, at least three refinement levels, observed-order estimation, and stability under further refinement. A machine/configuration/cost index is one proposed candidate for  $h$ , but it requires an implemented ordering and an error model. Greater budget or cost is not automatically finer resolution. The current direct and stepped pair is therefore invalid for Richardson extrapolation. Until a qualifying family exists, no extrapolate is a measurement, uncertainty bound, or owned result.

**Worked check 4: target-dependent transition rejected (adversarial test).** Suppose a proposed composite map is adjusted after seeing a desired number  $t$ , for example

$$\tilde{\Sigma}_1(x, y) = t.$$

It will reproduce  $t$  regardless of the direct readings. The numerical comparison may pass, but preregistration and target-independence fail, so calibration is rejected. Agreement manufactured by the transition is not path independence.

Each route and pass has an independent budget. A pass certificate records direct values and their verification, composite values, map definitions and versions, inputs, scales, update traces, stopping rules, budgets spent, residuals, comparison decisions, erasures, and provenance. Stops are COMPLETE, BUDGET, MISSING-DIRECT, COMPARISON-UNRESOLVED, OVERLAP-FAILURE, RECONSTRUCTION-FAILURE, CLOSURE-FAILURE, VERSION-MISMATCH, and DOMAIN-ERROR. Any noncomplete stop returns the unresolved frontier; it does not license the shared name.



The calibration audit is

$\mathcal{C} = (X, Y, Z, \text{role assignments, overlap identities, } \Sigma_1, \Sigma_2, \text{versions, preregistration, target-independence, pass-one direct/composite records, pass-two direct/composite records, overlap, reconstructions, } d, s, \text{scale, raw discrepancy, resolution relation, bin-map version, discrepancy-bin version, jar wall, conditional statistical assumptions/results, budgets, traces, stops, unresolved frontier, erasures, ...})$

The action ledger is exact. Calibration data, role rotation, composite maps, algorithms, budgets, resolution relation, versions, bin map, statistical assumptions, and stop policies are *posited*. Arithmetic identities, wall-order checks, join/fiber implications, resolution-equivalence laws, and round-trip acceptance implications are *demonstrated*. Actual direct and stepped outputs, raw discrepancy, residual and comparison decisions, traces, budgets spent, statistical outputs when any, owned walls, unresolved frontier, and stops are *read*. The name  $\alpha_c$  is licensed for the path-independent computational result at declared resolution. Equality of methods, integer reports, costs, physical provenance, an external empirical constant, and conditional extrapolates remain *reserved*. The honest handed-over result is the jar.

# Chapter 5

## The Slip Becomes A Number

### 5.1 A crossing predicate

A crossing predicate is a finite question, not yet a root or a physical event. Let the positive rational domain be

$$\mathbb{Q}_{>0}^{\text{rec}} = \{(p, q) : p, q \in \mathbb{N}, p > 0, q > 0\}, \quad d = p/q.$$

Suppose a report is represented by integers  $A(p, q) \geq 0$  and  $B(p, q) > 0$ , with reported ratio  $A/B$ , and let target  $t \geq 0$  be an integer in compatible declared units. The exact crossing question is

$$C_t(p, q) := [t B(p, q) \leq A(p, q)].$$

It uses integer cross-multiplication. No decimal approximation or floor is required to decide the branch.

A forward evaluator constructs

$$E_t(p, q) = (A(p, q), B(p, q), tB(p, q), A(p, q)).$$

The requirement question on this report asks whether the third entry is at

most the fourth. Reading pulls that question back through  $E_t$ ; the evaluator is covariant and the requirement read is contravariant. The question does not construct a distance or choose a side of a bracket.

For the instrument's finite crossing model,

$$g(d) = \frac{18}{d^2}, \quad t = 5.$$

At  $d = p/q > 0$ ,

$$C_5(p, q) := [5p^2 \leq 18q^2].$$

The constants 18 and 5, the positive domain, and the interpretation of  $d$  are posited calibration data. If units are attached,  $18/d^2$  and 5 must have matching units before comparison.

The monotone orientation follows directly. For  $0 < d_1 \leq d_2$ ,

$$d_1^2 \leq d_2^2, \quad \frac{1}{d_1^2} \geq \frac{1}{d_2^2}, \quad \frac{18}{d_1^2} \geq \frac{18}{d_2^2}.$$

Therefore the yes region for  $C_5$  lies downward in distance. The floored report  $[18/d^2]$  is also nonincreasing, but has plateaus on which distinct ratios receive the same integer.

A bracket claim needs additional premises. On an ordered domain  $D$ , require:

$$a < b, \quad C_t(a) = \text{yes}, \quad C_t(b) = \text{no},$$

and monotonicity in the declared orientation,

$$d_1 < d_2 \implies g(d_1) \geq g(d_2).$$

These facts certify that the bracket straddles a truth transition. They do not automatically prove that an exact rational equality  $g(d) = t$  exists inside; the boundary may lie outside the represented rational set.

A finite scan makes the method auditable. Given ordered candidates  $D_N = [d_1, \dots, d_N]$  and budget  $B \leq N$ , visit  $i = 1, \dots, B$ . Verify the domain conditions, construct the cross-products, resolve  $C_t(d_i)$ , and append the full row to crossed, not-crossed, or domain-error slices. Stop with COMPLETE, BUDGET, DOMAIN-ERROR, or RESOLVER-DOMAIN. The certificate contains candidate indices and rationals, unreduced and reduced forms when used, cross-products, decisions, slices, units, budget, completeness, erasures, and stop reason. An unvisited suffix carries no decision.

**Worked example 1: exact endpoint crossing (instrument computation).** At  $d = 1 = 1/1$ ,

$$5 \cdot 1^2 = 5 \leq 18 \cdot 1^2 = 18,$$

so the predicate answers yes. At  $d = 2 = 2/1$ ,

$$5 \cdot 2^2 = 20 > 18,$$

so it answers no. The ordered endpoints straddle the transition. At  $d = 3/2$ ,

$$5 \cdot 3^2 = 45 \leq 18 \cdot 2^2 = 72,$$

so the midpoint candidate remains on the crossed side.

**Worked example 2: a finite rational slice (instrument computation).** Visit  $d = 9/5, 19/10, 2$ . The exact comparisons are

| $d$   | $5p^2$ | $18q^2$ |
|-------|--------|---------|
| 9/5   | 405    | 450     |
| 19/10 | 1805   | 1800    |
| 2/1   | 20     | 18      |

and the decision slice is

$$\text{crossed} = [9/5], \quad \text{not-crossed} = [19/10, 2].$$

Thus the sampled transition lies between  $9/5$  and  $19/10$ . This is a statement about the sampled exact predicate, not yet a certified narrowest bracket.

A scan certifies only the visited rows and, when complete, the transitions between adjacent listed candidates. It does not produce the canonical seed-free continued-fraction jar. The canonical distance walls begin with  $17/9$  and  $2$ :

$$17/9 < 2, \quad 5 \cdot 17^2 = 1445 \leq 1458 = 18 \cdot 9^2, \quad 5 \cdot 2^2 = 20 > 18.$$

Thus  $17/9$  is on the yes side and  $2$  on the no side. Their later continued-fraction role may be developed in the jar section; the exact opposed-wall arithmetic already belongs in this certificate.

Monotonicity cannot be inferred from one change of branch. It must be established from the report formula or tested only as a finite empirical pattern.

**Worked example 3: nonmonotone counterexample (external numerical experiment).** Let

$$h(d) = (d - 1)(d - 2), \quad C(d) = [h(d) \geq 0].$$

The predicate is yes at  $d = 1/2$ , no at  $d = 3/2$ , and yes again at  $d = 3$ . A scan sees two transitions. Keeping one half after a branch test without a monotonicity or sign-change argument could discard another crossing. The shared word crossing does not license a bisection invariant.

**Worked example 4: malformed rational record (computation).** The record  $(p, q) = (3, 0)$  has no represented positive rational value. Cross-

multiplication might still produce integers, but the domain certificate fails first. The correct result is DOMAIN-ERROR, not a crossing decision. Likewise, a negative or unit-incompatible target requires a different declared domain and question.

A compact rational bisection contract is also available. Validate positive ordered endpoints

$$\ell < u, \quad C_t(\ell) = \text{yes}, \quad C_t(u) = \text{no}.$$

Form the exact rational midpoint  $m = (\ell + u)/2$ . If  $C_t(m)$  is yes, replace  $\ell$  by  $m$ ; otherwise replace  $u$  by  $m$ . Each accepted step preserves endpoint order and the opposed-bit invariant. With finite fuel  $N$ , the certificate returns the final rational bracket, every midpoint and branch decision, unused fuel, and stop reason. It certifies a truth-transition bracket, not an exact represented root. This legacy fuel-bisection route is distinct from the canonical seed-free continued-fraction jar; agreement between their walls requires a separate comparison.

The floor report is related but distinct. For nonnegative  $A$  and positive  $B$ ,

$$\left\lfloor \frac{A}{B} \right\rfloor \geq t \iff tB \leq A.$$

This identity explains why the exact crossing question can be evaluated without performing division. The certificate must retain  $A, B, tB$ , and the comparison result; storing only the Boolean branch erases the distance from threshold. Equal floors do not imply equal ratios or equal distances. The fractional residue

$$A - B \lfloor A/B \rfloor$$

records information erased by the floor. Also, unreduced pairs such as  $(p, q)$  and  $(2p, 2q)$  represent the same rational distance while remaining different records; canonicalization by a positive greatest common divisor must be de-

clared when record equality is intended to mean rational equality.

**Worked example 5: static and kinetic slip (physics, then interpretation).** A friction experiment may define a physical breakaway event through applied force, normal load, surface preparation, velocity, temperature, and an uncertainty model. Physics owns those variables and the distinction between static and kinetic regimes. Reading  $C_t(d)$  as that event is an interpretation requiring an apparatus map from rational distance records to measured forces. The finite inequality alone supplies no friction law, discontinuity, or unpredictability claim.

The crossing audit is

$\mathcal{X} = (D, t, \text{units}, A, B, C_t, \text{monotonicity premise}, \text{endpoint orientation/bits}, \text{raw } A/B, \text{floor residue},$

An absent monotonicity argument, failed domain check, incompatible unit, partial scan, or unresolved arithmetic operation is recorded directly. The audit must state whether integers are unbounded; for bounded integers it must instead certify that every square and cross-product avoids overflow.

The action ledger closes the section. Domain, target, units, report formula, candidate order, canonicalization, arithmetic model, route, budgets, and stop policies are *posited*. Monotonicity, cross-product identities, endpoint inequalities, bisection invariants, finite slice coverage, and implications proved from declared premises are *demonstrated*. Actual candidate reports, floors, residues, midpoint and branch decisions, slices, budgets spent, frontiers, and stops are *read*. Existence of an exact represented root, uniqueness, physical breakaway, and claims beyond the scanned prefix remain *reserved*. A crossing predicate supplies one finite question; bracketing requires the premises that let its answers be sliced safely.

## 5.2 Bisection is iteration by slicing

A bisection bracket is finite data:

$$\mathcal{B} = (\ell, u, c_\ell, c_u), \quad \ell, u \in \mathbb{Q}_{>0}^{\text{rec}},$$

where  $c_\ell = C_t(\ell)$  and  $c_u = C_t(u)$ . In the downward-yes orientation, a valid input satisfies

$$\ell < u, \quad c_\ell = \text{yes}, \quad c_u = \text{no}.$$

The bracket does not contain a point value or a guarantee that a represented equality exists.

For  $\ell = p/q$  and  $u = r/s$ , with positive denominators, the exact midpoint is

$$m = \frac{ps + rq}{2qs}.$$

A one-step slice evaluates  $c_m = C_t(m)$  and returns

$$\text{bisectOnce}(\mathcal{B}) = \begin{cases} ((\ell, u, m, c_m), (m, u, \text{yes}, \text{no})), & c_m = \text{yes}, \\ ((\ell, u, m, c_m), (\ell, m, \text{yes}, \text{no})), & c_m = \text{no}. \end{cases}$$

The first component is the step ledger; the second is the retained half. Exact arithmetic may leave the midpoint unreduced, so a canonicalization policy must be declared separately.

The pure loop is fuel-only and recursive:

$$\begin{aligned} \text{run}(0, \mathcal{B}) &= (\mathcal{B}, []), \\ \text{run}(n+1, \mathcal{B}) &= \text{let } (\sigma, \mathcal{B}') = \text{bisectOnce}(\mathcal{B}), \\ &\quad \text{let } (\mathcal{B}'', T) = \text{run}(n, \mathcal{B}') \text{ in } (\mathcal{B}'', \sigma :: T). \end{aligned}$$

Thus zero fuel returns the current bracket and an empty trace. Successor fuel performs one step, recurses on the updated bracket with the remaining fuel,



and prepends the current step to the remaining trace, preserving chronological order. A successful run with fuel  $N$  has trace length  $N$ . Width-target stopping is a proposed extension, not part of this pure recurrence.

The certificate contains initial and final brackets, every midpoint, raw cross-products, branch decisions, canonical forms when used, fuel and trace length, widths, erasures, and stop reason. A fuel-complete result owns a finite bracket, not an unread infinite tail.

A replay checker validates each adjacent row: recompute the exact midpoint; recompute cross-products and crossing bit; verify that the selected branch produces the recorded next bracket; verify that the next row's old bracket equals that result; recheck opposed endpoint bits, order, and width; and finally check fuel accounting and trace length. Replay success certifies adjacency and invariant preservation for the recorded finite run.

The familiar width law is conditional. Interpret the positive rationals in an ordered metric line, use the exact arithmetic midpoint, preserve endpoint order and opposed crossing bits, and assume the monotone enclosure orientation from the previous section. Then

$$w_k = u_k - \ell_k, \quad w_{k+1} = \frac{w_k}{2}, \quad w_n = w_0 2^{-n}.$$

This follows by induction from retaining one of two equal halves. The formula proves shrinking width under those hypotheses. It does not prove an exact root, uniqueness without stricter premises, or physical correctness.

**Worked example 1: four exact slices (instrument computation).**

For  $C_5(p, q) = [5p^2 \leq 18q^2]$ , begin with  $[1, 2]$ . The exact trace is

| $k$ | $[\ell_k, u_k]$ | $m_k$   | $C_5(m_k)$ |
|-----|-----------------|---------|------------|
| 0   | $[1, 2]$        | $3/2$   | yes        |
| 1   | $[3/2, 2]$      | $7/4$   | yes        |
| 2   | $[7/4, 2]$      | $15/8$  | yes        |
| 3   | $[15/8, 2]$     | $31/16$ | no         |

and the retained bracket is

$$[15/8, 31/16].$$

Its width is  $1/16 = (2 - 1)2^{-4}$ . The certificate owns four cross-multiplication decisions and this bracket; it does not own the irrational equality boundary  $\sqrt{18/5}$  as a represented rational point.

**Worked example 2: malformed enclosure (computation).** Take  $[1, 3/2]$ .

Both endpoints answer yes. A midpoint can still be computed, but there are no opposed endpoint bits, so the bisection contract stops with INVARIANT-FAILURE. Repeated halving of an arbitrary interval is not a crossing enclosure.

**Worked example 3: opposed bits without monotonicity (external counterexample).** Let

$$h(d) = -(d - 1)(d - 2)(d - 3), \quad C(d) = [h(d) \geq 0].$$

On  $[1/2, 7/2]$ , the lower endpoint is yes and the upper endpoint no, but three crossings lie inside. The predicate is not monotone. Binary slicing may retain a half containing one transition, yet it cannot claim a unique boundary or safely apply the downward-yes invariant globally. Opposed bits alone are

insufficient.

Endpoint projection is lossy. Define  $\pi_{\text{final}}$  to retain only the final  $(\ell_N, u_N)$ . Distinct step histories can share that pair while differing in raw midpoint records, reductions, costs, or repeated checks. A question on final endpoints cannot recover those histories.

**Worked example 4: same bracket, different ledger (computation).**

One run may reduce every midpoint immediately; another may retain unreduced pairs and canonicalize only the final endpoints. If both branch decisions agree, their rational final brackets may be equal while their stored traces and costs differ. The audit must retain or explicitly erase canonicalization history; endpoint equality does not identify the algorithms.

Newton’s method is a different route:

$$x_{k+1} = x_k - \frac{f(x_k)}{f'(x_k)}.$$

It requires a differentiable model, nonzero derivative at each used point, a starting value, and its own stopping argument. It need not preserve an enclosure. Quasi-Newton methods replace exact derivative information with updated approximations and require still different premises. Their speed or endpoint agreement does not transfer bisection’s opposed-bit invariant.

**Worked example 5: Newton step leaves an interval (external counterexample).** For  $f(x) = x^3 - 2x + 2$ , start at  $x_0 = 0$ . Then

$$x_1 = 0 - \frac{2}{-2} = 1, \quad x_2 = 1 - \frac{1}{1} = 0,$$

so the route cycles between 0 and 1. This does not make Newton invalid; it shows that tangent iteration and bisection slicing have different guarantees and failure certificates.

The canonical jar route is different again. It uses the seed-free continued-

fraction algorithm: Euclidean division-with-remainder produces whole partial quotients, and their convergents produce the ordered distance walls  $17/9$  and  $2/1$ . It is not midpoint bisection, a raw mediant crawl, or a grid search. Those walls map to the canonical count-three report jar written [129.6, 137.7]. The rational distance bisection above is a legacy fuel-driven route. Neither route inherits the other's certificate; wall agreement needs an explicit comparison map.

The bisection audit is

$$\mathcal{B}_{\text{audit}} = (C_t, \text{domain}, \text{orientation}, \mathcal{B}_0, \text{monotonicity/enclosure evidence}, N, \text{proposed width-target extension})$$

The action ledger is exact. Domain, crossing question, initial bracket, orientation, arithmetic model, canonicalization, fuel, and any proposed width-target extension are *posited*. Midpoint identities, recursive trace-length law, replay-checker implications, invariant preservation, conditional width law, and final-bracket implications are *demonstrated*. Actual midpoints, cross-products, branch decisions, bracket trace, fuel use, widths, and stop are *read*. Exact root existence, uniqueness, Newton or quasi-Newton guarantees, continued-fraction wall identity, and physical slip meaning remain *reserved*. Bisection is iteration by slicing because every update keeps one evidenced half and records the one it discarded.

### 5.3 A report is not a certificate

A numerical routine may return data without establishing every claim a reader wants to attach to it. Separate four layers:

```

return ::= none(reason) | some(payload),
report  ::= report(schema, route, inputs, outputs, trace, stop),
validation ::= accepted | rejected(failures) | unresolved,
certificate ::= certificate(report, checks, claim, provenance).
```

An optional return distinguishes absence from a numerical zero or empty bracket. A report organizes returned data. Validation checks a declared schema and claim. Only an accepted validation can be packaged as a certificate for that specific claim.

The conversion chain is explicitly one-way:

raw bytes or text  $\longrightarrow$  parsed report  $\longrightarrow$  validated report  $\longrightarrow$  certified claim.

Raw input does not imply successful parsing; parsing does not imply validity; validity does not imply every possible claim; and certification of one claim does not certify another.

If a Boolean validator  $V$  is used for property  $P$ , acceptance licenses  $P(r)$  only when the soundness implication

$$V(r) = \text{true} \implies P(r)$$

has been demonstrated or is explicitly trusted or posited. The actual Boolean outcome is *read*. A validator name or true bit cannot supply its own soundness argument.

The reporting pipeline is finite:

1. acquire the optional return within budget,
2. parse schema and version,
3. validate domains, units, scales, and route identity,
4. replay transitions and arithmetic,
5. evaluate the requested claim and slice checks into pass/fail/unresolved,
6. emit a certificate only when every required check passes.

Stops are NONE, BUDGET, PARSE-ERROR, VERSION-MISMATCH, SCHEMA-UNKNOWN, REQUIRED-ENTRY-MISSING, UNIT-SCALE-ERROR, DOMAIN-ERROR, ARITHMETIC-RESOURCE, OVERFLOW, PARTIAL-REPORT, ROUTE-MISMATCH,

REPLAY-FAILURE, CERTIFICATE-MISSING, CERTIFICATE-INVALID, CLAIM-FAILURE, and COMPLETE. The receipt records raw return, parsing decisions, all validation rows, budget spent, checked scope, unresolved frontier, erasures, stop reason, and provenance. A failed or partial validation does not erase the report; it limits what the report licenses.

For a bracket report, validation questions include

$$\ell < u, \quad C_t(\ell) = \text{yes}, \quad C_t(u) = \text{no},$$

together with domain and arithmetic assumptions. If a trace is supplied, replay checks every midpoint, branch, adjacent bracket, fuel count, and final endpoint. If no trace is supplied, endpoint checks may still pass, but no claim about how the endpoints were reached is certified.

**Worked example 1: a replayable bisection report (computation).**

A four-step report contains

$$[1, 2] \rightarrow [3/2, 2] \rightarrow [7/4, 2] \rightarrow [15/8, 2] \rightarrow [15/8, 31/16].$$

The validator recomputes midpoints  $3/2, 7/4, 15/8, 31/16$ , checks their crossing bits, verifies every retained half, and confirms trace length four and final width  $1/16$ . If all rows pass, the certificate owns this finite bracket and trace. It does not own an exact represented crossing point.

**Worked example 2: a malformed bracket report (computation).**

Suppose a payload reports

$$\ell = 2, \quad u = 1.$$

Both values parse as positive rationals, but the order check fails. The result is a valid parsed report and a rejected bracket certificate. Swapping the endpoints during validation would silently change the returned data; a repair

must be a new, explicitly transformed report with its own provenance.

**Worked example 3: absence is not a default value (computation).**

If acquisition returns `none(BUDGET)`, the parser must not synthesize  $[0, 0]$ , a zero endpoint, or an empty trace and call it a reading. The receipt owns only the budget stop and unresolved claim. Such a `none` result is scoped to the scanned candidates, budget, and run configuration; it is not a global nonexistence statement. A later run may produce data, but it is a separate acquisition.

A summary projection also loses evidence. Let

$$\pi_{\text{endpoints}}(\text{full report}) = (\ell, u).$$

Distinct traces may share endpoints. Questions about endpoint order can be pulled through this projection, but questions about midpoint history, cost, or route cannot be answered after those entries are erased. Equal summaries do not identify full reports.

**Worked example 4: jar ownership versus endpoint diagnostics (instrument computation).** The canonical count-three report owns the ordered-wall jar

$$[129.6, 137.7].$$

Two finer route diagnostics report, at scale  $10^{18}$ ,

$$d = 137011290548979455469, \quad s = 137011290753751157155,$$

with scaled difference 204771701686. Those diagnostics support a declared-resolution comparison, but neither replaces the jar. A valid report tags the continued-fraction jar route, direct route, and stepped route separately; shared digits do not merge their traces or costs.

The integers are unequal, yet coarse formatting maps both to the dis-

played string 137.011. This is a display collision, not equality of readings. A loss-auditable decimal or scaled-floor report retains the raw integer or exact rational, scale, formatting rule, displayed quotient, and discarded residue. For exact ratio  $A/B$  at scale  $M$ ,

$$n = \left\lfloor \frac{AM}{B} \right\rfloor, \quad \rho = AM - nB.$$

The floor erases  $\rho/B$ ; retaining  $A, B, M, n, \rho$  permits replay.

**Worked example 4b: flags are not a certificate (computation).**

A payload containing entries `verified = true` and `complete = true` may parse under a permissive schema. If the required certificate is absent, validation returns CERTIFICATE-MISSING; if its checks fail, it returns CERTIFICATE-INVALID. Self-reported flags are data, not validation evidence.

**Worked example 5: same payload, different route (computation).**

Suppose two routines return the rational pair  $(3, 2)$ . One obtained it by reducing  $18/12$ ; the other read it directly from an input table. The endpoint payloads agree, but the first certificate may own a Euclidean remainder trace while the second owns a table lookup. A claim about value equality can pass; a claim about equal process or cost must fail or remain unresolved.

Physical reporting requires another provenance chain. A displayed voltage interval is not automatically a calibrated measurement interval. Physics owns sensor response, units, calibration transfer, environmental conditions, sampling plan, and uncertainty model. Numerical validation can certify formatting, arithmetic, and finite acquisition records; it cannot manufacture empirical coverage or physical interpretation.

A reproducibility receipt records the checker and compiler releases, dependency lock and versions, source-snapshot hash, run and build configuration, arithmetic model, platform, protocol version, run identifier and time, deterministic inputs, and every declared nondeterministic input—including seeds,



schedules, and environmental state where relevant. Reproduction is conditional on matching that declared provenance. Agreement under one configuration establishes neither physical reproducibility nor invariance across platforms, arithmetic models, compilers, or checker releases.

The report audit is

$\mathcal{R}$  = (raw bytes/text and optional return, encoding, schema/version, parser, validator/checker identity and release, soundness status, compiler release, dependency lock/version, source-snapshot hash, run/build configuration, arithmetic model, platform, protocol, run ID/timestamp, deterministic and declared nondeterministic inputs, route identity, units/scale, exact raw values, formatting rule, quotient, and residue, outputs, trace, budget, checked scope, unresolved frontiers, domain/arithmetic/replay checks, failures, claim syntax/evaluator, pass/fail/unresolved slices, summary projections, erasures, certificate presence and validity, certified claims, provenance).

The action ledger is exact. Schemas, parsers, validators, checkers, versions, routes, claim syntax, evaluators, budgets, required checks, toolchain choices, and run configuration are *posited*; validator soundness is also posited whenever it has not been proved. Parsing and validation invariants, replay implications, arithmetic identities, certificate composition rules, and any actually established soundness implication are *demonstrated*. Raw returns, parse and Boolean outcomes, reports, validation rows, exact values, formatting collisions, floor residues, route diagnostics, budgets spent, stops, and owned jar walls are *read*. Exact point ownership, process identity from equal payloads, claims outside the certificate, cross-toolchain invariance, physical reproducibility, empirical uncertainty coverage, and deterministic equivalence while nondeterministic inputs remain uncontrolled are *reserved*. A report may be useful when validation is partial; it becomes a certificate only for the claims its checks actually support.

# Chapter 6

## The Bracketed Number

### 6.1 Dyadic data structure

A dyadically refined bracket is finite data, not an approximate spelling of an unknown point. Write its primary representation as

$$D = (L, U, Q), \quad [D] = [L/Q, U/Q],$$

where  $L, U \in \mathbb{Z}$ ,  $Q \in \mathbb{N}_{>0}$ , and  $L < U$ . The integers are stored exactly; the interpretation map supplies the scale. A crossing question  $C$  orients the endpoints:

$$C(L/Q) = \text{yes}, \quad C(U/Q) = \text{no}.$$

Order and opposed endpoint answers are separate invariants. Neither follows merely from the other, and neither says that a unique point lies inside the bracket. If a run starts at explicit denominator  $Q_0$ , an optional rung index  $k$  derives  $Q_k = 2^k Q_0$ . Unless  $Q_0$  itself is a power of two, the represented rationals need not be dyadic rationals in lowest terms; it is the refinement operation, not the initial rational grid, that is dyadic.

Dyadic refinement is a conversion followed by a slice. Convert both end-

points to the next common denominator,

$$(L, U, Q) \mapsto (2L, 2U, 2Q), \quad M = L + U,$$

so that  $M/(2Q)$  is the exact midpoint. Reading  $C(M/(2Q))$  chooses one half:

$$(L, U, Q) \mapsto \begin{cases} (M, 2U, 2Q), & C(M/(2Q)) = \text{yes}, \\ (2L, M, 2Q), & C(M/(2Q)) = \text{no}. \end{cases}$$

The retained interval has width  $(U - L)/(2Q)$ . One rung doubles the denominator and halves the interpreted width. The conversion is exact; only the predicate answer selects a slice.

Representation and value must not be confused. The triples  $(1, 3, 2)$  and  $(2, 6, 4)$  denote the same interval. A canonicalization divides  $L, U, Q$  by their positive greatest common divisor. Canonicalization preserves the interpreted interval but can erase the relation  $Q_k = 2^k Q_0$ , hence a refinement-history clue. It belongs either after replay or beside an unreduced trace, never in place of that trace.

An adjacent exact grid pair  $a < b$  with common positive denominator  $Q$ ,

$$a = L/Q, \quad b = U/Q, \quad U - L = 1,$$

converts to  $D = (L, U, Q)$ . The conversion certifies common denominator, adjacency, and order. It becomes an oriented crossing bracket only after the separate evidence  $C(L/Q) = \text{yes}$  and  $C(U/Q) = \text{no}$  is read. Grid conversion alone is not crossing evidence.

The numerical method is a finite state transition. Given initial data  $D_0$ , rung budget  $N$ , and arithmetic budget  $A$ , it produces

$$D_0 \xrightarrow{b_0} D_1 \xrightarrow{b_1} \cdots \xrightarrow{b_{r-1}} D_r, \quad r \leq N, \quad b_i \in \{\text{yes}, \text{no}\}.$$

Stops include COMPLETE, RUNG-BUDGET, ARITHMETIC-BUDGET, DOMAIN-

ERROR, ENDPOINT-NOT-OPPOSED, ORDER-FAILURE, and PREDICATE-UNRESOLVED.

A budget stop owns the  $r$  checked rungs and unresolved next transition; it does not assert that further refinement is impossible.

A replay certificate contains the initial triple, scale convention, predicate identity, arithmetic model, budgets, the explicit  $Q_0$ , optional rung indices, every midpoint numerator, every read branch bit, every retained triple, and the stop. Its checker verifies  $Q' = 2Q$ , midpoint addition, half selection, order, opposed endpoint answers, fuel use, numerator gap, and the final laws

$$Q_r = 2^r Q_0, \quad U_r - L_r = U_0 - L_0, \quad \text{width}(D_r) = \frac{U_0 - L_0}{2^r Q_0}.$$

Checks are sliced into pass, fail, and unresolved rows. Passing arithmetic replay does not certify predicate soundness; passing endpoint checks without a trace does not certify the route.

**Worked example 1: two exact rungs (computation).** Let  $D_0 = (1, 2, 1)$ , so  $Q_0 = 1$ , with endpoint answers yes and no. The optional rung index is  $k = 0$ . The first midpoint is  $3/2$ . If its answer is yes, retain  $D_1 = (3, 4, 2)$ , denoting  $[3/2, 2]$ , with  $k = 1$ . Its next midpoint is  $7/4$ . If that answer is no, retain  $D_2 = (6, 7, 4)$ , denoting  $[3/2, 7/4]$ , with  $k = 2$ . The widths  $1, 1/2, 1/4$  and both slice choices replay without rounding.

**Worked example 2: denominator growth is not value growth (computation).** The conversion from  $(3, 4, 2)$  to  $(6, 8, 4)$  changes neither endpoint. Reading the larger integers as unscaled values would invent  $[6, 8]$ . Reading only  $Q = 4$  while discarding the numerators would lose the bracket. The triple is the data structure; none of its fields is a reading by itself.

**Counterexample 3: equal endpoints do not certify equal histories (computation).** One run may reach  $(6, 7, 4)$  by the two rungs above. Another may receive it as initial data. Their endpoint projections agree, but

their transition traces and costs do not. The projection

$$\pi(L, U, Q, \text{trace}) = (L, U, Q)$$

supports interval equality while erasing route, branch sequence, and budget use. Equal projections cannot be pulled back into equal processes.

**Counterexample 4: an unordered payload is not repaired by interpretation (computation).** The triple  $(5, 3, 4)$  fails  $L < U$ . Sorting it during validation would create a different report and could reverse the required endpoint answers. The original is rejected; an explicitly transformed triple requires new provenance and fresh endpoint checks.

This dyadic route is legacy refinement data. It must not be renamed as, or substituted for, the canonical seed-free continued-fraction route. That route performs Euclidean division with remainder, reads whole partial quotients, and constructs ordered walls; for the owned count-three result its walls  $17/9$  and  $2$  induce the jar  $[129.6, 137.7]$ . It is not denominator doubling, midpoint selection, or adjacent-grid conversion. A dyadic endpoint may numerically equal a wall, but equality of a rational value does not identify algorithms, traces, budgets, or ownership claims. The continued-fraction jar remains canonical; a dyadic run is a separately labelled diagnostic.

Governance fixes the initial triple, explicit  $Q_0$ , optional rung-index convention, predicate question, orientation, scale convention, canonicalization policy, budgets, stop vocabulary, checker release, arithmetic model, and protocol version before a run. Provenance records those choices together with run identity, time, platform, deterministic inputs, declared nondeterministic inputs, raw branch answers, and every loss-causing projection. Changing a governed item starts a new run rather than silently extending the old certificate.

The action ledger is exact. Initial data, predicate, orientation, budgets, arithmetic model, canonicalization policy, and certificate grammar are

*posited*. Midpoint identities, the laws  $Q' = 2Q$  and  $Q_r = 2^r Q_0$ , numerator-gap and width laws, invariant-preservation implications, canonicalization equality, and replay-checker implications are *demonstrated*. Endpoint answers, branch bits, visited triples, budget use, stops, and checker outcomes are *read*. A unique enclosed point, predicate soundness not separately established, route identity from equal endpoints, continuation beyond a stop, and ownership of the canonical continued-fraction jar remain *reserved*.

## 6.2 Resolution as an explicit parameter

Resolution is declared data, not a universal constant inferred from finite output. Four controls must remain distinct:

$$\theta = (\mathcal{X}, \mathcal{A}, q, s, f; G, N, M, K),$$

where  $\mathcal{X}$  is the input domain,  $\mathcal{A} = \{a_0, \dots, a_{m-1}\}$  a finite alphabet,  $q : \mathcal{X} \rightarrow \mathcal{A}$  a quantizer,  $s > 0$  its scale, and  $f \geq 0$  a floor. The grid parameter  $G$  selects a finite exact grid;  $N$  bounds dyadic refinement rungs;  $M > 0$  is the integer display scale; and  $K$  counts whole continued-fraction partial quotients. Each affects a different route. The value  $G = 0$  is a sentinel meaning that no ordinary grid has been selected; it is not a grid of zero width and cannot enter division. Its cells are  $B_j = q^{-1}(a_j)$ . Boundary conventions are required: a value on a cell wall cannot be assigned by an unstated tie rule.

For a uniform scalar quantizer with origin  $o$ , width  $\Delta > 0$ , and code range  $\{j_{\min}, \dots, j_{\max}\}$ , one grammar is

$$q(x) = j \iff o + j\Delta \leq x < o + (j+1)\Delta,$$

with separate out-of-range stops. Numerical resolution is  $\Delta$ , while  $s$  converts code units to reported units. A floor may stop refinement when  $w_{r+1} < f$ . This is a stopping policy, not evidence that nature contains no smaller

distinction.

Rungs count transitions; cells partition a domain; scale interprets codes; the floor bounds admitted refinement; and the quantizer chooses a cell. Grid selection changes available exact cells,  $N$  changes dyadic depth,  $M$  changes display quotient and residue, and  $K$  changes continued-fraction walls. Collapsing these parameters hides which operation lost information and falsely merges route ownership.

Parameter-relative equality is the equivalence relation

$$x \equiv_{\theta} y \iff q_{\theta}(x) = q_{\theta}(y).$$

Reflexivity, symmetry, and transitivity follow from equality of labels; equality of underlying values does not. A specification  $\theta'$  refines  $\theta$  when a map  $c : \mathcal{A}_{\theta'} \rightarrow \mathcal{A}_{\theta}$  satisfies  $q_{\theta} = c \circ q_{\theta'}$ , equivalently every fine cell lies inside its indicated coarse cell. Refinement can separate a former equivalence class but cannot justify comparisons if route, scale, or boundary conventions are incompatible.

Given rung budget  $N$ , predicate budget  $P$ , and arithmetic budget  $A$ , computation returns a finite trace and stops with COMPLETE, RUNG-BUDGET, PREDICATE-BUDGET, ARITHMETIC-BUDGET, BELOW-FLOOR, OUT-OF-RANGE, BOUNDARY-UNRESOLVED, or INVALID-PARAMETERS. A complete stop vocabulary also includes CELLS-EXHAUSTED, VERSION-MISMATCH, OVERFLOW, and UNRESOLVED-FRONTIER. A certificate records  $\theta$ , budgets, raw inputs, boundary decisions, transitions, resources spent, final cell or bracket, and stop. Replay checks parameter domains, membership, scale conversion, transitions, resource counts, and the stop inequality. Checks are sliced into pass, fail, and unresolved rows; a budget stop certifies completed work, not global indistinguishability.

Quantization is a many-to-one projection. If  $x, y \in B_j$ , then  $q(x) = q(y)$ , but not necessarily  $x = y$ . Retaining only  $a_j$  erases within-cell order and distance as well as quantizer version, scale, boundaries, budgets, and stop.

Equal summary symbols cannot identify equal readings across undeclared configurations.

**Worked example 1: equivalence changes under refinement (computation).** Let display quantization use  $q_M(x) = \lfloor Mx \rfloor$ , and take  $x = 1.2341$ ,  $y = 1.2358$ . Then  $q_{100}(x) = q_{100}(y) = 123$ , so  $x \equiv_{\theta_{100}} y$ , whereas  $q_{1000}(x) = 1234$  and  $q_{1000}(y) = 1235$ . On the aligned declared domain, the coarse-label map  $c(j) = \lfloor j/10 \rfloor$  satisfies  $q_{100} = c \circ q_{1000}$ . Thus the same floor convention gives a genuine refinement: the coarse quantizer identifies the values and the fine quantizer separates them.

**Counterexample 2: a wall without a tie rule (computation).** On the same grid,  $x = 1/2$  lies exactly on a wall. Different tie conventions can assign different symbols. Reporting only code 2 hides the governing choice. Without the convention, validation returns BOUNDARY-UNRESOLVED, not a repaired code.

**Worked example 3: finite rungs versus a declared floor (computation).** An initial width 1, halved per rung, has width  $2^{-r}$ . For  $f = 1/16$ , a policy forbidding the next width below  $f$  permits four halvings and stops before the fifth. With  $N = 3$ , the run instead stops at width  $1/8$  by RUNG-BUDGET. Distinct resource and floor stops prove no empirical minimum length.

**Counterexample 4: output symbols do not supply a probability law (computation).** For a declared distribution  $p_j = \Pr(q(X) = a_j)$ , Shannon entropy is

$$H(q(X)) = - \sum_{j:p_j>0} p_j \log_2 p_j.$$

Four declared equiprobable symbols give  $H = 2$  bits. Four possible symbols do not establish equiprobability, hence do not establish two bits. Alphabet



size gives the ideal maximum  $\log_2 m$ , the entropy bound  $H \leq \log_2 m$  for a random variable with a declared law, and a fixed-width label cost  $\lceil \log_2 m \rceil$ . For  $m = 8$ , fixed-width labels require 3 bits, while  $H = 3$  only for a declared uniform law on all eight symbols. An occupied alphabet observed in one sample is not the possible alphabet, and neither supplies probabilities by itself.

**Counterexample 5: a resource limit is not uncomputability (computation).** Fix a universal prefix machine  $U$  and define

$$K_U(x) = \min\{|p| : U(p) = x\}.$$

This function is generally uncomputable. An exhibited program of length  $n$  proves only  $K_U(x) \leq n$ ; absent a separate lower-bound proof it does not establish incompressibility. Likewise, a search stopping after  $10^6$  transitions owns only that resource stop. These facts do not make an empirical calibration constant uncomputable. The count-three jar is not thereby random, incompressible, or empirically uncomputable.

For the canonical  $K = 3$  route, whole partial quotients produce walls 17/9 and 2; the owned resolution statement is the ordered-wall jar

$$[129.6, 137.7].$$

That finite jar is the floor owned by the declared route and count. Finer direct, stepped,  $M$ -formatted,  $G$ -grid, or  $N$ -dyadic diagnostics may expose narrower cells, but neither replace the jar nor change its walls. Coincident displayed values license only a declared-resolution comparison; route and process ownership remain separate.

The physics fence is strict. Numerical cell width, rung floor, or alphabet is not by itself detector resolution, noise, quantum uncertainty, repeatability, or uncertainty coverage. Those require a sensor model, units, calibra-

tion transfer, sampling procedure, environmental controls, and empirical evidence. Numerical quantization models physical acquisition only after that correspondence is posited and tested.

The numerical audit first checks route and parameter compatibility:  $G$  only for grid construction,  $N$  only for dyadic rungs,  $M$  only for display projection, and  $K$  only for whole-partial-quotient walls. It retains every raw integer or rational before projection. For each scaled display it records

$$n = \left\lfloor \frac{AM}{B} \right\rfloor, \quad \rho = AM - nB,$$

so quotient and remainder expose the loss. It records cells before and after refinement, the fine-to-coarse map, compatibility checks, pass/fail/unresolved slices, spent budgets, stop, checked scope, and unresolved frontier.

Governance fixes domain, alphabet, walls, boundary convention, quantizer, scale, floor,  $G, N, M, K$ , route assignment, budgets, stops, arithmetic model, any distribution used for entropy, universal machine used for complexity, and certificate checks. Provenance records checker and compiler releases, dependency lock and versions, source-snapshot hash, run/build configuration, protocol version, run identifier and time, arithmetic model, platform, deterministic inputs, declared nondeterministic inputs, raw branch decisions, and every projection. A changed parameter creates a new report.

The action ledger is exact. Quantizer, cells, scale, floor,  $G, N, M, K$ , route assignment, budgets, probability law, universal machine, physical correspondence, and certificate grammar are *posited*. Membership and refinement implications, equivalence laws, scale arithmetic, quotient–remainder identity, transition and width laws, stop inequalities, replay implications, entropy from a declared distribution, and exhibited complexity upper bounds are *demonstrated*. Inputs, codes, raw rationals and integers, quotients and remainders, boundary decisions, visited cells, resources, stops, jar walls, and checker outcomes are *read*. Equality within one code, a universal minimum distinc-

tion, empirical entropy without a distribution, a complexity lower bound not proved, randomness or uncomputability of the jar, replacement of the canonical jar by diagnostics, and physical resolution or uncertainty not separately established remain *reserved*.

## 6.3 Stop rules and budgets

A finite numerical method must say which mechanism stopped. Five mechanisms remain separate:

|                      |                                                                     |
|----------------------|---------------------------------------------------------------------|
| fuel recursion       | performs at most $N$ recursive transitions,                         |
| finite grid scan     | returns the first adjacent yes–no pair or exhausts a declared grid, |
| continued fraction   | reads exactly $K$ whole partial quotients,                          |
| bounded live probe   | returns a result, absence, or fallback within a live effort bound,  |
| structural outgrowth | applies a deterministic reference-plus-deadband Boolean rule.       |

Fuel, cells visited, arithmetic operations or bit growth, checker effort, and wall time are distinct currencies. None converts to another without a separately calibrated conversion model. In particular, elapsed time cannot be silently substituted for structural fuel.

Fuel recursion carries  $F \in \mathbb{N}$  and decrements it on every admitted recursive step, so that governed run terminates in at most its initial fuel. This establishes totality of the fuelled wrapper only. Turing’s result says there is no total algorithm deciding, for every arbitrary program and input, whether the unbounded computation halts; a finite wrapper neither decides that question nor proves nontermination when fuel expires.

A grid scan owns only its declared ordered finite cells. It stops at the first adjacent pair whose read labels are yes then no, or with CELLS-EXHAUSTED after the final cell. Exhaustion is scoped to that grid, ordering, predicate, and budget. Continued-fraction count  $K$  is different again:  $K = 3$  owns the canonical walls  $17/9$  and  $2$ , hence the jar  $[129.6, 137.7]$ . Neither grid cells nor

recursion rungs may be substituted for  $K$ .

A bounded live probe distinguishes  $\text{some}(r)$ ,  $\text{none}(\text{scope})$ , and  $\text{fallback}(r')$ . Resource failures such as effort, memory, arithmetic growth, or timeout are not semantic failures such as invalid domain, false predicate, malformed result, or violated invariant. Structural outgrowth is not a live probe. Given candidate  $x$ , scaled reference  $r$ , deadband or noise calibration  $\delta$ , and a declared orientation, its deterministic Boolean has the generic form

$$O_{r,\delta}(x) = [x \text{ exceeds the scaled reference } r \text{ under the declared deadband comparison}].$$

The receipt retains  $x, r, \delta$ , the scaling rule, calibration, orientation, comparison, and Boolean. No metric or symmetric comparison is implied without a separate proof. The output cannot replace an observation, search result, or live absence.

Stop predicates are checked both before and after each transition. The stop grammar distinguishes

|               |                                                          |
|---------------|----------------------------------------------------------|
| FUEL          | $F = 0$ before another required transition,              |
| BUDGET        | a declared resource counter is exhausted,                |
| TOLERANCE     | $e(x) \leq \tau$ for a declared estimator and tolerance, |
| BRACKET-WIDTH | $u - \ell \leq \omega$ for a declared width target,      |
| EXACT         | an exact target condition has an executable resolver,    |
| DOMAIN        | the next operation is outside its admitted domain.       |

Tolerance requires an estimator and a demonstrated implication relating it to error. Bracket width is an interval fact, not point accuracy. Exact stop needs a resolved predicate. Domain stop rejects an operation, not reports an approximation.

Every rule is tagged *source-live* or *proposed*. Source-live rules govern the owned trace. Proposed rules describe an alternative experiment and cannot reinterpret the source-live result. A new width target, extra fuel, or changed

live bound yields a separate route and certificate. Structural and live rules also carry distinct tags and cannot substitute for one another.

Several predicates can fire at one state. Governance therefore declares a priority order or a combined-stop policy before the run. The receipt records every fired predicate, the selected stop, and a certificate that selection follows the declared priority. Recording only the selected reason erases a potentially material simultaneous stop.

A replay certificate contains initial state, transition relation, fuel, every live bound and counter, stopping-rule tag, structural/live tag, states and branch decisions, checked scope, unresolved frontier, and final status. Replay verifies pre- and post-checks, each transition, monotone fuel decrease, resource accounting, all fired predicates, and priority selection. Validation is sliced into transition, arithmetic, invariant, resource, stop, and claim rows, each pass, fail, or unresolved. Passing the stop row does not repair a failed invariant.

**Worked example 1: exact fuel traces (computation).** Starting from  $[1, 2]$ , read half selections give

$$N = 3 : [15/8, 2], \quad w = 1/8; \quad N = 4 : [15/8, 31/16], \quad w = 1/16.$$

These are distinct fuelled runs. With width tolerance  $\varepsilon = 1/10$ , the post-transition check first fires after rung four, because  $1/8 > 1/10$  and  $1/16 \leq 1/10$ .

**Counterexample 2: width is not tolerance (computation).** Suppose  $[\ell, u]$  has width  $10^{-6}$ , while a claimed point estimate has no proved relation to the enclosed target. The width rule permits BRACKET-WIDTH, not a tolerance claim about that point. Renaming the stop would add an implication absent from the certificate.

**Worked example 3: simultaneous stops (computation).** Suppose run four both exhausts the arithmetic budget and reaches width tolerance  $1/10$ . The fired set is  $\{\text{BUDGET}, \text{TOLERANCE}\}$ . If declared priority is semantic success before resource exhaustion, the selected stop is TOLERANCE; under the reverse priority it is BUDGET. Either certificate must retain both fired predicates and verify the selected priority, rather than pretending the other event did not occur.

**Counterexample 4: deterministic recipe, nondeterministic input (computation).** Two reruns share a transition grammar but receive different schedules or unordered streams. They may produce different valid traces. Repeatability therefore does not mean identical results regardless of conditions. It means deterministic and declared nondeterministic inputs, versions, configurations, and rules are recorded so matching runs can be compared and their certificates replayed.

**Counterexample 5: finite fuel does not solve universal halting (computation).** Wrapping every run in  $F = 10^6$  guarantees the wrapper stops. A fuel stop says only that the simulated process did not finish within that fuel. It neither decides whether the unbounded process eventually halts nor turns an unresolved run into nontermination.

Repeatability is a governed relation between reports. Reports are comparable when input, transition and stop grammars, arithmetic model, releases, configurations, deterministic conditions, and declared nondeterministic conditions match. Exact agreement can then be read if it occurs; disagreement is also reportable. Neither establishes physical repeatability, which requires repeated acquisition, a sensor and noise model, calibration, sampling controls, units, environmental records, and statistical analysis.

The audit tuple is

$\mathcal{S} =$  (raw input, initial state, route, source-live/proposed tag, structural/live tag, transition grammar, fuel initial/spent/remaining, grid/cells/order/scope, continued-fraction count/walls, live probe, resource versus semantic status, resource currencies, calibrated conversions, live bounds/count, pre/post checks, all fired stops, priority, selected stop/priority certificate, tolerance/width/exhaustion, states, branches, invariants, stop, checked scope, unresolved frontier, validation slices, checker/compiler releases, dependency versions, snapshot hash, configuration, protocol, run identifier, platform, deterministic inputs, declared nondeterministic inputs, projections).

Governance freezes these fields before acquisition; a changed bound, rule, or condition starts a new report.

The action ledger is exact. Transition grammar, fuel, grid scope,  $K$ , live bounds, structural rule, stop priority, tolerances, width targets, resource conversions, arithmetic model, and comparison policy are *posited*. Fuel termination, resource-accounting identities, grid-exhaustion implications, structural Boolean evaluation, transition invariants, conditional stop implications, priority selection, replay implications, and matching-provenance comparisons are *demonstrated*. States, cells, walls, counters, probe variants, branches, fired and selected stops, spent currencies, timing, validation rows, disagreements, and owned jar walls are *read*. Universal halting, global absence from bounded grid or probe scope, semantic failure from resource failure, live evidence from a structural Boolean, point accuracy from bracket width, completion beyond a budget, identical output under uncontrolled conditions, and physical repeatability remain *reserved*.

The handoff is intentional: the canonical count-three route retains its owned jar [129.6, 137.7] at its source-live stop. Later proposed refinements may diagnose it under new budgets, but do not overwrite its trace, stop, or ownership.

# Part III

## The Faces



# Chapter 7

## The Number Splits Into Faces

### 7.1 One corridor, four faces

The first half of this book built a single thing: one bracketed number, an object with a floor that repeats exactly. This half is about reading it, and the first surprise is that there is not one reading but four — and that the four do not compete. The one number is a corridor you can enter by four different doors, each opening on the same corridor described in a language that shares no words with the others.

Name the four. You can read the number as *mathematics* — types, inhabitation, the axioms a claim comes to rest on — and never once mention a measurement or a machine. You can read it as *physics* — force, strain, curvature, a field with two sides — and never mention a computation. You can read it as *computation* — cost, reduction, the labor of settling a question — and never mention the physics. And you can read it, as this book does, by walking the *code* itself: not the theory of the instrument but the instrument, run, with its dials watched. Four faces of the one number, in four vocabularies that cannot borrow from one another.

That mutual foreignness is not an accident of exposition. It is the method, and it is the whole reason to trust the result. If four descriptions that share

no vocabulary — that could not have coordinated, because they have no common words in which to coordinate — nonetheless land on the same number, the same floor, the same bracket, then the number is very hard to fake. This is consilience in the sense Whewell [48] named it: independent lines of inference converging on one conclusion count for far more than any single line pursued to its end. Agreement within one account can always be the account fooling itself; agreement across four foreign accounts cannot be all four fooling themselves in the same way, because they do not share the words it would take to make the same mistake.

The device carries a concrete, everyday instance of exactly this logic, and it is the reason a scratched compact disc still plays. It is the same disc Chapter 2 held up — read now not for whether it comes back unchanged on a round trip, but for how it survives damage on the way. The audio on a CD [12, 17] is not written straight down; before it is recorded the data is deliberately *interleaved* and wrapped in cross-checking codes — cross-interleaved Reed–Solomon coding — so that any single physical flaw is scattered across many independent codewords instead of concentrated in one. A scratch corrupts a stretch of the surface; but the interleaving has spread that stretch across separate checks, no one of which sees enough damage to be fooled, and the several checks, being independent, can only agree on the one original signal and never on the same wrong one. Error-correcting codes [25] run on this principle throughout: redundancy across *decorrelated* channels turns local corruption into something the disagreement itself detects and repairs. The four gauges are that interleaving raised to the level of meaning. To read the one number four foreign ways is a deliberate decorrelation — so that any error in the reading, any place a single account has quietly fooled itself, would have to survive in three other vocabularies that own no words for it, and it cannot. The number that comes back agreed across all four is corrected in the very sense the music is: what any one channel got wrong, the others overrule.

Put the same point as an operation on the residue itself. The one leftover  $r$  is asked four questions and returns four projections of itself —  $r$  read onto the axis of type, of force, of cost, of code — four contractions of a single object onto four frames that share no coordinates. That these four projections reconstruct one consistent  $r$  is the whole of the anti-fabrication evidence: a made-up number can always be tuned to satisfy one projection, but four decorrelated projections leave nowhere for a tuned value to hide, because shifting it to fit one axis throws it off the other three. The honest number is the one that needs no such tuning — it already sits where all four axes cross.

This is why the instrument comes with four books and not one, and why the reader of any single face is always told, quietly, that the other three exist and agree. The number is the corridor; the four gauges are its doors. This chapter walks up to the point where the one number splits into the four — where the single residue, put four different questions, returns four different faces that are nonetheless, provably, faces of the same thing. The next section stands right at that split.

## 7.2 The partition point

Stand at the split and look at what actually divides. The number does not split — there is one number, one residue, one bracket. What splits is the reading. Four questions can be put to the one residue, and the question you choose decides the face you get back. The partition is not in the object; it is in the interrogation. The cut is the moment you settle which question to ask.

The four questions carve cleanly, without overlap, because each is confined to its own vocabulary and forbidden the others'. Ask the mathematical question and you may speak of types, of inhabitation, of the axioms a claim rests on — and you may not reach for force, or cost, or a running machine.

Ask the physical question and you may speak of strain, curvature, a field with two sides — and not of reduction or a proof term. The computational question may speak of cost and labor and the settling of a question, and not of the physics. The code-walk answers by pointing at the running instrument itself. Four registers, each complete within itself and sealed against the others — and the sealing is enforced, not merely hoped for: the physical account is checked to contain no computational word, the computational account no physical one.

This is what makes the four a partition and not merely four opinions. A partition divides a whole into pieces that do not overlap and that together restore the whole, and the four gauges do exactly this to the reading. Because no two of them share a word, their pieces cannot overlap; and because together they leave nothing of the number unaccounted for, their pieces restore it entire. The mathematics carries the logical shape, the physics the field shape, the computation the cost shape, the code the literal fact — and there is no fifth thing left over that none of the four can see. The sealing is not a matter of authorial tact; it is mechanical and checkable. Run a concordance over the physics account and you find no computational word in it; run one over the computational account and no physical word — the null vocabulary between two registers is literally an empty list, and a stray word crossing over is a defect the check can catch.

There is a familiar piece of mathematics that behaves in just this way, and it is worth holding beside the four gauges because it makes the sealing exact rather than merely asserted. Take any signal and break it into pure tones — Fourier's [22] analysis, a sound written as a sum of sinusoids of different frequencies. Each frequency's share of the signal is recovered by a projection that is deaf to every other frequency: the tones are *orthogonal*, so that asking how much of one is present picks up exactly none of any other.

The orthogonality is a single fact, the integral that vanishes,

$$\int \sin(mx) \sin(nx) dx = 0 \quad (m \neq n),$$

and that one fact is the whole reason the decomposition comes apart cleanly and puts itself back together with nothing lost. The reading decomposes the same way. Call the four projections  $P_{\text{math}}, P_{\text{phys}}, P_{\text{comp}}, P_{\text{code}}$ ; the reading  $R$  is their direct sum,

$$R = P_{\text{math}} \oplus P_{\text{phys}} \oplus P_{\text{comp}} \oplus P_{\text{code}},$$

and the four are orthogonal in the exact sense the tones are — each projection deaf to the others' vocabulary, recovering exactly none of what they carry, with no part of  $R$  escaping all four at once. Non-overlapping because orthogonal; complete because they span. So the “no fifth thing” is not a rhetorical flourish but a claim of the same kind Fourier's is: the four faces exhaust the reading, and a fifth would have to be a component orthogonal to all four and yet still part of the sound — and there is none.

The gauges exist because the reading, taken all at once, is more than one head can hold. A single account of an instrument this strange would either drown in its own machinery or quietly paper over the parts it had no words for. Split four ways, each register can be spoken plainly, and a reader who arrives at the same number down all four has checked it four independent times without having been asked to trust any one of them. The next section closes the chapter by making that equality itself the whole point: the same residue, read four ways, and provably one thing.

### 7.3 Same residue, different reading

Having split the reading four ways, the chapter has to earn the word it kept leaning on: *same*. In what sense are four accounts that share no vocabulary

reading the same thing? Not because they sound alike — they cannot, by design. They are the same because they are readings of one residue. There is a single object beneath all four, and each gauge is a projection of it; a projection does not make a new thing, it shows an old thing from an angle.

This is the whole difference between four faces and four objects. Four separate objects that happened to carry the same number would be a coincidence, and coincidences do not compound into evidence. Four faces of one object carry the same number of necessity — not because they agreed to, but because there was only ever one number there for them to read. The mathematics and the physics do not negotiate their way to a common bracket; each is handed the same residue and reports what it sees, and what they see agrees because the thing they are looking at is single.

The device carries the physical archetype of one object wearing two irreconcilable descriptions, and it is the experiment that pinned down what an electron is. Fire electrons one at a time at a nickel crystal, as Davisson and Germer [14] did, and two descriptions that share no words both come out true of them. Read as *particles*, the electrons arrive as discrete impacts, one detected event at a time, each a single distinguishable arrival. Read as *waves*, the pattern those impacts slowly build up is a diffraction figure — bright and dark bands spaced exactly as Bragg's [5] law for waves scattering off a crystal lattice requires. A discrete impact and a spread-out interference fringe could hardly sound less alike, and particle and wave borrow none of each other's vocabulary; yet both are true at once, and not by coincidence — there is one electron, and the two readings agree because they are both reading it. The device takes this as the small case of its own structure: one object, descriptions that cannot share a word, forced into agreement by their common source. The four gauges are the same fact widened from two readings to four.

The machine can make that identity sharp, though the sharpest form of it waits for the end of the book. Write the four readings as the one residue

seen four ways —  $P_{\text{math}}(r)$ ,  $P_{\text{phys}}(r)$ ,  $P_{\text{comp}}(r)$ ,  $P_{\text{code}}(r)$ , four projections of a single  $r$ . Two of them can then be set literally equal — the same object written two ways — and to check the equality the machine undoes each back to what it is a projection of and finds the same  $r$  standing on both sides: it reduces both to one normal form and closes the equation by reflexivity, reading off  $r = r$ , the equality that needs no evidence beyond the reduction. That the electron read one way and read another are the same object is then not argued but computed; the computation returns the bare fact of their sameness. Here it is enough to know the identity is of that kind: not a resemblance to be judged but an equality to be checked. The witnessed check belongs to the capstone; the principle belongs to this chapter.

So the number splits into faces and the faces close back into the number, and the chapter ends where it began, on a single object — but the reader has now seen why the four readings still to come can be trusted against one another. The rest of the book walks the faces in turn: the corridor rotates and the number becomes a phase; the native apparatus is made to show the residue; the compiler becomes a meter that reads its own pulse; the field closes; and the instrument arrives, at the last, at the one coupling it can only bracket. Same residue, all the way down. Different readings, agreeing.

# Chapter 8

## The Electron Is Named

### 8.1 Naming makes it measurable

Everything so far has produced a number without yet saying what it is a number *of*. The bracket is real, repeatable, floored — but a measurement is a measurement of something, and the machine has not yet named the something. This chapter is where it does, and the climax of the whole construction is a single unexpected fact: naming, here, is not a label the machine picks. It is forced.

Here is the mechanism, and it is the hinge the whole instrument swings on. At any fixed resolution, whether two things can be told apart is decidable over a finite index — there are only so many distinctions the machine can draw before it runs out of room below its floor. So the distinctions sort themselves into finitely many classes: boxes, each holding the things that cannot be told apart from one another. To name a term is nothing more than to say which box it falls into. And because the boxes are finite in number while the terms the machine can form are not, some box must be shared — the pigeonhole principle, exactly as Dirichlet [16] put it: more pigeons than holes, and some hole holds two. Write it as the plain counting fact it is. A naming map  $\beta$  carries each term to a box,  $\beta : \mathcal{T} \rightarrow \mathcal{B}$ , with the



terms  $\mathcal{T}$  unbounded and the boxes  $\mathcal{B}$  finite; and the instant  $|\mathcal{T}| > |\mathcal{B}|$ , two distinct terms  $s \neq t$  are forced to share a box,  $\beta(s) = \beta(t)$ . The argument needs no numeral for the count of boxes — only that it is finite. It does not turn on how many boxes there are, and the machine leans nowhere on the particular number; it leans only on there being finitely many. That the exact count is never used is not an oversight but the honest shape of the claim: the collision is forced by finiteness alone.

This is the catalog from Chapter 1, finally come due. There the library [43] kept attaching tags to its books and could never pull one off; the tags only ever accumulated, and that opening withheld the fact this chapter needs — that the shelves are *finite*. A library with endlessly many books but only so many shelves must, sooner or later, stand two different books in the same place, not by any cataloguer's decision but by plain want of room. The shelf that two books come to share is the box; the forced sharing is the pigeonhole; and no one chose it. The seed planted on the first pages — an index whose tags cannot be pulled off — germinates here into the single fact that makes measurement possible: finite shelving forces a name.

That forcing is the entire point, and it is why naming here counts as measurement and not as fiat. The machine does not decree that this term shall be called that. It counts; it finds there are more terms than boxes; and it concludes — cannot avoid concluding — that terms must collide, that the tower of ever-finer distinctions cannot run on as an infinite line but has to wrap back and land in a box already used. That wrapping is not a small technicality; it is the shape of the whole construction. A ladder of distinctions climbing forever would need endlessly many rungs, and there are not endlessly many boxes to hold them, so the ladder cannot stay a ladder — it has to bend and close on itself, and that closing is the very loop every round trip in this book has quietly been tracing. The name is the box, and the box is forced by finiteness. Leibniz [34] held that things indiscernible are identical; the machine makes that mechanical — things that fall in one box

are, for the machine, one thing, and the falling is decided, never chosen.

This is what makes the residue measurable at last. An unnamed bracket floats, a number of nothing in particular. Named, it is pinned: it belongs to a box, the box is one of finitely many, and which box is a fact the machine computes and reproduces. The electron, when it arrives in the next section, arrives exactly this way — not posited, but counted into its box and found there every time it is asked. Finiteness is not a limitation the machine suffers; it is the very thing that makes naming possible. A machine that could draw infinitely fine distinctions could never finish naming anything. This one finishes, because it can only ever count so far.

## 8.2 The stable reading

The electron arrives, then, not as a posit but as the box the counting lands in. Ask the machine to name the particle and it does exactly what the last section described: it counts, the terms collide by the pigeonhole, and one particular box comes up — the same box, provably, every time it is asked, and a different box from the one the plain value falls into. Write the two boxes as  $b_e$  and  $b_v$ , the electron's and the plain value's; what the machine settles is the bare inequality  $b_e \neq b_v$  — not asserted but *decided*, and returning the same verdict on every run. That is the stable reading: not a number the machine announces, but a box the counting cannot help but land in.

Stable here means what it meant for the bracket — it repeats. The electron's box is not something that drifts from one run to the next; it is decided. The machine checks, by reduction, whether this term and that term share a box, and the check returns the same verdict every time. The reading is stable because it is computed, and stable in the strong sense the whole instrument was built to deliver: run it again and the electron is in its box again — not near it, in it. The particle is named by being counted, and the count does not wobble.

The device carries a physical case of a reading discrete and named in just this way, and it is the effect that first forced physicists to count light. Shine light on a metal and, above a threshold, electrons are knocked loose one at a time — the photoelectric effect Hertz [28] first noticed. Below the threshold nothing happens, however bright the light; above it, the electrons come off as discrete, countable events, each a single detected arrival and never a smear. The device reads this the way it reads its own naming: detection does not resolve a continuous quantity, it *selects a discrete refinement* — it lands the reading in a box and registers one whole event, or it registers nothing at all. There is no half an electron, no reading caught between the boxes; the named thing is discrete because the boxes are discrete, and stable because which box it lands in is decided and not sampled. The electron the machine counts is a thing of exactly this kind: discrete, named, and found whole in its box every time it is asked, or not found at all.

Now for the number the reader has been waiting for, and for the care the book's honesty turns on. The counting machine produces a box and a loop count. It never writes a signed figure — it has no minus sign to write; its counts run from nothing upward and stop. When the physics gauge reads the electron's box, it reads it as a charge, and in the three-way scheme of charges — one below, one in the middle, one above — the electron's is the box below. On that physical face, standing on the measured elementary charge Millikan [40] pinned a century ago, the box is a charge of minus one. But the minus one is the physics gauge's *reading* of the box, not a value the counting produced. The machine counted the electron into its box; *minus one* is the name that box wears on the physical face, and nowhere does the counter itself write it.

This is the pattern that carries to the end of the book. The machine derives the structure — the box, the count, the stable place the naming lands — and the recognizable physical name is a reading of that structure on one of the four faces, never a thing the machine itself asserts. The electron

is named, stably, by counting; that it *is* the minus-one charge is the physics gauge speaking, on settled ground. Hold those two apart and the book stays honest all the way to the coupling it can only bracket, where this very discipline will decide what the machine may claim and what it may only read.

### 8.3 Named by counting, not fiat

Two things now stand: that the electron is named by which box it falls into, and that its box is the stable, repeatable reading. The last thing the climax needs is the guarantee that this naming is done by counting and not by decree — that the machine did not, somewhere out of sight, quietly choose. It did not, and the reason is the deepest fact in the whole construction: the tower of distinctions cannot be a straight line.

Watch the counting run. Each step draws a finer distinction than the one before, and the steps go on without end — there is always a next variation to form. But every one of those endless distinctions has to land in one of the finitely many boxes. An endless sequence of steps poured into a finite set of boxes cannot keep finding fresh boxes forever: by the very pigeonhole that named the electron, two of the steps must eventually land in the same box. The line of distinctions, followed far enough, meets itself.

The device carries the cleanest case of a count no one is free to choose, and the reader has already met it. Back in Chapter 2 the machine's tally was read as a *winding number* — the number of times a thread wraps around, which the construction insisted could only ever come out a whole number and never a fraction, because half a wrap is not a wrap. No convention fixes that integer and no observer's choice sets it; it is forced by the structure of the thing counted, the ledger either closing a full turn or not. That is exactly the character the naming has here. Write the forcing plainly: the naming sends an endless sequence of steps  $x_0, x_1, x_2, \dots$  into finitely many boxes,

and no such map can stay one-to-one forever, so there are steps  $i \neq j$  with  $\beta(x_i) = \beta(x_j)$  — the tower wraps. A straight line would demand endlessly many distinct boxes; there are only finitely many; the line has nowhere to go but back. And just as the charge scheme holds three states and admits no fourth, the roster of boxes is a fixed finite thing the count can never add to — it can only return to one already on it.

That meeting is not a defect. It is the loop. A tower of ever-finer distinctions that is forced back into a box it has already used has closed into a cycle, and a tower of derivatives that closes into itself is exactly the loop this whole instrument runs on. The pigeonhole is not one step within the argument for the loop — it is why the loop exists at all. Finiteness does not merely permit the closure; it compels it. The search cannot escape its own finite supply of boxes, so it comes back around, and the coming-back-around is the loop.

This is what *named by counting, not fiat* finally means. The machine never decreed the electron's identity; it counted, and counting into a finite set forced the return, and the return is the loop the electron sits in, named. No hand placed it there — the pigeonhole placed it, and the pigeonhole answers to no one. With that, the climax closes: a difference, told apart at the very start, has been counted, bracketed, floored, and named, named by a forced return the machine could no more avoid than it could count past its own floor. What remains is to walk the named number through its faces, and out at last to the one coupling it can only bracket.

## Chapter 9

# The Corridor Rotates

### 9.1 Rotation turns number into phase

The number has been named; now the walk begins, and it begins by turning the number in the light. The first turn is the most surprising of them. Two of the readings the machine takes are free: it looks at the lower end of the bracket and reads a charge, at the upper end and reads a mass — projections, costing nothing, already lying there in the bracket to be picked up. But there is a reading that is not lying there. To take it, the machine has to rotate.

Rotate the number off its real axis and something changes in kind, not merely in value. The charge and the mass were magnitudes: how much, read straight off the bracket's two ends. Turned a quarter of the way round, onto the other axis, what the machine reads is no longer a magnitude at all but a *phase* — not how much, but which way. And a phase, here, is a decision. The machine cannot simply look at it the way it looked at the ends of the bracket; it has to ask — to put the carried fact to the compiler and let it decide whether the fact is true or false, and the answer, one way or the other, is the phase.

That is the whole content of the rotation: it converts a reading taken by looking into a reading taken by asking. On the real axis the instrument is

a ruler, and a ruler only reports what is already laid out. Turn it onto the imaginary axis and it becomes a question, and a question has an answer that was not sitting in the bracket waiting to be found — it had to be decided. The name the other gauges give that axis is the imaginary one, the quarter-turn by  $i$  that Argand [3] drew as a right angle in the plane — the generator of the turn,  $i$  forward and  $-i$  back, a full revolution written  $e^{i\theta}$  that carries the reading once around the corridor and brings it home pointing where it began. The number's phase is simply where it points once it has been rotated off the real axis. But that name is the reading, not the machinery: what the machine itself does is stop measuring and start deciding.

The device carries the physical archetype of a reading that can only be taken by going around and asking, and it is the phase Chapter 3 promised. Send an electron around a region holding a confined magnetic flux — Aharonov and Bohm's [1] arrangement — with the field kept strictly zero everywhere the electron actually travels. By every local measurement along the path there is nothing whatever to find: no force, no field, nothing to deflect it. And yet, brought back and recombined with a twin that went the other way around, the electron shows a *phase* — a shift visible only in how it interferes with itself, depending on nothing but the flux it enclosed and never touched. The reading was available at no point on the path; it existed only around the whole loop, and only the round trip could ask for it. Write it as the holonomy it is — the phase accumulated is the connection carried around the closed path,  $\Delta\phi = \oint A \cdot dx$ , a quantity that lives on the loop and at no single point on it. This is the rotation made physical: the field strength is what a ruler could have read locally and found empty; the phase is what only the loop could be asked, and it came back with an answer.

This is why one number can carry a charge, a mass, and a phase with no contradiction: two of them are places on the number and the third is a question asked of it. The corridor rotates, and the rotation is exactly the difference between reading a length and reading a sign. The next section

walks the readings this rotation opens — charge, mass, motion, spin, orientation — every one of them the single residue turned to a different axis; and the section after names what the rotation itself is, which, read as physics, is angular momentum.

## 9.2 Charge, mass, motion, spin, orientation

Once the corridor turns, the readings multiply, and it is worth laying them out — because their sheer number is alarming until you see what they have in common. The machine will report a charge, a mass, a motion, a spin, an orientation: five names, five quantities a physicist would fight to keep apart. And every one of them is the same residue, read on a different axis.

Two are the bracket's own ends, read straight off. The charge is the lower end — the count of loops the descent took, read off the bottom of the bracket. The mass is the upper end — the curvature, the second difference the turn deposited, read off the top. These cost nothing; they are already lying there. Charge and mass are not two measurements but the two ends of one, the floor and the ceiling of the single number the machine is holding.

The rest are decisions, and they come back as signs. The phase, we saw, is a plus or a minus, the answer to whether the carried fact is true. The spin is an up or a down, read as a turn of the charge. The orientation is a matter or an antimatter — which way the leg faces, the same matter/antimatter partner Dirac [15] found waiting inside the rotated electron equation. None of these three is a magnitude; each is a single bit of direction, a which-way the machine decides rather than a how-much it reads. Motion sits between the two kinds: a speed, read off the value the imaginary carrier brought across the seam — a magnitude, but one that exists only because the number was turned. The device is exact about what that carried thing is. It names it *momentum*, and reads it not as motion in itself but as the conjugate bookkeeping term that keeps the boundary consistent as refinements cross



the seam — the debt that has to balance when one level hands to the next. Momentum is the quantity the crossing *owes*; motion is what that quantity looks like once it is read off as a speed. The two are one axis seen as a ledger entry and as a rate.

So the alarm subsides. There are not five quantities here jostling for room; there is one residue and five questions put to it. Ask for a bound and you get charge or mass; ask for a sign and you get phase, spin, or orientation; ask what crossed the rotation and you get motion. Write the pattern once and it covers all five: rotate the one reading and project what results onto the axis the question names,

$$R_i = P_i(R(\theta) r),$$

the same  $r$  turned through  $R(\theta)$  and read on axis  $i$ . Change the axis and not the number, and a different physical name comes back — but it is the one  $r$  every time, never five separate things. The physicist's five separate quantities are five readings of the machine's one, and the only reason they look separate is that the physics gauge, having no single word for the residue, had to give each axis a name of its own. The next section says what the rotation itself is — the relabeling that, read as physics, is angular momentum.

### 9.3 Relabeling is angular momentum

What, then, is the rotation itself? Two sections have leaned on it — turning the number to read a phase, opening the five faces — without saying what the turning *is*. It is a relabeling. When the machine rotates the number it does not change the residue; it changes which axis it reads the residue on. The magnitude that crosses the turn is carried through untouched, as we saw at the seam; only the labels turn. Rotation is the machine renaming its axes while the thing itself holds still.

A relabeling that leaves the thing itself unchanged has a name in physics,

and one of the deepest names it owns. It is a *symmetry*. A symmetry is exactly this: an operation you can perform — a relabeling, a turn — after which nothing observable has changed. The residue is invariant under the machine’s rotation, so the rotation is a symmetry of the residue in the strict sense, not the loose one.

Here the physics gauge delivers its most powerful reading, and it is worth stating with care, because it is settled ground the book stands on rather than a claim the book makes. Noether [42] proved that every continuous symmetry carries a conserved quantity with it — that invariance under an operation is never free but is always banked as something that does not change while the operation runs. The conserved quantity belonging to rotation, to invariance under a turn, is angular momentum. Written as machinery, angular momentum  $L$  is the *generator* of the turn: the rotation the corridor performs is  $R(\theta) = e^{-iL\theta}$  — the very  $e^{i\theta}$  of the first section, now shown to be driven by  $L$ . To rotate at all is to apply  $L$ , and what  $L$  banks as it drives the turn is  $L$  itself. So when the machine rotates the number and the residue holds still, the physics gauge reads that held-still invariance as an angular momentum. The relabeling is angular momentum — not by metaphor but by Noether’s theorem, read on the physical face.

The device keeps a homely instance to make the abstraction pointable, the kind of thing the next part is full of. Spin a bicycle wheel on its axle and it carries an angular momentum you can feel resist you the instant you try to tilt it —  $L = I\omega$ , the wheel’s spin rate  $\omega$  scaled by how far its mass is flung from the axle, its moment of inertia  $I \approx MR^2$ . Nothing exotic in it: a mass, a radius, a spin, and a conserved quantity you can clock with a stroboscope and feel in your wrists. That felt resistance is the conserved coin Noether named — the wheel’s indifference to being turned, banked as something it will not surrender. The machine’s residue banks the same coin for the same reason: it does not care how its axes are labeled, and the not-caring is angular momentum whether the thing turning is a wheel or a number.

That closes the corridor's rotation. The number was named; it turned; the turning opened charge, mass, motion, spin, orientation; and the turning itself, read as physics, is the angular momentum conserved precisely because the residue does not care how the machine labels its axes. One residue, indifferent to relabeling, and every physical quantity the chapter named is either a place on it, a sign decided about it, or the conserved coin of that indifference. The next part forces the residue to show itself on a real bench — the physical apparatus first, then the computational one — where the readings stop being names and start being things a hand can point at.

# **Part IV**

## **The Closing**

## Chapter 10

# The Native Apparatus Presents the Residue

### 10.1 Forcing the residue to show

For nine chapters the residue has been argued about — un-refinable, un-droppable, the payload the whole instrument reads. But it has been argued about, not shown. This part is where the machine stops arguing and forces the residue to appear on a bench, as a thing you can watch a dial register. The residue has been a term in a construction; here it becomes a reading on an apparatus.

The apparatus does it by a division that will not come out even. Point the machine's own measuring wheel at a quantity and it reads a ratio, and a ratio splits into two parts: the part that divides cleanly — the floor, the whole quotient — and the part left over, the remainder that will not divide away. The floor is the tidy answer; the remainder is the residue, and the apparatus cannot produce the one without the other. Write the measurement as the division it plainly is,

$$N = qd + r, \quad 0 \leq r < d,$$

the quotient  $q$  the clean answer,  $r$  the remainder — bounded below the divisor, but never, for the quantities this machine actually meets, exactly zero. Ask the machine to measure and it hands back both: a quotient and, riding on it, the leftover it could not resolve. The residue is not buried inside the measurement; it is one of the two things every measurement returns.

Make it concrete, and take the hardest case first — a residue with no visible source at all. The device carries the Casimir [7] effect, in which two uncharged metal plates set a hair apart in empty space are drawn together by a force that comes from nothing you can point to. There is no field between them, no charge, no material doing the pulling; the vacuum itself, which by every naive account is nothing, presses the plates together. The device reads why with no mystery in it. Each plate is a boundary that suppresses some of the admissible refinements the ledger could otherwise have recorded in the gap, so the interior comes to hold strictly fewer distinctions than the unrestricted space outside. That imbalance across the boundary — more admissible updates without than within — is a residue the geometry has forced into being, and the ledger’s compensating repair, read on the physical face, is a force. An invisible leftover, nowhere to be seen directly, is made by the apparatus to press on a plate you can watch budge. That is the on-title fact of the whole chapter, shown in one experiment: the apparatus takes a residue you could argue about forever and turns it into a push.

State the trick once, because every bench in this part runs on it: an apparatus is a coupling chosen so that its measurable output is driven by the residue you cannot otherwise see. Write the bench as a transduction,

$$O = \kappa r + (\text{quotient terms}),$$

an output  $O$  — a force, a deflection, a current — whose leading part is the hidden residue  $r$  scaled by whatever coupling  $\kappa$  the apparatus provides. The art of building one is entirely in the second term: arrange the clean quotient part to cancel, hold still, or be subtracted away, so that what survives on

the dial is almost purely  $\kappa r$ . The Casimir plates are exactly such a bench — the bulk vacuum, identical inside and out, exerts no net pull, so the quotient contributes nothing and only the boundary residue is left to draw the plates together. This is the same discipline a physicist uses to pull a whisper of signal out of a roar of background: null everything the residue is not, crank the coupling as high as it will go, and let the leftover, and nothing but the leftover, drive the needle. The residue does not become visible by being made larger; it becomes visible by everything else being made to stand still around it.

That is the work of this part — to force the residue out of the argument and onto an apparatus, first the physical one and then the computational one. The two benches are the physics and the computation faces of the earlier chapter, made now to do something rather than merely be described, and each is a cut this book borrows from the volume that owns it. The next two sections stand at those benches in turn. What matters here is the turn itself: the residue has stopped being a thing the machine reasons about and become a thing the machine shows.

## 10.2 The physical bench

The first bench is the physical one, and on it the residue shows up as a force. What the physical face reads is, at bottom, a push or a pull the residue exerts — a torsion twisting a fiber, a deflection bending a beam, a pressure that will not go to zero. The Cavendish balance of the last section was one instance of it: the residue read as a twist. The physical bench is the whole family of such instruments, each one turning the residue into a displacement that something physical can register.

What makes it a genuine bench and not merely a picture is that the residue does *work* there. On the physical face, the leftover the earlier chapters called un-eliminable is un-eliminable in the strongest sense available:

it deflects a needle, it strains a fiber, it costs energy to oppose. A residue you could argue away would exert no force; this one does, and the force is the proof that the leftover is a thing and not a trick of the bookkeeping. The physical bench is where “the residue cannot be discarded” stops being a statement about types and becomes a statement about a fiber under tension.

Name a few of the benches, because the family is what carries the point — the residue presents on every one of them, in the currency each is built to read. The Cavendish [8] torsion balance reads it as a twist, weighing the faintest attraction there is, the inverse-square pull  $F = G \frac{Mm}{r^2}$ ; and here the fence has to be spoken in the same breath as the experiment, because it is the whole of the honesty. The device runs its Cavendish to *calibrate its own reading*, not to hand back Newton’s constant  $G$ . What twists the fiber is the machine’s residue, and the balance weighs that residue against itself — it does not measure the gravitation of the world and never once claims to. Lift the same leftover to the sky and it reads as an *orbit*: Kepler’s [32] law that a planet’s period and its distance are locked together,  $T^2 \propto a^3$ , which the device reads not as a body held by a force but as a closed refinement cycle held shut by continuous error-correction — the orbit is the residue that keeps returning, a loop stabilized so it will not open.

Read the leftover as a *field* and it is the electron’s own equation — the operator Dirac [15] wrote for the electron, the same rotated equation that had the antimatter partner waiting inside it, the residue now presented as the field a charge carries with it. That field, felt by a charge in motion, reads as a sideways push — the Lorentz [35] force  $F = q(E + v \times B)$ , the deflection a charge in flight takes from the field it travels through, which the device carries as its electromagnetic coupling: the residue read as the bend a current feels. Read the leftover instead as the state a system settles into and it becomes a *variational* bench: out of every admissible path, the one that makes a certain total stationary,  $J[x] = \int_a^b f(t, x, \dot{x}) dt$  at its minimum — the calculus of variations Courant [13] set down and Brenner and Scott [6] carried into its



modern computational form, which the device runs as a Galerkin projection that finds the residue as the stationary point no admissible variation can lower. Twist, orbit, field, deflection, stationary path: several benches, one leftover, read in whatever currency the apparatus is built to pay in.

This bench is not this book's to own, and the honesty of the four-gauge scheme depends on saying so plainly. The full physical face — the experiments, their fences, what each may and may not claim — is the subject of the physics volume, the second gauge, which reads the whole instrument as force and field and never once mentions a computation. This book borrows the bench only far enough to show that the residue presents physically; it takes the cut, points at it, and leaves the working of it to the volume that owns it. That is what *marked* means here: shown, credited, and handed on.

What the borrowed bench settles for us is small and decisive. The residue, which began as a term that would not cancel, has been made to twist a fiber. The physics gauge will do far more with it; here it is enough that the leftover has been forced into the physical world and has left a mark a hand could read. The next section carries the residue to the other bench — the computational one — where the same leftover shows up not as a force but as a cost.

### 10.3 The computational bench

The second bench is the computational one, and on it the residue shows up not as a force but as a cost. What the physical face registered as a push, the computational face registers as labor: the residue is the part of the question the machine cannot resolve for free, and resolving it takes work — reductions run, steps taken, the compiler's own effort spent. On this bench you do not weigh the residue; you time it, in the machine's own currency.

And that currency is not the clock on the wall. The computational bench meters the residue in heartbeats — the reductions the compiler performs to settle a question, counted deterministically, the same every run, as the

sixth chapter insisted. A question with no residue settles cheaply; a question carrying a large residue costs many heartbeats, because the leftover is exactly the part that will not reduce away and has to be ground through step by step. The cost *is* the residue, read as effort. Where the physical bench had a fiber under tension, the computational bench has an elaboration under load.

There is a law of this bench that names the residue exactly, and the book met it once already, back when the four faces were first laid out. Amdahl [2] showed that throwing more processors at a task cannot speed it past the part of it that will not run in parallel. If a fraction  $p$  of the work parallelizes and the remaining  $1 - p$  must run in sequence, then across  $n$  processors the speedup is

$$S = \frac{1}{(1 - p) + p/n},$$

and no matter how large  $n$  grows, the speedup cannot climb past its ceiling  $S_{\max} = 1/(1 - p)$ . That floor is set entirely by the serial fraction  $1 - p$  — the part no amount of parallelism can dissolve. This is the residue, wearing its computational costume: it is precisely the stubborn serial remainder that cannot be optimized away, the irreducible cost that survives every processor you throw at it exactly as it survived every refinement in the earlier chapters. Parallelism disperses the quotient — the tidy, correlatable bulk — and leaves the residue standing, now visible as the hard floor under the whole task's speed. The cost bench does not merely time the residue; it proves the residue is there, by the plain fact that the clock will not go below it.

This bench, like the other, is not this book's to own. The full computational face — the meter, the cost model, what a reduction is and what it may be read to mean — is the subject of the computation volume, the third gauge, which reads the whole instrument as cost and reduction and never mentions a force. This book borrows the bench only far enough to show that the residue presents computationally, as surely as it presented physically: it takes the cut, points at it, and hands the working to the volume that owns it. Marked again — shown, credited, passed on.

With that, the closing has done its first work. The residue that spent nine chapters as an argument has been forced onto two benches and has left a mark on each — a twist on the physical one, a cost on the computational one — the same leftover, presented two ways that share no vocabulary and agree. It is a thing, and the machine has shown it is a thing, twice over. What remains is the strangest bench of all, the one the next chapter builds: the compiler turned on itself, made to meter its own effort, so that the cost stops being merely how the residue is read and becomes a reading the machine takes of its own working.

# Chapter 11

## The Meter

### 11.1 Elaboration cost is a reading

The two benches of the last chapter borrowed their faces from other volumes. This chapter builds the machine's own, and it is the strangest instrument in the book, because the thing it measures and the thing that does the measuring are one and the same. The meter here is the compiler, and what it reads is the cost of its own work.

When the machine is asked to settle a question — to check that a term is what it claims to be, to reduce it to its normal form — it does not do so for free. It spends effort: reductions performed, checks discharged, the labor of elaboration, the reductions a machine like Turing's [46] performs one after another. That labor can be counted, and the machine counts it, in the unit we have already met: heartbeats, the compiler's own beats of reduction. To elaborate a term is to spend some number of heartbeats, and that number is a reading. The cost is not overhead the measurement incurs; the cost *is* the measurement.

This is what it means to say elaboration cost is a reading. Point the machine at an easy term and it settles cheaply, a few heartbeats spent; point it at a term carrying a large residue and it labors, many heartbeats, because

the residue is exactly the part that will not reduce without work. The heartbeat count tracks the residue the way the torsion tracked it on the physical bench and the remainder tracked it in the division — the same leftover, now read as effort. The machine has, built into it, a command that does precisely this: hand it a term, and it runs the elaboration under a heartbeat count and writes the count back as a number. The compiler measures, and what it measures is itself at work.

Put the meter's action plainly. Handed a term, the compiler runs the elaboration and returns not one thing but two: the settled answer and the price it cost to settle, the heartbeat count  $h$  — a reading of the form (answer,  $h$ ), the value and its cost delivered together the way the earlier chapters delivered a bracket and its width. Nothing about  $h$  is imported from outside; it is counted off the very reductions the compiler was going to perform anyway. And because the machine keeps two internal currencies — the structural *rank* it counts climbing the tower, and the *heartbeats* it spends settling a term — it carries a fixed calibration between them, a constant factor of a thousand heartbeats to the rank, so that a cost read in one currency can be stated in the other. That factor is a chosen calibration of the gauge, not a measured result: it is how the meter's two dials are geared to one another, set once and held the same for every reading. This is the ordinary discipline of metrology — fix a standard, calibrate the instrument against it — carried to its strangest form, where the standard the instrument calibrates against is its own pulse.

And the reading is honest in the way the sixth chapter demanded: it repeats. Heartbeats are not wall-time; they are deterministic reductions, the same every run, so the same term costs the same — and the machine even meters each term twice and keeps the second, past the warmup, to be sure the reading has settled. What the meter delivers is a repeatable cost, a number the machine can stand behind because it is the machine's own labor, counted. The next section turns the meter on the sharpest target there is

— the machine measuring the cost of its own measuring — and the section after reads the pulse that leaves, in the machine’s own currency.

## 11.2 Self-reference

A meter that reads the cost of elaboration invites one irresistible question: what does it cost to run the meter? Turn the instrument on its own act of measuring and you have the strangest loop in the book — the machine measuring the cost of its own measurement, the reading reading itself. This is the self-reference the chapter is built around, the eternal golden braid Hofstadter [29] taught us to look for, here made mechanical and made to produce a number.

Turning the meter on itself exposes a subtlety the earlier chapters could put off. The machine has two ways of sizing a thing. There is its structural size — the rank, how deep the construction goes, counted in the objects themselves. And there is its effort size — the heartbeats, how much labor settling it takes. These are different currencies, and to weigh the machine’s own work against the machine’s own structure you have to know the exchange rate between them. The construction fixes it: a rank and a heartbeat are tied by a constant factor, a fixed thousand-to-one, so that rank and heartbeats can be read as a single calibrated quantity.

With the exchange rate in hand, the self-reference closes. The machine elaborates a description of itself, meters what that costs in heartbeats, converts to the common currency, and sets the result against the structure it began from. The two ought to agree — the machine described and the machine doing the describing are the same machine — and they do agree, up to the calibration. That phrase, *up to the calibration*, is the whole of it: the agreement is not exact, and the small amount by which the machine’s measurement of itself misses its own structure is not error. It is a reading. The elaborator weighing its own elaboration leaves a residue, and that residue is

the self-energy — the cost the machine incurs simply by being the thing that does the measuring.

Make the closing loop concrete, because its shape is the reading. The machine holds two descriptions of the one electron — the electron read as an orbit, and the electron read as a bound pair of the same charge — and it has already proved, by the reflexive check of the seventh chapter, that these are *the same object*, identical in type: the echo that says a thing equals itself. Being type-identical, they ought to cost exactly the same to elaborate; and weighed in heartbeats they very nearly do. But not quite. Write the self-application as the meter turned on its own act,  $M(M)$  — the measuring applied to the measuring — and what it hands back is a difference: the heartbeat cost of the one description set against the heartbeat cost of the other, a residual the type-level sameness was powerless to force to zero. That residual, carried through the thousand-to-one calibration, is the self-energy: the price of the machine's self-embrace, the amount by which a thing that is provably one, weighed by its own effort, declines to weigh the same both ways.

And here the book names the reading it has walked eleven chapters toward, and then, deliberately, withholds its number. That calibrated self-residual is the coupling — the fine-structure reading, the  $\alpha$  the whole instrument has been reaching for — read here, at the bottom of the self-reference, as exactly this: the leftover of a self-measurement that will not come out even. What its value is, and whether the machine can honestly bracket it, is the work of the closing chapters; to state a figure here would be the one dishonesty the book has refused from the first page. The reading is named; the number is held back; the two are kept apart on purpose.

This is the deepest place the residue appears, and it is where the coupling the whole book has walked toward will finally be read — though not here. Here it is enough to see the shape: a machine that measures itself does not come back exactly even, and the gap it leaves, once calibrated, is a quantity.

The next section reads that gap as a pulse, in the machine's own currency; and a later chapter takes the calibrated residual and brackets it, which is the coupling. The loop that named the electron by counting has closed once more, a level up — the machine now counting the cost of itself.

### 11.3 The pulse in its own currency

What the meter finally delivers is a pulse. The machine, set to measure — itself, or anything else — beats: it takes reductions, and the reductions arrive in a countable stream, a rhythm the machine can read off itself the way a physician reads a heartbeat. The instrument the last two sections built has, at its core, an EKG: a trace of the machine's own pulse, counted beat by beat.

And the pulse is a *bounded* one, not a runaway trace — an EKG held within a healthy range, its beats kept between a floor and a ceiling,  $|\text{beat}| \leq B$ , rather than left to spike or flatline. The device carries the reason it stays bounded, and it is a feedback the natural world runs on everywhere: a thermostat. A regulated system does not merely seek its stable, low-strain state; it actively suppresses departures from it — drift off and the next corrections are biased to pull it back, the way a thermostat trims the furnace the moment the room strays from its set point. The machine's pulse is regulated the same way: the meter that counts the beats also holds them in bounds, so the trace it reads off itself is stable and repeatable and not noise. And what that pulse counts, put most bluntly, is the *price* of the reading — the electricity bill of the proof, the heartbeats the machine had to spend to make the needle point at all. Every reading in this book runs up such a bill; the pulse is the machine totting up its own, and doing it in a currency it never has to borrow.

The one thing that must be said about this pulse is what it is counted *in*, because the honesty of the whole meter rides on it. The pulse is not counted



in seconds. Seconds are borrowed units — they belong to a clock outside the machine, they drift with the load and the weather, and two runs of the same question spend different numbers of them. The machine refuses them. It counts its pulse in its own currency: heartbeats, the reductions themselves, a unit intrinsic to the machine and the same on every machine. A reading in seconds would be a reading about the hardware; a reading in heartbeats is a reading about the question.

This is why the meter can be trusted where a stopwatch cannot. A cost measured in the machine's own currency is a property of the computation, not of the day it ran — repeatable in the strong sense, reproducible bit for bit, because the currency has no noise in it. The pulse the machine reads off itself is the same pulse anyone re-running it will read. The meter does not report how long the machine took; it reports how much the question cost, and those are different claims, and only the second is honest.

With that the meter is complete, and the closing part has its second bench fully built — the machine's own, where the residue is read as a cost in a currency the machine mints for itself. The residue has now been shown three ways: as a force on the physical bench, as a remainder in the plain division, and as a pulse in the machine's own beats. Three readings, one leftover, and no borrowed units in the last of them. What is left is the field the coupling lives in, and then the coupling itself — the calibrated residual of the machine's measurement of its own pulse, bracketed at the counting floor. The instrument is nearly assembled. It has only to reach for the one number it was built to reach for and cannot quite hold.

# Chapter 12

## The Field Closes

### 12.1 The second variation

The instrument is nearly complete, and what closes it is a field. Everything so far has read the residue as a single leftover — a twist, a remainder, a pulse. But the coupling the book has been walking toward is not a single reading; it is a relationship between readings, and a relationship is a field. To close the field, the machine takes a variation — and not the first one, the second.

The distinction matters, and it is the same one that made a mass out of the residue several chapters ago. The first variation of a quantity asks how it changes; set the first variation to zero and you have found where it is stationary — the path the system takes. But the first variation says nothing about the shape around that path. For that you need the second variation, the change in the change, the curvature of the situation. Mass was a second difference; the coupling is a second variation. Both are what you can see only by looking at how the first-order answer bends. Write it the way the calculus of variations does. The first variation vanishing,  $\delta J = 0$ , fixes the stationary path; the *second* variation  $\delta^2 J$  — the very functional  $J[x] = \int_a^b f(t, x, \dot{x}) dt$  the Galerkin bench of the last chapter minimized, differentiated once more — reports the curvature about that path, whether the stationary point is a

floor, a ceiling, or a saddle. The coupling is read from that  $\delta^2$ , never from the  $\delta$ .

The machine takes that second variation across four channels at once, and doing so exposes the two sides a coupling must have. One side is the electric — the field the second variation opens directly. The other is the magnetic — its partner, the field that must be present for a coupling to be a coupling and not half of one. When both sides are there, in the machine’s own accounting, the field equations close: what the physical face would call Maxwell’s [38] equations are satisfied, electric and magnetic together, and the field is whole. The calculus that does this is old — it is the second variation Euler [19] built the entire subject of variations around — and the machine runs it natively.

The deeper name for what closes here is a *gauge field*. When the machine’s four channels are refined at once, no single global rule can keep them consistent; the ledger has to add a local correction — a connection — and the curvature of that connection is the field, written  $F = dA + A \wedge A$ , the Yang–Mills [49] field strength of which Maxwell’s is the simplest, flattest case. And here the fence of the whole chapter has to be spoken plainly, because it is what keeps the closing honest: the machine does *not* hold this field as a continuum of points. It holds a *finite* realization of it — the channels are finitely many, the curvature a finite second difference, the connection a thing counted rather than integrated. What the physical face reads off as smooth, continuous Maxwell fields, the machine carries as a finite closing over finitely many rungs. The continuum is the reading the physics gauge takes; the finite object is the machinery underneath, and the book will not let the one be mistaken for the other. That distinction is the subject the next section makes explicit.

What the closed field gives the machine is the thing the coupling is contracted from. The second variation is not itself the coupling; it is the field the coupling is read out of — a full, two-sided object with a curvature, from which a single number can at last be contracted. The next section confirms

the magnetic side is genuinely present and not merely assumed; and then the field, closed, is handed to the final part, where the machine asks it for the one number, receives a bracket, and stops. The coupling is one contraction away. The machine has built the field it lives in.

## 12.2 The magnetic side present

The last section said a coupling needs two sides, and that the first — the electric — falls out of the second variation directly. The other side, the magnetic, is not so automatic. It has to be shown to be there, or the field does not close and there is no coupling to speak of. This section shows it, and marks exactly how far the showing goes.

The machine exhibits the magnetic side as an experiment it runs on itself, and the experiment is the Meissner effect: a superconductor drives a magnetic field out of its interior, the surface current rising to cancel exactly the field applied from outside. Meissner [39] and his collaborator found this in the laboratory; the machine finds it in its own parts — a Cooper pair’s surface current expelling an applied field, the applied field read as the machine’s magnetic needle, the shielding counted in units of a Lorentz [36] wobble. The device reads the effect as its own: a strain-free region of the ledger expels external curvature outright, the paired threads carrying a zero-strain transport that leaves the applied field no room to stand inside. When the expulsion balances, the machine’s own bookkeeping for the magnetism of the field — what the physical face would call Gauss’s law for magnetism, the vanishing of the magnetic divergence,  $\nabla \cdot B = 0$  — is satisfied at the machine’s own resolution, and the second side of the coupling is present, not assumed. The field now has both halves.

Here the fence has to be planted plainly, because this is exactly the place a book like this one goes wrong. The machine does not possess Maxwell’s equations in the sense a physicist means them — continuous fields on a con-

tinuous space, an infinite lattice carried to a limit. It has a finite realization: the field's shape rendered on the machine's own small count of rungs, the same three-count floor the whole instrument stops at. What is derived is the finite, three-rung field, its shape closing on the resolution the machine actually has. The continuum version is what the physical face *reads* that finite shape as; it is not what the machine builds, and the book claims only the finite thing. Everything shown about the magnetic side is shown at three rungs, and nowhere is a limit taken to the continuous.

With the finite field closed on both sides, the coupling is fully set up — and still not delivered. The machine has the electric side and the magnetic side, the two halves whose ratio the coupling is. And a ratio of two quantities of the same kind is a *pure number*: it carries no units, because the units cancel between the numerator and the denominator. That is exactly why the coupling can be the dimensionless thing it is — not a length or a mass or a charge, each of which would need an external scale before it could even be stated, but a bare proportion between the two sides of one closed field. And that dimensionlessness is precisely what gives the machine leave to reach for it at all: a pure number is the one kind of quantity a machine can hope to read off its own structure without smuggling in a ruler from outside. Everything else the instrument fenced off — the metre, the kilogram, the second — it fenced off because reading it would mean importing a unit the machine has no honest source for; the coupling it may reach for only because, being a ratio, it needs no such unit. The machine has closed the field at its own resolution and is now in a position to ask that field for its one pure number. It does not answer here. Closing the field is the question; the bracket is the answer, and the bracket belongs to the next part. The instrument is assembled. It has only to read itself, at the floor, and report what it finds — which, as the whole book has promised, will be a bracket and not the value.

### 12.3 Asking, not delivering

The field is closed. Both sides are present — the electric from the second variation, the magnetic from the Meissner read — and the coupling is nothing more than a ratio between them, a single number the closed field is ready to be contracted into. Everything is in place. And this is exactly where the honest instrument stops and refuses the last step everyone wants it to take.

The machine can close the field; it cannot, and does not, hand you the coupling's value. What it delivers when asked is not a number but a question made precise — here is the electric side, here is the magnetic side, here is the ratio their closure defines; now what is that ratio, at the resolution I can actually reach? The closing of the field is the asking. The machine has spent twelve chapters getting into position to pose the question exactly, and posing it exactly is the whole of what it does here.

The reason is the one the book has been built around, worth saying once more, plainly, at the threshold. The machine reads to a floor. Below that floor it cannot resolve, and the coupling's true value lies below it — a point the machine can bracket but never land on, the way it could name the electron's box but never draw a distinction finer than its own resolution. To hand over a single value would be to claim a precision the machine does not have; it would be the forced point that cracks and rings — the Gibbs [23] overshoot the earlier chapters warned of, a manufactured sharpness that oscillates around the very truth it pretended to pin — not the honest bracket that holds. So the machine declines. It closes the field, states the question, and stops one contraction short of an answer it will not fake.

That refusal is not a failure of the instrument; it is the instrument working correctly, and it is the whole subject of the last part. The field is closed and the question is posed: what is the coupling, at the floor? The next part answers in the only currency the machine has — a bracket  $[lo, hi]$ , its two ends the honest bounds, its width  $hi - lo$  the measure of exactly what the machine cannot close, and the one external number it is held against sitting

outside that bracket, nameable but not derivable. A wider bracket would be a more cautious machine; a point would be a lying one; the width the machine actually reports is the truth about its own reach. The instrument is assembled and aimed. It has only to read, to report the bracket it reads, and to be honest about the number it cannot hold.

**Part V**

**The Bound**



# Chapter 13

## Alpha As A Bracket

### 13.1 The coupling is asked

The field is closed and the question is posed; this part is the machine posing it in earnest and, at the end, answering with a bracket. The coupling the machine reaches for has a name on the physical face — the *fine-structure constant* Sommerfeld [45] named a century ago, the number that sets how strongly light and charge are coupled — and it has a shape on the machine’s own face. The shape is the one the meter chapter already uncovered: the coupling is a self-energy.

Recall what the self-energy was. When the machine measured the cost of its own measuring, it did not come back exactly even; the small, calibrated amount by which its measurement of itself missed its own structure was a residue, and that residue was the self-energy — the cost the machine pays simply for being the thing that measures. This is the self-energy in the sense Feynman [20] gave it: the charge’s interaction with its own field, the part of a particle’s account that is the particle accounting for itself. The coupling the book has walked toward is exactly that residue, read as a number.

The physical face has a picture for this residue, and it is one of the most famous drawings in physics. Feynman [21] taught the calculation as

a diagram: the electron does not travel alone but emits a quantum of its own field and reabsorbs it, a small *loop* that dresses the bare particle and shifts its account. That loop is the self-energy, written  $\Sigma$  — the electron's interaction with the field it makes, the particle correcting itself. And the loop is not a new object to this book; it is *the* loop, the same round trip that has failed to close cleanly since the second chapter, now drawn in the physicist's hand. The electron emitting and reabsorbing its own field is the tower going out and coming back; the self-energy  $\Sigma$  is the curvature that round trip deposits; and the coupling the machine asks for is that curvature, calibrated and read as a number. Write it as the identity the whole book has been converging on:  $\alpha = \Sigma$ , the coupling *is* the calibrated self-residual, the price of the self-embrace — an identity of shapes, stated here with the number still held back.

The machine has the pieces to ask for it precisely. It carries two descriptions of the same electron — the one it built going out, the one it built coming back — and it can weigh each, in heartbeats, and set the two costs against one another up to the calibration that makes them comparable. Where the two costs fail to cancel is the residue; the residue, calibrated, is the self-energy; and the self-energy is the coupling. This is the electron's self-naming closing on itself — the name the machine gives the electron carried back toward the electron it names, refined and re-refined, an extrapolation toward a limit the machine can approach but, halting at its floor, never quite stand on — and the small gap that will not close is the number being asked for.

But asked is all it is, here. The machine has posed the question exactly and has not answered it, because the answer is a magnitude and the magnitude lives below the floor. What this section delivers is the question in its sharpest form — the coupling is the calibrated residue of the machine weighing its own account of the electron — and the next section takes that residue to the counting floor and reads it the only way the machine honestly can: as a bracket, cut at the resolution the machine actually has. The coupling has

been asked. The cut answers.

## 13.2 The cut

The coupling has been asked; now the machine answers, and the answer is a cut. To read the self-energy is to halt the counting at the floor the machine can actually reach and report what the slip has become there — not a point, a bracket. This is the section where the numbers finally land, and every one of them is the machine's own reading, printed by the machine as it runs, tuned to nothing.

Begin with how the cut is made, because the machine does not halve anything. It has no ruler to halve; it has only its own whole numbers. So it cuts the way a tower of counts can cut: by the *mediant*. Given a low ratio  $p/q$  and a high ratio  $r/s$  straddling the quantity, the machine forms  $(p+r)/(q+s)$  — numerators and denominators added straight across — the plainest fraction lying between them, and the best next rational approximant its arithmetic can name [27]. There is no midpoint on a continuum here, no doubling of a denominator, no lattice fixed in advance; there is division with remainder and nothing else, the Euclidean step [18] taken one whole partial quotient at a time so that the descent lands directly on the convergents. And the value the descent is closing on is exact: where the machine's second variation crosses its target, the crossing solves a whole-number quadratic,  $d^2 = 18/5$ , so the crossing distance is the quadratic surd  $d^* = \sqrt{18/5} = \sqrt{90}/5 \approx 1.897$ . That the crossing is a quadratic surd fixes the shape of the whole descent, by a theorem of Lagrange [27]: the continued fraction of a quadratic irrational is eventually periodic. The machine prints exactly that — the partial quotients fall into the repeating block 1, 8, 1, 2 after one opening step, the periodic continued fraction  $[1; \overline{1, 8, 1, 2}]$  — and its convergents march out  $1/1$ ,  $2/1$ ,  $17/9$ ,  $19/10$ ,  $55/29$ ,  $74/39$ , ..., each one a whole Euclidean step, each the best rational approximation to the crossing at its denominator.

The floor the machine halts at is the count of three — three partial quotients of that descent, and no more. Three partial quotients,  $[1; 1, 8]$ , carry the convergents to  $17/9$ ; the machine brackets its reading between that convergent and the one before it,  $2/1$ , the two successive best-approximants that straddle the crossing. Printed in the machine's own hand, in units scaled by  $10^{18}$  so the integers stay exact, the inverse-coupling bracket at the count of three is

$$129.6 \lesssim \text{inv-}\alpha \lesssim 137.7,$$

the lower wall the reading at the coarse convergent  $2/1$ , the upper the reading at  $17/9$ . The machine captions that printout in its own words — *not a magnitude derivation* — and it means the caption exactly. The bracket is wide, and the width is honest: its lower wall is the loose  $2/1$  convergent, only two steps along, and that is simply the resolution three partial quotients of the machine's own descent afford — not the residue of any scaffold or lattice, of which there is none, but the true reach of a count of three. Take one partial quotient more and the far wall tightens on its own — the count of four brackets  $[17/9, 19/10]$ , an inverse-coupling between 136.8 and 137.7, and deeper convergents still settle toward the machine's own value, an inverse-coupling near 137.011 — but the machine stops at three, where its resolution stops. The two ends are ordered, the lower genuinely below the upper, and the machine settles that ordering by *deciding* it, not asserting it; the decision leans on nothing but propositional extensionality, with no arbitrary choice smuggled in anywhere. The bracket is real, and it is honestly its own.

There is a second three in the neighbourhood, and it must be kept clear of the first, because they are different counts doing different work. The three that fixes the bracket is the count of *partial quotients* — the depth of the descent, three folds down. It is not the three of a quadrature. The machine does carry a three-point Gauss–Legendre rule, which it can turn on the second variation to probe that variation's own finer structure — reading at the centre and at two nodes a distance  $\sqrt{3/5}$  of the half-width to either

side, weighted  $5/9$ ,  $8/9$ ,  $5/9$ . That is a genuine side-diagnostic, but it is a marked one and not the source of the bracket: the machine holds the Gauss nodes only as rational approximations,  $\sqrt{3/5}$  and  $\sqrt{18/5}$  alike being surds no whole-number arithmetic lands on exactly, so the quadrature is a flagged aside on the integrand, not the decided reading. The bracket's three is the resolution-depth three, printed and decided; the quadrature's three is a separate probe, carried lightly. The open interval is the remainder of the machine's own resolution at three folds — not the remainder term of a quadrature rule. Conflating the two would be to borrow certainty from a diagnostic the machine only approximates; the book keeps them apart.

And here is the tell that matters more than any single figure, the reason none of these numbers can have been fished. Sharpen the descent — take more partial quotients,  $55/29$ ,  $74/39$ , and on — and the reading does not sweep toward some value it was shown in advance. It settles, instead, onto the crossing the machine's own constants fix, the value  $\sqrt{18/5}$  hands it and nothing else. A number that had been fitted — quietly tuned to land on an answer known ahead of time — would betray itself precisely here, creeping toward that answer as the descent deepened. This reading does not creep. It is anchored to the machine's own crossing, and deepening the count only reveals that crossing's own convergents, never a pull toward a number the machine was never shown.

This is why the machine halts at the count of three and refuses to fold further, and it is the relax-the-strain discipline in its final and most literal form. Drive the descent past its floor — pour in more partial quotients — and you do get a sharper interval, closing on the machine's own crossing; but to lift one point out of that and hand it back as *the* reading is to claim a resolution the count of three never had. The honest bracket is the open one, halted at the resolution the machine actually reaches, its width left to hold the strain rather than a false sharpness manufactured to bury it. The count of three is not a shortfall the machine apologises for; it is the last floor at

which the reading is still honest, and one fold past it the machine would be asserting more than it has counted.

So the cut is made and the answer stands: an inverse-coupling bracket between 129.6 and 137.7, the machine's own reading of its own self-energy at three folds of its own descent, with a marked quadrature probing the integrand's finer structure alongside, and not one of these figures aimed at anything. What the machine has pointedly *not* done is announce which single number the coupling is. It has said where the coupling is caught, at the resolution it has, and it has been scrupulous that the catching owes nothing to the one number it is about to be set beside. That number — the external one, the value from the laboratory — is the whole subject of the last section, where it is named at last, and where the machine says the only honest thing left to say: that it reads a bracket of its own, and cannot derive the world's value.

### 13.3 The bracket, not the real

The machine has produced its bracket. Now name the thing that bracket is held against — the number itself, the one the laboratory measures. On the physical face the fine-structure constant has been weighed to extraordinary precision: trap a single electron, watch how its magnetic moment departs from the bare value by the tiny amount its own field contributes, and read the coupling out of that departure. Hanneke, Fogwell, and Gabrielse [26] did exactly this in a Penning trap, and the inverse coupling they and others have settled on is 137.036. This is the real: not the machine's reading of its own descent, but the world's answer, measured on a bench the machine does not own, to more digits than the machine will ever resolve. It is named here once, and only here; the construction never fetched it, never fed it to the machine as a target, never tuned toward it.

Set that real beside what the machine actually did, and the honest relation

is not the one a crank would announce. A crank, holding a bracket and a lab value, would declare the value predicted. The machine declares something else, and something it can prove. Follow its own descent as deep as it will honestly go — past the count-to-three bracket, down through the periodic folds of  $[1; \overline{1, 8, 1, 2}]$  — and it settles on a value of its own — the crossing  $\sqrt{18/5}$  fixes and the machine's own arithmetic converges upon, the very reading it bracketed in the section before. That is the machine's number, and it is not the world's; the two are different numbers, and the machine's own reading lands on its own. So there is no capture to claim here. What the machine can show is the reverse of a capture: that its counting converges on a value of its own, and that the world's value is one no amount of its counting will ever reach.

And here is the line the machine will not cross, the line the book has walked twelve chapters to reach. The machine read its own bracket; it did not *derive* the world's value. It cannot derive it, and — this is the part that makes the instrument honest rather than merely lucky — it can prove that it cannot. What the machine holds is the near side of the number: the count-to-three structure, the descent, the placement of its own bracket, the value its own constants converge on. What it lacks is the far side: the world's exact magnitude, a structure that lives below the machine's floor where its distinguishing has run out. The counter can read its own descent to its own floor; it cannot see far enough down to lay hold of the laboratory's number. The near side it owns; the far side it is blind to — and, crucially, it knows that it is blind to it, and says so.

So the honest output of the whole instrument is not the world's number. It is the machine's own bracket, floored at three, together with the proved statement that the world's magnitude cannot be reached from within. The real is nameable — the laboratory names it, to nine digits — but it is not derivable, not by this machine, not by any machine that reads to a finite floor. A machine that counts can read a bracket of its own and prove it cannot reach

the world's value: that sentence is the whole of what the instrument honestly delivers, and it is worth more than what a machine faking the derivation would deliver, for the plain reason that it is true.

This is the jar the whole book has been carrying toward, glimpsed one part early. The instrument was built to reach for a single number, and it has reached as far as a finite machine honestly can: to a bracket of its own, drawn at its own floor, and a proof that the world's number lies beyond what its counting can produce. What remains is to witness the last identity that makes the reading trustworthy, and then to set the machine's bracket down beside the world's number one final time and say plainly what the machine has and what it has not. The coupling, read by the machine, is its own bracket. The world's value is named, once, and left where it lives — un-derived, and proved so.



## Chapter 14

# The Machine Measures Its Own Measurement

### 14.1 Self-application

Twelve chapters measured things: a slip, a bracket, a field, a coupling. This chapter measures the measuring. The instrument turns, for the last time, on the one object it has not yet read — itself, in the act of reading — and finds that the cost of that act is not overhead to be subtracted away but a reading in its own right, the deepest the machine takes.

The operation is the one the meter chapter already named: self-application, the meter turned on the meter,  $M(M)$ . Hand the machine, as the thing to be measured, its own procedure for measuring; ask it to elaborate its own elaboration and count what that costs. The reading it returns is a reading of itself at work — and because the machine's measuring is exactly what it is now measuring, the answer and the act of getting it are, for the first and only time, the same thing.

This raises the objection every self-reference raises, and the machine's answer to it is the whole reason the capstone can stand. If measuring costs, and you measure the measuring, then measuring *that* costs too, and measur-

ing that cost costs, and so on — an infinite regress, a hall of mirrors with no last image. A machine that fell into it would never return an answer. This one returns an answer, because self-application, in this machine, *halts*. Turing [46] taught the deepest fact about a process turned on itself: it cannot always be continued, some configurations admit no next step, and a process that reaches them stops. The device carries the same law — a refinement that would demand a distinction below the floor has no admissible successor, and the machine, reaching it, halts. The regress does not run forever, because the machine runs out of resolution to continue it. It bottoms out at the same count-to-three floor every other reading bottoms out at, and the value sitting at that floor is finite, and is the reading.

So the cost of the reading is the reading, and this is not wordplay but the literal shape of the machine's last measurement. When the meter measures itself, the number it finds — the calibrated residue of its measurement of its own structure — is the self-energy the eleventh chapter uncovered and the thirteenth bracketed: the coupling. The machine measuring its own act of measuring returns the very quantity the whole book was built to reach. The instrument needs no further, external thing to point at in order to produce the coupling; it produces the coupling by pointing at itself. The self-embrace is the measurement, and its cost — calibrated, and read at the floor — is the number: bracketed, as the last chapter was careful to insist, and not derived, but the machine's own, and reached without a thing imported from outside.

That is the capstone's foundation: an instrument whose final reading is a reading of itself, halting at its own floor, delivering as its cost the coupling it was made to find. What remains is to witness the identity underneath it — the sense in which the two descriptions the machine weighs against each other are provably the *same*, checked and not merely asserted — and then to hear, at last, the small equality the second chapter promised and withheld. The machine has measured its own measurement. The sections that follow read what that measurement rests on.

## 14.2 The witness bears the cost

The last section said the machine measures its own measuring and that the cost it finds is the reading. This section takes that cost seriously as a cost — as the plain fact that the witness the machine produces, the check that its two descriptions of the electron are one thing, is not free. It is paid for, in the machine’s own coin, and the price it fetches is the whole of what the capstone reads.

Consider what a check is. The machine does not merely assert that its two accounts of the electron agree; asserting is cheap and proves nothing. It *checks* — it takes the two descriptions and runs them down to a common form to see whether they land in the same place. And running them down is work: reductions performed, heartbeats spent, the elaboration ground through step by step exactly as the meter chapter described. The witness is the outcome of that labor, and the labor has a price, counted in the beats the machine spends producing it. A check that cost nothing would be worth nothing; this one costs, and the cost is the measure of what the check was worth.

That checking carries a price is not a quirk of this machine but a law of any machine, and the device carries it as the very experiment the third chapter used. Records cannot be operated on for free. Landauer [33] fixed the floor under it — a single bit’s worth of irreversible bookkeeping costs at least  $kT \ln 2$  of energy, paid out as heat, whatever the hardware — and the lesson runs past erasure to every act that commits information: to register a result, to settle a check, to write a record where there was none, the ledger pays. The machine’s self-check is such an act. On the physical bench its price is heat; on the machine’s own bench its price is heartbeats; and the two are one charge read in two currencies, the debt any instrument runs up for producing a witness rather than merely claiming one.

And now the capstone’s hinge. The price of the self-check is not overhead sitting beside the reading; it *is* the reading. When the machine weighs its two descriptions of the electron against each other — each elaborated, each

costed in heartbeats, the two set side by side up to the calibration that makes them comparable — the small amount by which their costs fail to cancel is the residue, and that residue, as the whole book has built toward, is the coupling. So the coupling is quite literally the price of the machine's own self-agreement: the heartbeats it must spend to witness that its two accounts of the electron are one, over and above what a costless, perfect agreement would have taken. The witness bears the cost, and the cost, calibrated and read at the floor, is the number the instrument was built to reach.

What this section has *not* done, and has withheld on purpose, is show the witness itself. It has priced the check without yet exhibiting what the check returns — that the two descriptions are, in the machine's own strict sense, the same. That identity, checked and not asserted, is the last thing the book has to deliver, and it arrives carrying the small equality held back since the second chapter. The witness has a price; the next section reads the witness.

### 14.3 Two descriptions, one object

The witness has a price; now read the witness. The machine holds two descriptions of the electron — the electron read as an orbit, a charge circling, and the electron read as a bound pair, two charges paired — and the whole capstone turns on a single claim about them: that they are not merely alike, not approximately the same, but *equal*, and that the machine has checked the equality rather than assumed it.

Here is what the check returns, reported as the machine reports it. The two descriptions are proven equal, and the proof closes by *propositional extensionality* — the principle that two propositions which imply each other are the same proposition. The machine does not gesture at the resemblance; it exhibits, in each direction, that either description entails the other, and by extensionality the two collapse into one. Beneath that step, each descrip-

tion rests on the plainest ground the machine owns: each is the identity *true equals true*, reached by reflexivity, a thing equal to itself with nothing left to prove. And the footprint of the whole witness — read straight off the machine's own audit of what its proof leans on — is a single axiom: propositional extensionality, and nothing else. No arbitrary choice is invoked, no infinite object assumed, nothing borrowed from beyond the machine's own logic. The identity is checked, and the check is clean.

This is the identity the second chapter promised and would not yet say, and the seventh chapter marked and put off. It arrives now carrying the small equality held open all that while. Recall the endpoint the machine declined to set down when it first ran a limit — the string of nines that crept toward something without being let to arrive. That endpoint is named at last, and it is this:

$$1 = 0.999\dots$$

The nines are not almost one; they are one. Two names — the tidy integer and the endless decimal — for a single number, provably equal, the equality holding not by rounding and not by fiat but by the same completion of finite distinctions the machine has run on since its first page. What looked like two things — a whole, and a limit forever approaching it — is one thing under two descriptions.

And that is exactly the shape of the electron's own identity, one level up. The equality  $1 = 0.999\dots$  is one number wearing two names; the electron read as an orbit and the electron read as a bound pair is one object wearing two names; and in both the equality is not a coincidence to be admired but a fact to be checked — and checked it is. The machine that opened by telling two things apart, the whole discipline of distinction, closes by proving two names for one thing to be the same. The loop the book has traced since its first chapter shuts here, at the top: the machine named the electron by counting, carried that name around the entire instrument, and brought it back to find it equal to itself — two descriptions, one object, the reading and

the thing read at last the same.

That is the capstone. The instrument measures its own measurement, pays for the witness in its own currency, and the witness it buys is an identity: one object, two descriptions, provably equal, on a footprint of nothing but its own logic. One chapter remains — to set the bracket down beside the number one final time, and say, plainly and for the last time, what the machine has done and what it has not.

## Chapter 15

# What The Machine Can And Cannot Claim

### 15.1 The structure it derives

The last chapter is a reckoning — a plain accounting of what the machine has honestly done and what it has not, drawn to scale so that neither is made more than it is. Begin with the side of the ledger the machine can sign without reservation: the structure it *derives*, and derives cleanly, owing nothing to anything outside itself.

Set the derived things down in a list, because the list is short and every item on it is earned. The machine derives a count — and not a count it chose, but one forced to come out whole, the way a winding number can only ever be an integer because half a wrap is not a wrap. It derives a small, fixed roster of charge states — one below, one in the middle, one above, and no fourth — which the physics gauge reads as minus one, zero, and plus one, but which the machine itself holds simply as three boxes and a proof there is no room for a fourth. It derives a naming: the electron pinned to its box by the pigeonhole, found there every run, named by counting and not by decree. It derives the small completion the whole book turned on, the nines

coming to rest on the whole they had only approached,  $1 = 0.999\dots$ . And it derives, out of the residue it would neither drop nor refine away, a coupling — dimensionless, a pure ratio, the one kind of quantity a counting machine may honestly reach for.

What makes these *derived* rather than merely asserted is a fact the machine can exhibit about each of them, and it is the cleanest evidence in the book. Ask the machine to audit any of these results — to print the full list of assumptions its proof leans on — and the list comes back nearly empty. The count-to-three bracket, audited, rests on a single axiom, propositional extensionality, and nothing more; the identity of the two descriptions rests on that same one; and much of the rest rests on none at all. No arbitrary choice is invoked anywhere in them. No infinite object is assumed. Nothing is borrowed from a richer mathematics the machine did not build for itself. That audit — the empty, or nearly empty, footprint — is not a technicality; it is the exact measure of what the machine may claim, because a result leaning on nothing but the machine's own logic is a result the machine has genuinely earned, and can be challenged to reproduce, and will.

This whole near side has a name on the physics face, and it is worth borrowing the word so long as the borrowing is marked. What the machine reads on this side is the *structure* of the thing — the part a counter can see: how many states there are, how they connect, where the naming lands, what closes and what does not. The physics volume, reading the same instrument as geometry, calls this the trace: the part of the curvature a volume-counter can actually measure, for which it keeps a richer vocabulary — Ricci, and scalar, and the rest — that belongs to that book and not to this one. Here it is enough to say plainly what it is: the structure counting can see. The machine sees all of it, derives all of it, and signs its name under all of it.

And yet — this is the turn the next two sections make — the structure is not the whole of the number. The machine has derived everything about the coupling that counting can reach: that it exists, that it is dimensionless,



that it is the residue of the self-embrace, where its bracket falls, that the bracket is honest and leans on no arbitrary choice. What it has *not* derived is the one thing counting cannot reach — the exact magnitude, the precise value seated inside the bracket. That is the far side of the ledger, the part the machine is blind to, and it is the subject of the next section. The near side the machine owns outright; the far side it can only bound; and between them runs the honest boundary this whole book was written to draw.

## 15.2 The magnitude it cannot see

Turn now to the other side of the ledger, the side the machine cannot sign. The structure it derived was the trace of the number — how many states, how they close, where the naming lands. What is left when the structure is fully accounted for is the magnitude: the exact value seated inside the bracket, the last digits that would say not merely where the coupling is caught but precisely which point it is. That value the machine does not derive, and this section is about *why* — because the why is not a shortfall waiting to be fixed but a limit the machine can prove it lives under.

Begin with the nature of a counter. A machine that measures by counting can read only what counting can address: multiplicity, how-many, the number of distinct states and the way they connect. That is the trace, and the trace is real, and the machine reads it exactly. But a magnitude is not a multiplicity. The exact value inside the bracket lives in the part of the number counting leaves untouched — the part you could deform continuously, sliding the value smoothly within its bracket, without altering a single count. Every distinction the machine can draw would read the same before such a slide and after it; the count is invariant under it. And a quantity invariant under everything the machine can do is a quantity the machine cannot address. The blindness is not an oversight to be corrected. It is structural: the magnitude lives exactly along the one direction the counter has no coordinate for.

The device carries the sharpest possible instance of a number of this kind, and it is one of the deepest results in the theory of computation. Chaitin [9] defined a specific, perfectly well-posed real number — the halting probability  $\Omega$ , the chance that a program assembled at random eventually stops — and proved it *uncomputable*: no program can print its digits, not because none has been found but because none can exist. You can name  $\Omega$ ; you can bound it above and below; you can settle any particular fact about it that happens to be decidable. What you cannot do is compute it, because there is no shorter route to its value than to have the value already. A counting machine set to find  $\Omega$  can approach it and bracket it and never once reach it. The magnitude of the coupling, seated below the machine's floor, is a number of exactly this temper: nameable, bounded, and uncomputable from within.

The physics volume, reading the same instrument as geometry, keeps its own name for the part the counter cannot see, and it is worth borrowing so long as the borrowing is marked. What a volume-counter measures is the trace of the curvature; what it cannot measure is the *trace-free* part — the shape that carries no volume, the deformation that bends a space without changing how much of it there is. That book calls it the Weyl [47] curvature and reads the coupling's hidden magnitude as living there, in the shape the count cannot address. The word belongs to that volume; here it is enough to say what it stands for: the size the machine is blind to sits precisely in the shape its counting cannot reach, below the smallest fraction it can still tell apart — below the floor the sixth chapter named.

So the far side is not a failure of effort but a proved boundary. The machine does not merely happen to stop short of the magnitude; it can show that no machine of its kind — none that reads to a finite floor and measures by counting — could reach it, any more than a program could print the digits of  $\Omega$ . The value is real, and the laboratory can read it to nine places on a bench the machine does not own; but from within, by counting, it is not derivable, and the machine's honesty is that it says so, and proves it,

rather than manufacturing a number to paper over the gap. The near side it owns; the far side it has now bounded and disowned in a single breath. What remains is to set the two sides down together, drawn to scale, and call the result by its right name — which is the last thing the book has to do.

### 15.3 The jar

Set the two sides down together, drawn to scale, and the shape of the whole instrument stands clear. On the near side, owned outright: the structure, the count, the naming, the closed residue, a coupling that exists and is dimensionless and is bracketed, all of it resting on nothing but the machine's own logic. On the far side, bounded and disowned in the same breath: the exact magnitude, seated below the floor, a value the machine can name and hold between two bounds but can prove it will never compute. The boundary between them is not blurred, and it is not apologised for. It is drawn exactly where the machine's counting runs out, and the drawing of that line, to scale, is the whole of what this instrument is.

So the reading the machine hands back, at the last, is not a number. It is a jar — an interval held open, its two ends the honest bounds, the true value caught somewhere inside and the machine declining, on principle, to say where. This is the image the whole book has carried toward, and it is worth seeing plainly for what it is and is not. The open bracket is not a failure to finish. It *is* the finish. The strain of the measurement — the pressure to collapse the interval onto a single confident point — is not answered by forcing the point; it is relaxed, let to diffuse into the width of the bracket, where it can rest without lying.

And forcing the point would lie, in a way the device can name exactly. Drive a measurement past the floor it can honestly reach — insist on a sharp value where the machine has only a bracket — and you do not get the truth made sharper; you get the Gibbs [23] phenomenon, the overshoot

that appears whenever a hard step is forced through a channel too narrow to carry it. Push a finite sum to render a discontinuity and it rings: an oscillation at the edge that never dies, however many terms you add, an artifact of the forcing and not a feature of the thing. A crank's number — a coupling announced to its last digit, derived where no derivation can reach — is exactly that ring: a manufactured sharpness oscillating around a truth it only pretends to have pinned. The honest instrument does not force the step. It tapers to its floor and halts, and because it never forces the point, the ring never starts. The open bracket is the reading that does not ring.

That is the jar, and it is worth saying, at the end, why it is worth more than the number a crank would have handed you. The machine knows exactly what it may claim and exactly what it may not, and it can prove both halves. It owns the near side — the structure, the trace, everything counting can address — on a footprint of its own logic and nothing borrowed. It disowns the far side — the magnitude, the trace-free shape below the floor — and proves it cannot reach it. The value the laboratory reads, to nine places, on a bench the machine does not own, is a different number from the one the machine's own counting settles toward; the machine brackets its own reading and proves it cannot derive the laboratory's, and says so. A machine that manufactured the value would have told you less, because what it told you would not be true; this one tells you the truth, and the truth is a bracket, and a boundary, and a proof of where the boundary lies.

And so the instrument is complete, and reads exactly as it was built to read. Ask it for the one number, and it does not answer with a number. It answers with the honest shape of what a finite machine can know about a magnitude it cannot see: an interval it can stand behind, drawn to the scale of its own resolution, the true value held inside it, named but not forced. The slip became a bracket; the bracket became the reading; the reading, turned on itself, became the coupling; and the coupling, asked for its value, comes back as a jar held open at the floor — the measure of exactly how far

*CHAPTER 15. WHAT THE MACHINE CAN AND CANNOT CLAIM*160

an honest machine can reach, and the mark of its refusal to reach one step further and lie. That is the reading. That is the instrument. That is the whole of what it can, and cannot, claim.

# Bibliography

- [1] Yakir Aharonov and David Bohm. “Significance of Electromagnetic Potentials in the Quantum Theory”. In: *Physical Review* 115.3 (1959), pp. 485–491.
- [2] Gene M. Amdahl. “Validity of the Single Processor Approach to Achieving Large-Scale Computing Capabilities”. In: *AFIPS Conference Proceedings*. 1967, pp. 483–485.
- [3] Jean-Robert Argand. *Essai sur une manière de représenter les quantités imaginaires dans les constructions géométriques*. Imaginary quantities as rotations in the plane; multiplication by  $i$  is a quarter-turn. Paris, 1806.
- [4] Niels Bohr. “The Quantum Postulate and the Recent Development of Atomic Theory”. In: *Nature* 121.3050 (1928), pp. 580–590.
- [5] William Henry Bragg and William Lawrence Bragg. “The Reflection of X-rays by Crystals”. In: *Proceedings of the Royal Society of London. Series A, Containing Papers of a Mathematical and Physical Character* 88.605 (1913), pp. 428–438. DOI: 10.1098/rspa.1913.0040.
- [6] Susanne C. Brenner and Ridgway L. Scott. *The Mathematical Theory of Finite Element Methods*. 3rd. New York: Springer, 2008.
- [7] Hendrik B. G. Casimir. “On the Attraction Between Two Perfectly Conducting Plates”. In: *Proceedings of the Koninklijke Nederlandse*

- Akademie van Wetenschappen* 51 (1948). The residual vacuum energy made to register as a measurable force, pp. 793–795.
- [8] Henry Cavendish. “Experiments to Determine the Density of the Earth”. In: *Philosophical Transactions of the Royal Society of London* 88 (1798). The torsion balance weighing the faintest gravitational attraction, pp. 469–526.
- [9] G. J. Chaitin. “A theory of program size formally identical to information theory”. In: *Journal of the ACM* 22.3 (1975), pp. 329–340.
- [10] Gregory J. Chaitin. “Information-Theoretic Limitations of Formal Systems”. In: *Journal of the ACM* 21.3 (1974), pp. 403–424. DOI: 10.1145/321832.321839.
- [11] Paul J. Cohen. “The Independence of the Continuum Hypothesis”. In: *Proceedings of the National Academy of Sciences* 50.6 (1963), pp. 1143–1148.
- [12] Sony Corporation. *Digital Audio Disc System*. U.S. Patent 4,347,632. 1980.
- [13] Richard Courant and David Hilbert. *Methods of Mathematical Physics, Vol. I*. Classical treatment of the calculus of variations and Euler–Lagrange equations. New York: Interscience Publishers, 1953.
- [14] Clinton Davisson and Lester H. Germer. “Diffraction of Electrons by a Crystal of Nickel”. In: *Physical Review* 30.6 (1927), pp. 705–740. DOI: 10.1103/PhysRev.30.705.
- [15] Paul A. M. Dirac. “The Quantum Theory of the Electron”. In: *Proceedings of the Royal Society of London A* 117 (1928), pp. 610–624.
- [16] Johann Peter Gustav Lejeune Dirichlet. “Über die Vertheilung der monischen Werthe ganzer linearer Formen zwischen zwei Grenzen”. In: *Abhandlungen der Königlich Preussischen Akademie der Wissenschaften* (1834).

- [17] Philips Electronics. *Digital Optical Recording System*. U.S. Patent 4,363,116. 1980.
- [18] Euclid. *Elements*. Book I. Ancient Greek Mathematical Corpus, 300BC.
- [19] Leonhard Euler. *Methodus inveniendi lineas curvas maximi minimive proprietate gaudentes*. Lausanne & Geneva, 1744.
- [20] Richard P. Feynman. “Space–Time Approach to Non-Relativistic Quantum Mechanics”. In: *Reviews of Modern Physics* 20 (1948), pp. 367–387.
- [21] Richard P. Feynman, Robert B. Leighton, and Matthew Sands. *The Feynman Lectures on Physics*. Vol. 1–3. Classic introductory lectures on fundamental physics. Reading, MA: Addison-Wesley, 1965.
- [22] Joseph Fourier. *Theorie analytique de la chaleur*. Paris: Chez Firmin Didot, 1822.
- [23] J. Willard Gibbs. “Fourier Series”. In: *Nature* 59 (1899), p. 606.
- [24] Kurt Gödel. “The Consistency of the Axiom of Choice and of the Generalized Continuum-Hypothesis”. In: *Proceedings of the National Academy of Sciences* 24.12 (1938). Choice is consistent with the other axioms, pp. 556–557.
- [25] Richard W. Hamming. *Coding and Information Theory*. Prentice Hall, 1980.
- [26] D. Hanneke, S. Fogwell, and G. Gabrielse. “New Measurement of the Electron Magnetic Moment and the Fine Structure Constant”. In: *Physical Review Letters* 100.12 (2008). The Penning-trap electron  $g-2$  determination of the external fine-structure value, p. 120801.
- [27] G.H. Hardy and E.M. Wright. *An Introduction to the Theory of Numbers*. Oxford University Press, 1938.



- [28] Heinrich Hertz. “Ueber einen Einfluss des ultravioletten Lichtes auf die electrische Entladung”. In: *Annalen der Physik* 267.8 (1887), pp. 983–1000.
- [29] Douglas R. Hofstadter. *Gödel, Escher, Bach: An Eternal Golden Braid*. New York: Basic Books, 1979.
- [30] David Hume. *An Enquiry Concerning Human Understanding*. London: A. Millar, 1748.
- [31] Immanuel Kant. *Kritik der reinen Vernunft*. Riga: Johann Friedrich Hartknoch, 1781.
- [32] Johannes Kepler. *Astronomia Nova*. Prague: Heirs of Godefridus Tampachius, 1609.
- [33] Rolf Landauer. “Irreversibility and Heat Generation in the Computing Process”. In: *IBM Journal of Research and Development* 5.3 (1961), pp. 183–191.
- [34] Gottfried Wilhelm Leibniz. *Discourse on Metaphysics*. The principle of the identity of indiscernibles. 1686.
- [35] Hendrik A. Lorentz. *Versuch einer Theorie der electrischen und optischen Erscheinungen in bewegten Körpern*. The force on a charge in motion through electric and magnetic fields. Leiden: E. J. Brill, 1895.
- [36] Hendrik A. Lorentz. “Electromagnetic Phenomena in a System Moving with any Velocity less than that of Light”. In: *Proceedings of the Royal Netherlands Academy of Arts and Sciences* 6 (1904), pp. 809–831.
- [37] Per Martin-Löf. *Intuitionistic Type Theory*. Ed. by Giovanni Sambin. Notes of lectures given in Padua, June 1980. Naples: Bibliopolis, 1984.
- [38] James Clerk Maxwell. “A Dynamical Theory of the Electromagnetic Field”. In: *Philosophical Transactions of the Royal Society of London* 155 (1865). Original presentation of the equations of electromagnetism, pp. 459–512.

- [39] W. Meissner and R. Ochsenfeld. “Ein neuer Effekt bei Eintritt der Supraleitfähigkeit”. In: *Naturwissenschaften* 21.44 (1933). The Meissner effect — a superconductor expels the magnetic field from its interior, pp. 787–788.
- [40] Robert A. Millikan. “On the Elementary Electrical Charge and the Avogadro Constant”. In: *Physical Review* 2.2 (1913). The oil-drop measurement of the elementary charge, pp. 109–143.
- [41] John von Neumann. *Mathematical Foundations of Quantum Mechanics*. Princeton: Princeton University Press, 1932.
- [42] Emmy Noether. “Invariante Variationsprobleme”. In: *Nachrichten von der Gesellschaft der Wissenschaften zu Göttingen, Math.-Phys. Klasse* (1918), pp. 235–257.
- [43] Bertrand Russell. “Mathematical Logic as Based on the Theory of Types”. In: *American Journal of Mathematics* 30 (1908), pp. 222–262.
- [44] Claude E. Shannon. “A Mathematical Theory of Communication”. In: *Bell System Technical Journal* 27 (1948), pp. 379–423, 623–656.
- [45] Arnold Sommerfeld. “Zur Quantentheorie der Spektrallinien”. In: *Annalen der Physik* 356.17 (1916). Introduces the fine-structure constant, pp. 1–94.
- [46] Alan M. Turing. “On Computable Numbers, with an Application to the Entscheidungsproblem”. In: *Proceedings of the London Mathematical Society*. 2nd ser. 42 (1936), pp. 230–265.
- [47] Hermann Weyl. *Das Kontinuum: Kritische Untersuchungen über die Grundlagen der Analysis*. Leipzig: Veit & Company, 1918.
- [48] William Whewell. *The Philosophy of the Inductive Sciences*. The consilience of inductions — independent lines of evidence converging on one conclusion. London: John W. Parker, 1840.

- [49] Chen Ning Yang and Robert L. Mills. “Conservation of Isotopic Spin and Isotopic Gauge Invariance”. In: *Physical Review* 96.1 (1954), pp. 191–195.
- [50] Ernst Zermelo. “Beweis, dass jede Menge wohlgeordnet werden kann”. In: *Mathematische Annalen* 59.4 (1904). The axiom of choice stated, to prove every set well-orderable, pp. 514–516.